

CS3103: Computer Networks Practice

Network Applications & Demo

Principles of Network Applications

HTTP & WWW

E-mail Protocols (SMTP, POP, IMAP)

P2P Applications

Dr. Anand Bhojan

COM3-02-49, School of Computing

banand@comp.nus.edu.sg ph: 651-67351

<https://pollev.com/banand>

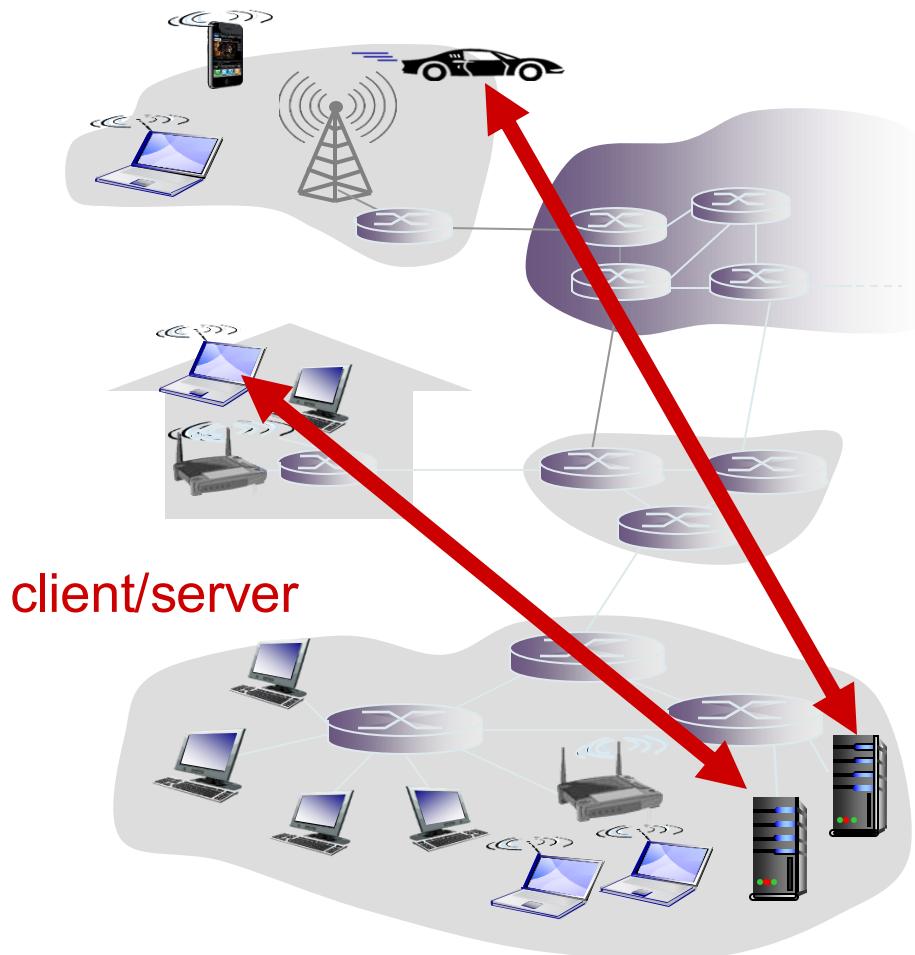
Make sure you
LOGIN using
your
NUSNET ID.



possible structure of applications:

- ▶ client-server
- ▶ peer-to-peer (P2P)

Client-server architecture



Which of the following are common properties of a server?

- ▶ always-on host
- ▶ uses ephemeral ports client
- ▶ uses permanent IP address
- ▶ initiates communication client
- ▶ uses well-known ports
- ▶ waits to be contacted

P2P architecture

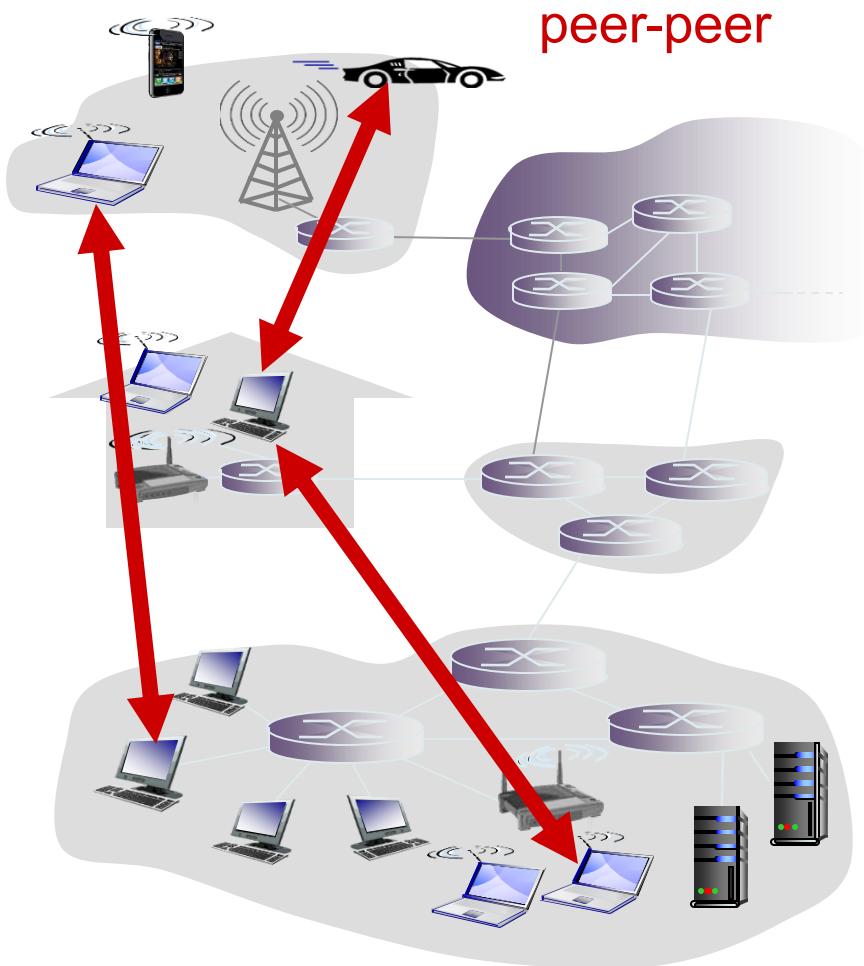
- ▶ no always-on server
- ▶ arbitrary end systems directly communicate and serve each other.
 - ▶ Each host runs both server and client processes

Q1: What is the advantage of P2P?

no need to set up

Q2: How the hosts find the IP addresses of each other?

Will discuss in next Section...



Socket

▶ Socket

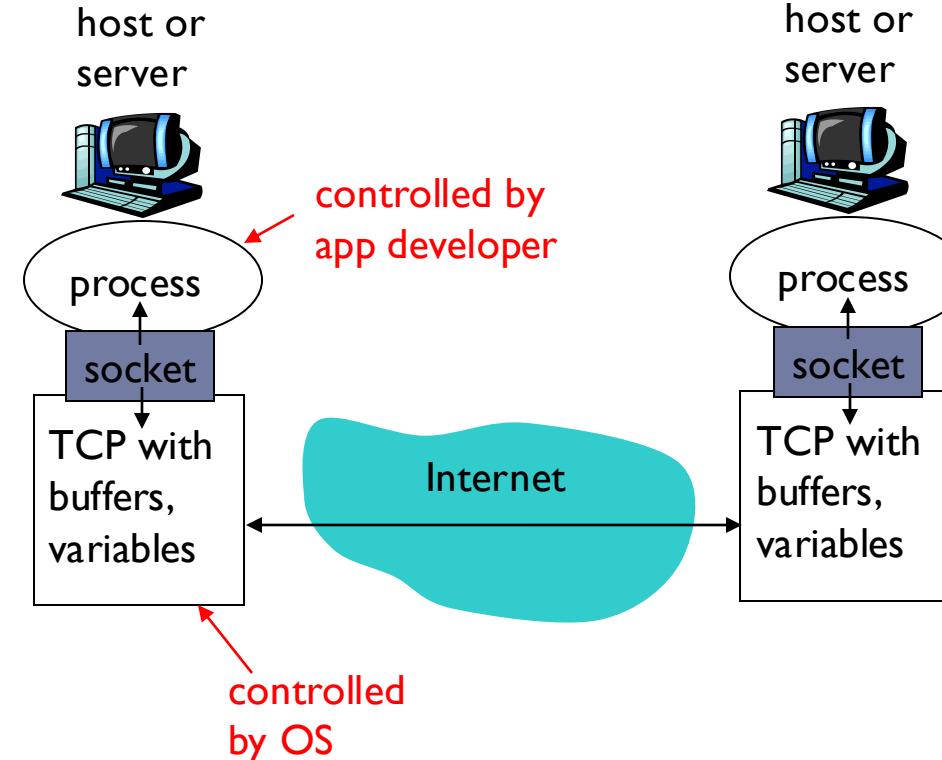
- ▶ Interface between the application layer and transport layer
- ▶ OS-controlled interface (a “door”) into which application process can both send and receive messages

▶ Addressing

- ▶ Host address + process identifier
- ▶ Eg. IP address + port number
 - ▶ Eg HTTP – port 80, SMTP – port 25, ftp – port 21

Tool:

- Use **ifconfig** (**ipconfig** in windows) to see edit your laptop’s IP address.
- Use ‘**cat /etc/services**’ to see all well known port numbers.
- Use ‘**netstat -a**’ to show all active connections.
- Use ‘**lsof -i -n**’ to show the COMMAND, PID (process ID), and USER (user ID) of open Internet sockets.



Socket: endpoint for comm. A combination of host IP address and port number.

Internet transport protocols services

TCP service:

- ▶ ***reliable transport*** between sending and receiving process
- ▶ ***flow control***: sender won't overwhelm receiver
- ▶ ***congestion control***: throttle sender when network overloaded
- ▶ ***does not provide***: timing, minimum throughput guarantee, security
- ▶ ***connection-oriented***: setup required between client and server processes, maintains ***connection state***

UDP service:

- ▶ ***unreliable data transfer*** between sending and receiving process
- ▶ ***does not provide***: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup

udp header to identify application, have port number

Q: why bother? Why is there a UDP?

Securing TCP

TCP & UDP

- ❖ no encryption
- ❖ cleartext passwords sent into socket, traverse Internet in cleartext

SSL

- ❖ provides encrypted TCP connection
- ❖ data integrity
- ❖ end-point authentication

SSL is at app layer

- ▶ Apps use SSL libraries, which “talk” to TCP

SSL socket API

- ❖ cleartext passwords sent into socket, traverse Internet encrypted

CS3103: Computer Networks Practice

Network Applications & Demo

Principles of Network Applications

HTTP & WWW

E-mail Protocols (SMTP, POP, IMAP)

P2P Applications

Dr. Anand Bhojan

COM3-02-49, School of Computing

banand@comp.nus.edu.sg ph: 651-67351



HTTP

- ▶ Uses **TCP Service.**
- ▶ HTTP is **stateless** (Note:- TCP maintains connection state, application layer is not maintaining the application/user state)
 - ▶ server maintains no information about past client requests
 - ▶ if server/client crashes, their views of “state” may be inconsistent, must be reconciled
- ▶ Default HTTP Server port: **80** [HTTPS – port 443]
- ▶ HTTP 1.0 (RFC 1945)
- ▶ HTTP 1.1 (RFC 2616)
- ▶ HTTP 2 (RFC 7540, RFC 9113)
- ▶ HTTP 3 (RFC 9114)

protocols that maintain “state” are complex!

- ❖ past history (state) must be maintained
- ❖ if server/client crashes, their views of “state” may be inconsistent, must be reconciled

HTTP

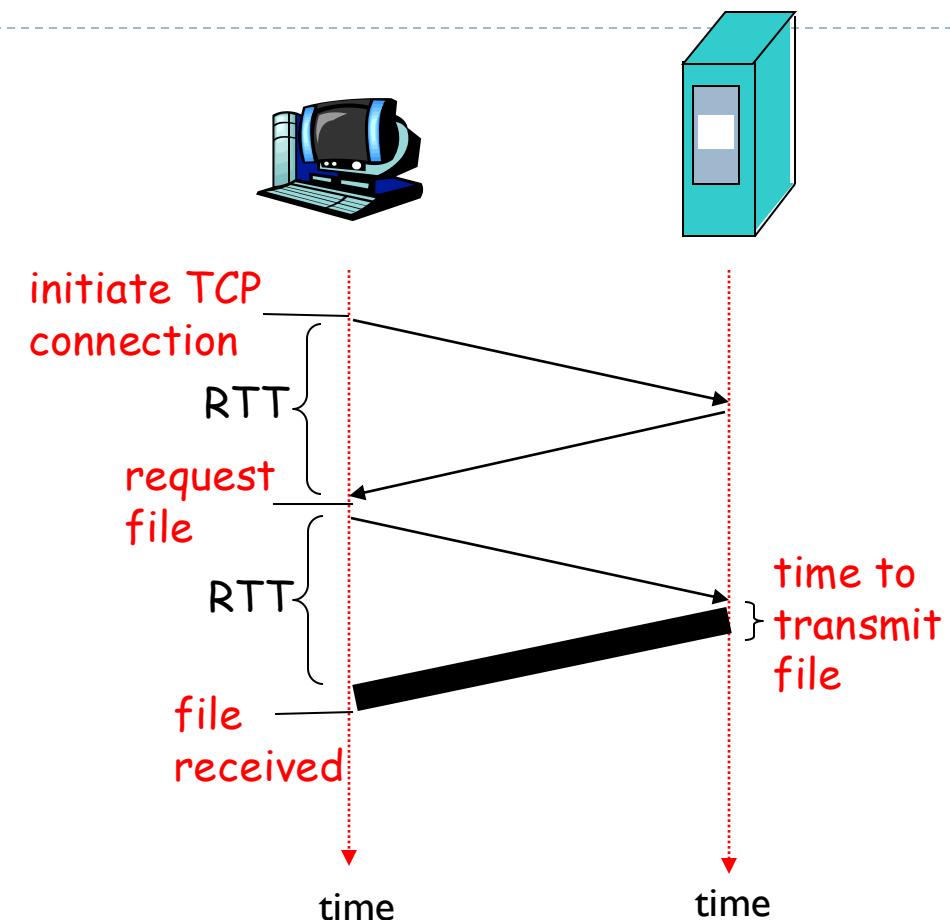
▶ Persistence vs non-persistence

▶ Non-persistence

- ▶ Server terminates connection after transferring the file/object
- ▶ HTTP 1.0 (RFC 1945)
- ▶ New connection is established each time (new set of TCP variables and buffers)
- ▶ browsers often open parallel TCP connections to fetch referenced objects

★ Response time (Non-persistence)

- one RTT (Round Trip Time) to initiate TCP connection
 - one RTT for HTTP request and first few bytes of HTTP response to return
 - file transmission time
- total = 2RTT+transmit time**



RTT- Time taken for a small packet from client to server and then back to the client.

▶ Persistence:

- ▶ Server leaves the TCP connection open after sending the response. Subsequent requests (referenced objects) between the same client and server will use the same connection.
- ▶ **Without pipelining** (one RTT for each referenced object)
- ▶ **With pipelining** (as little as one RTT for all the referenced objects)
- ▶ Default in HTTP 1.1 (RFC 2616): persistence with pipelining

Without Pipelining (one RTT per referenced object)

Client (left, blue) \leftrightarrow Server (right, green)

With Pipelining (HTML first, then 4 objects in 1 RTT)

Client (left, blue) \leftrightarrow Server (right, green)

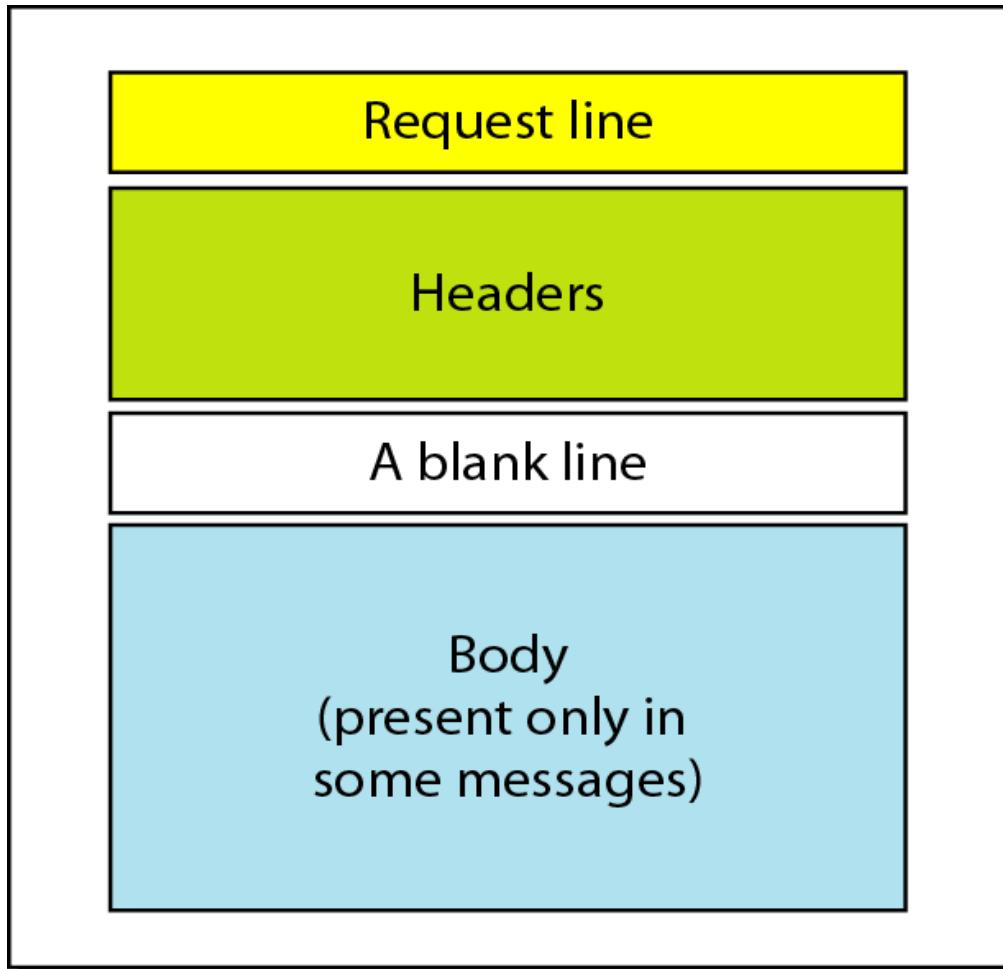
Legend: Request lines are thin; response lines are thicker. Batched response thickness equals the sum of the independent object lines (e.g., 4 objects \times 8px each \rightarrow 32px).

Sides: Client bubbles are blue (left), Server bubbles are green (right). TCP connection is kept open and reused for subsequent messages.

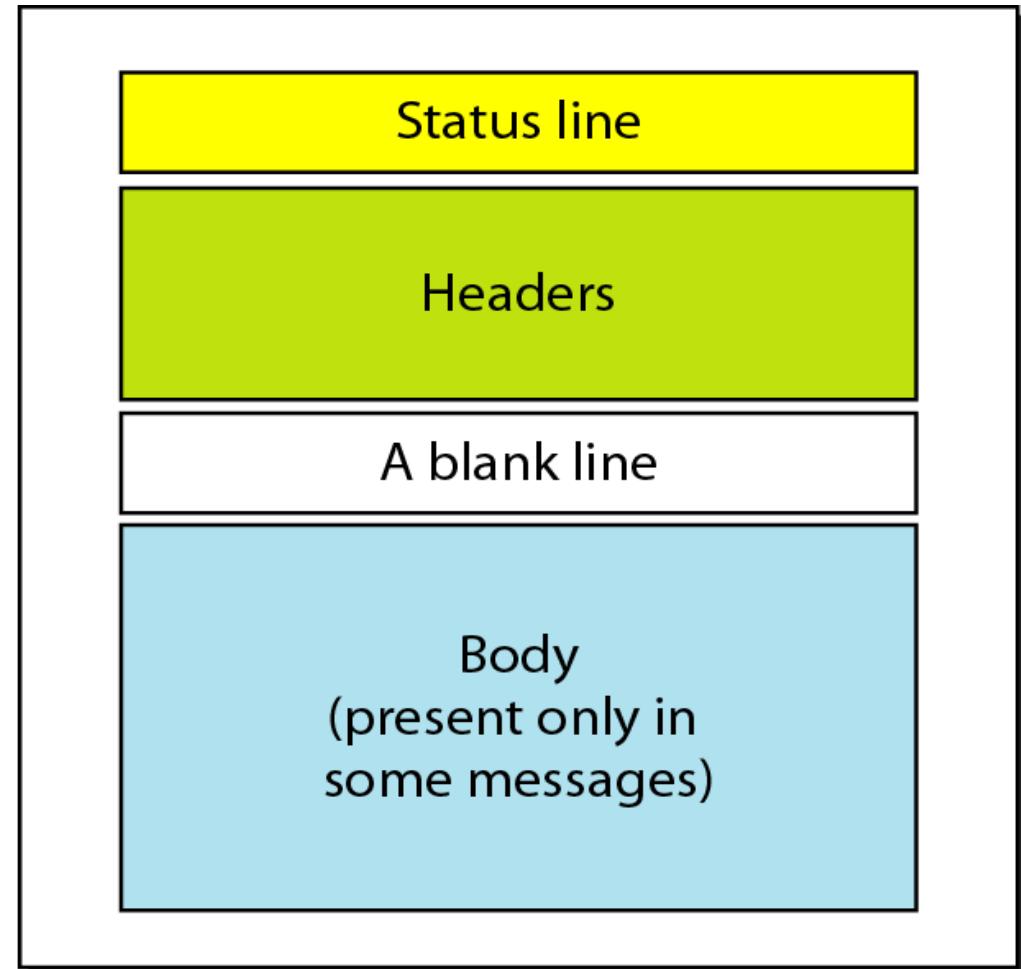
RTT: Highlighted in red inside brackets.

SENDS html file first,
then the objects, if pipelining, sends all at once,
if not then one at a time

HTTP Message Format

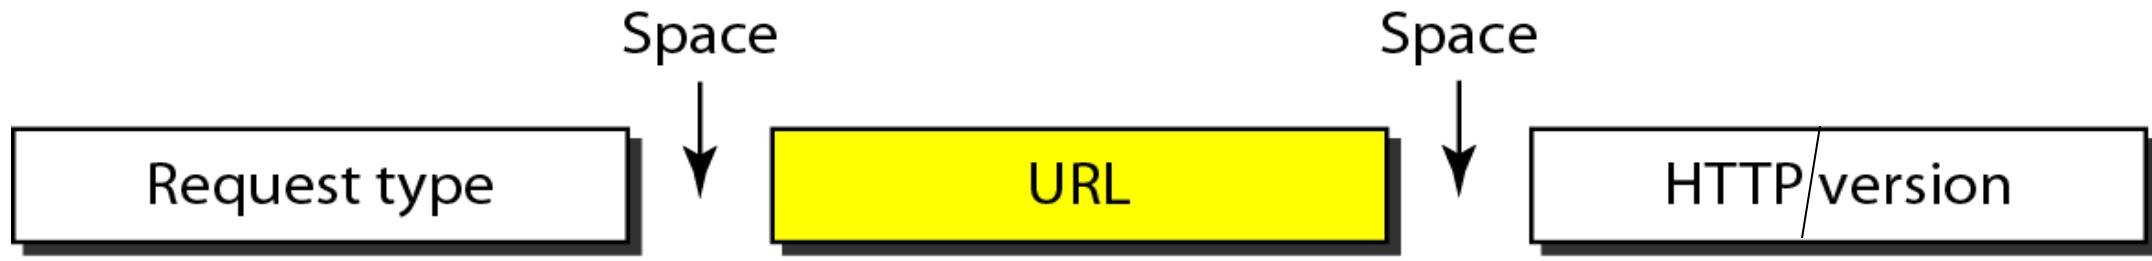


Request message

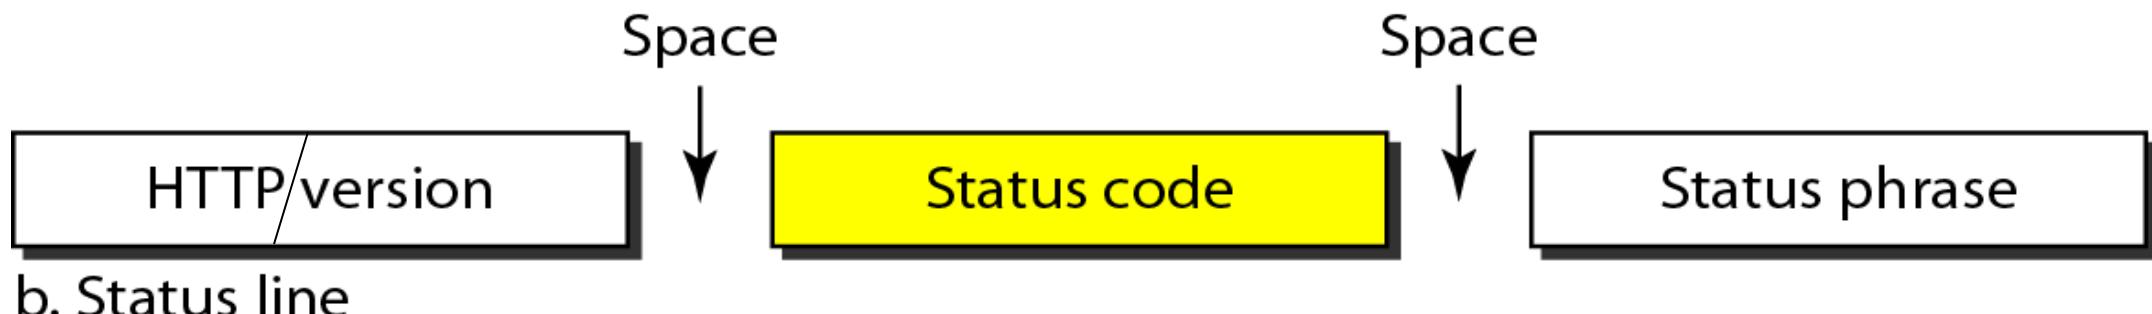


Response message

HTTP - Request line and Response status line



a. Request line



b. Status line

HTTP request message

- **HTTP request message:**

- ASCII (human-readable format)

request line (GET,
POST,
HEAD commands)

header
lines

carriage return, line feed
at start of line indicates
end of header lines

```
GET /index.html HTTP/1.1\r\n
Host: www.google.com\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
```

carriage return character
line-feed character

Other HTTP request messages

POST method:

- web page often includes form input
- user input sent from client to server in entity body of HTTP POST request message

GET method (for sending data to server):

- include user data in URL field of HTTP GET request message (following a '?'):

www.somesite.com/animalsearch?monkeys&banana

PUT method:

- uploads new file (object) to server
- completely replaces file that exists at specified URL with content in entity body of POST HTTP request message

DELETE method:

- Delete specified object (file) in the server

HEAD method:

- requests headers (only) that would be returned if specified URL were requested with an HTTP GET method.



HTTP response message

status line (protocol
status code status phrase) → HTTP/1.1 200 OK\r\nDate: Sun, 26 Sep 2010 20:09:20 GMT\r\nServer: Apache/2.0.52 (CentOS) \r\nLast-Modified: Tue, 30 Oct 2007 17:00:02
GMT\r\nETag: "17dc6-a5c-bf716880"\r\nAccept-Ranges: bytes\r\nContent-Length: 2652\r\nKeep-Alive: timeout=10, max=100\r\nConnection: Keep-Alive\r\nContent-Type: text/html; charset=ISO-8859-
1\r\n\r\ndata data data data data ...

header
lines

data, e.g., requested
HTML file →



HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

200 OK

- request succeeded, requested object later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (in Location: field)

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

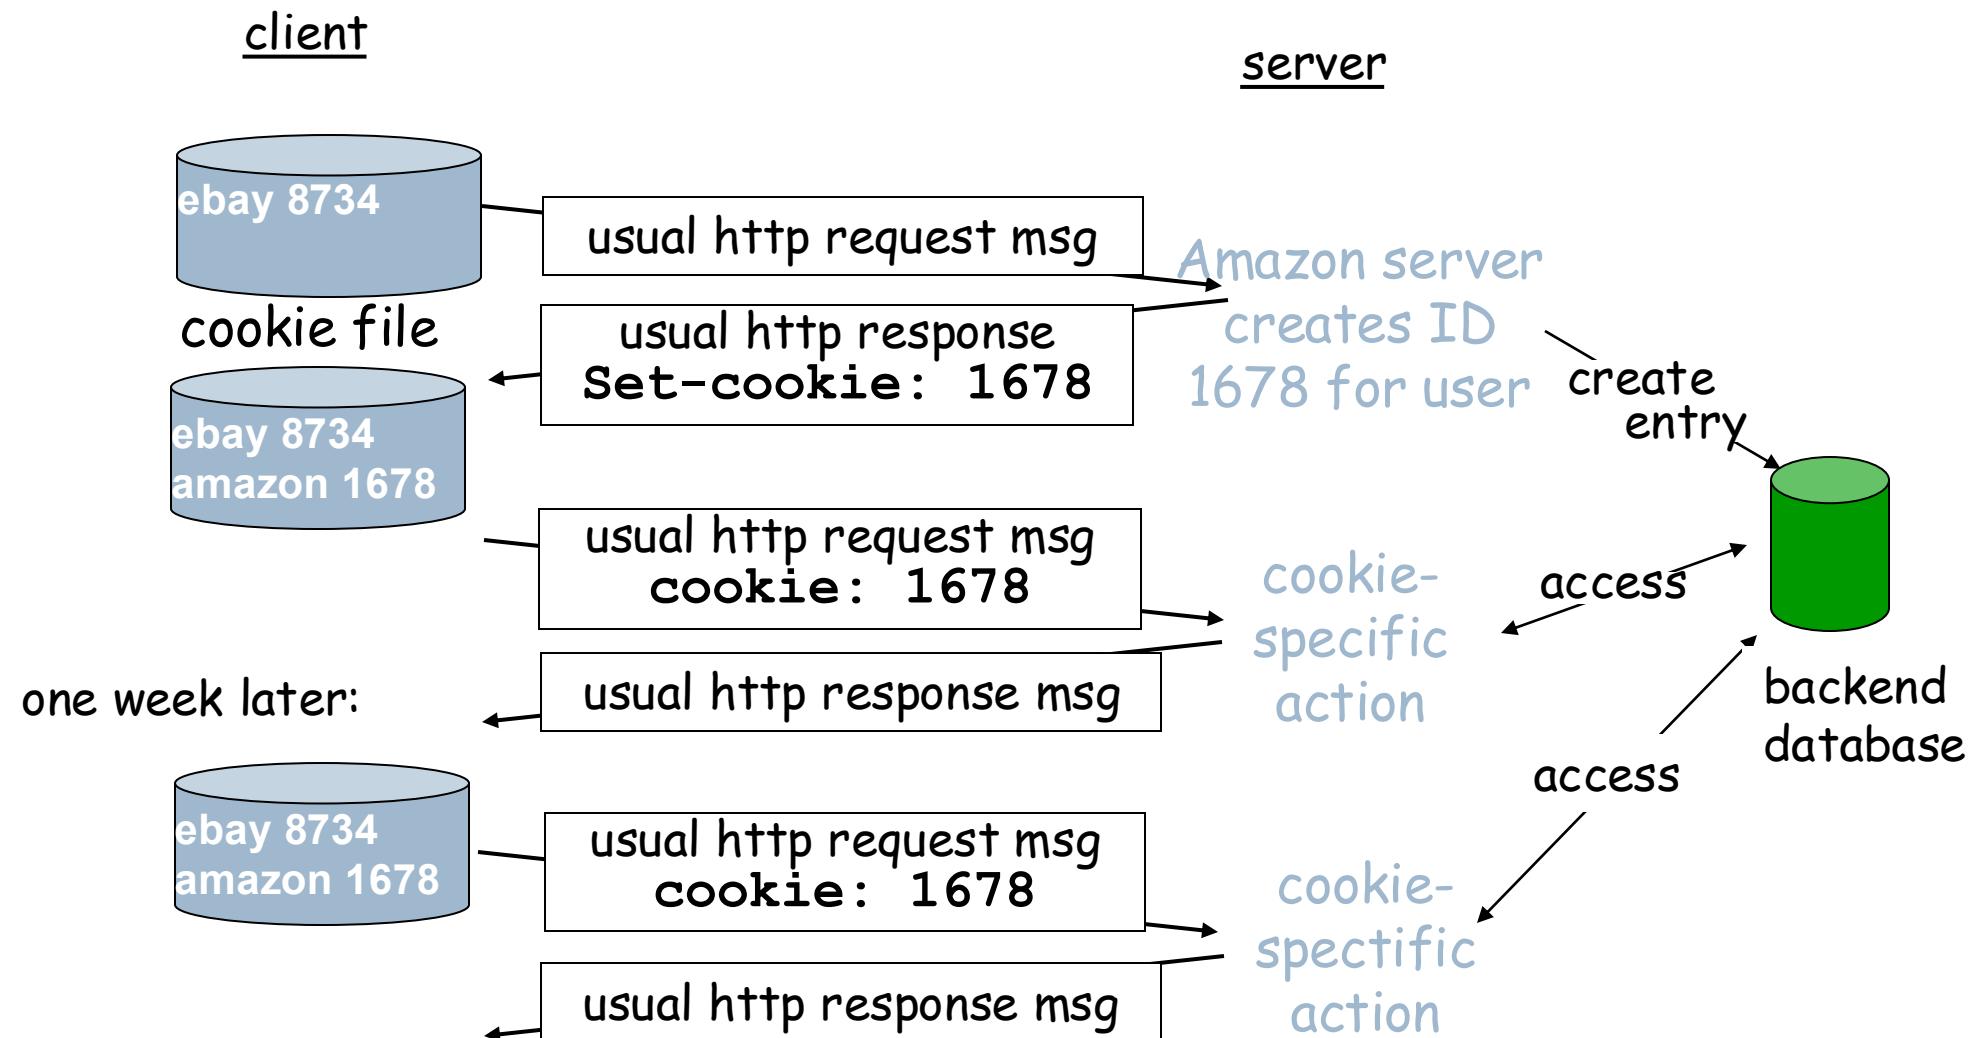
505 HTTP Version Not Supported



HTTP - User State Management : Cookies

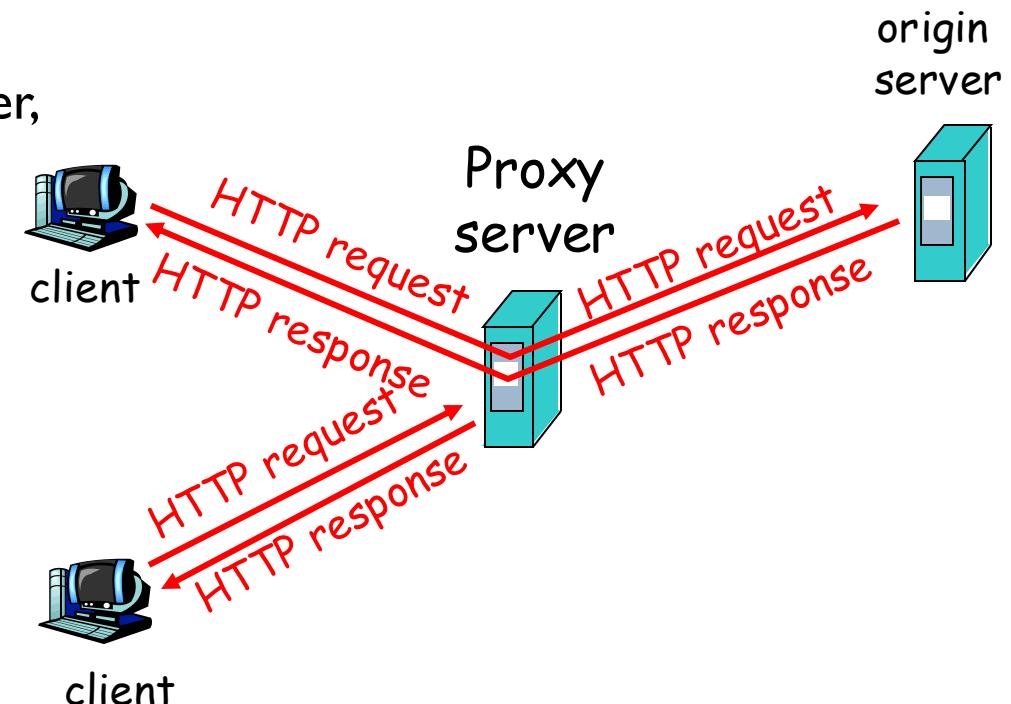
- ▶ HTTP server is **stateless**.
- ▶ HTTP uses **cookies** to manage user identity and state
- ▶ Cookies have four components.
 - ▶ HTTP **response header**. (by web server upon initial request from client)
 - ▶ Set-cookie: 1232341
 - ▶ HTTP **request header**. (every subsequent request by client)
 - ▶ Cookie: 1232341
 - ▶ **Cookie file** is maintained by the browser (client)
 - ▶ **Back-end database** at the web site (web server)

HTTP - User State Management : Cookies



HTTP - Web caches (proxy server)

- ▶ **Goal:** satisfy client request without involving origin server. Improve response time; reduce traffic.
 - ▶ user sets browser: Web accesses via cache
 - ▶ browser sends all HTTP requests to cache
 - ▶ object in cache: cache returns object
 - ▶ else cache requests object from origin server, then returns object to client



Conditional GET

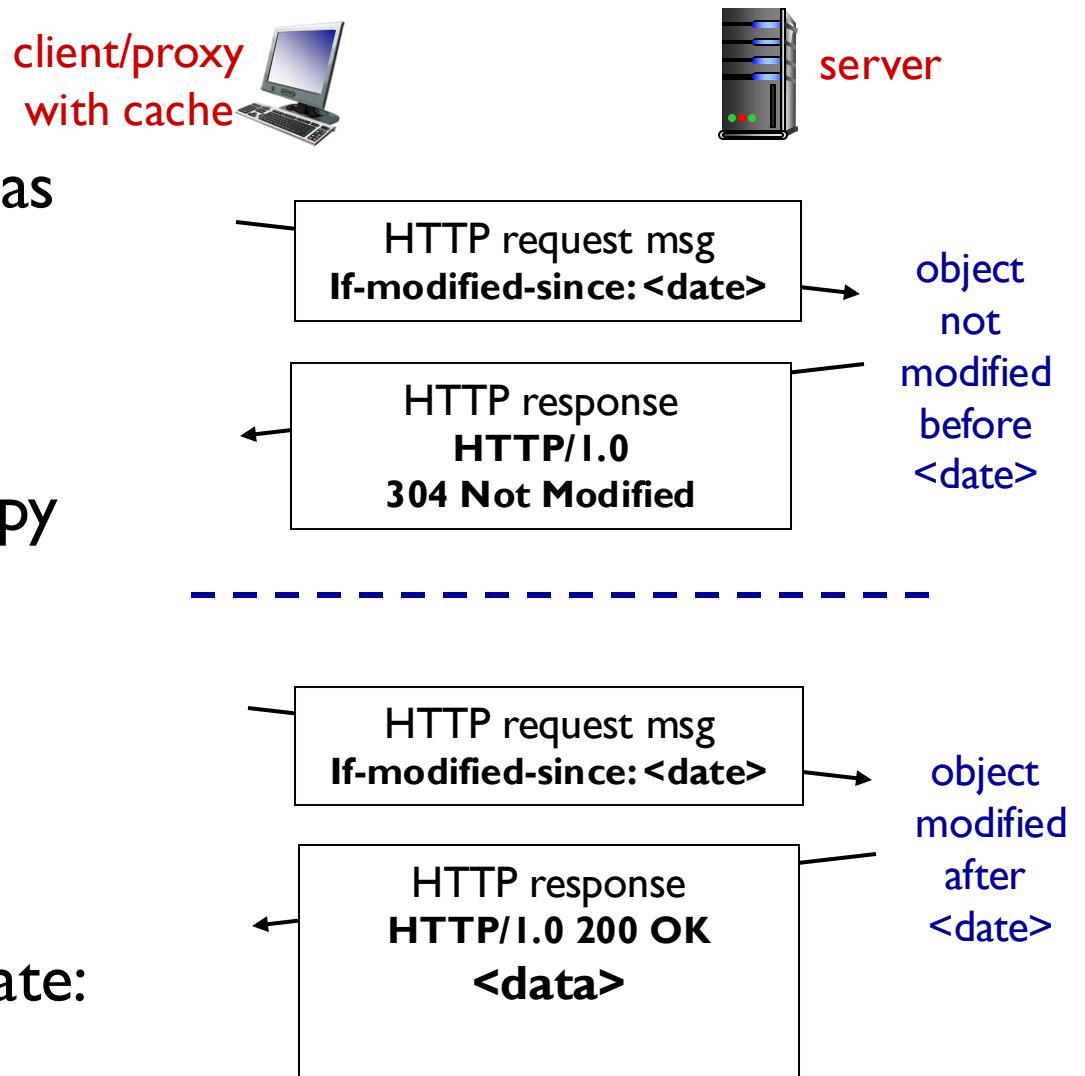
Goal: don't send object if cache has up-to-date cached version

- no object transmission delay
- lower link utilization
- **cache:** specify date of cached copy in HTTP request

If-modified-since: <date>

<date> - taken from **Last-modified:**

- **server:** response contains no object if cached copy is up-to-date:
HTTP/1.0 304 Not Modified



HTTP demo

bhojan — nc -v example.org 80 — 104x55

```
(base) bhojan@adminsoc5-MacBook-Pro-9 ~ % nc -v example.org 80
Connection to example.org port 80 [tcp/http] succeeded!
GET / HTTP/1.1
Host: example.org
-v: verbose
HTTP/1.1 200 OK
Accept-Ranges: bytes
Age: 218964
Cache-Control: max-age=604800
Content-Type: text/html; charset=UTF-8
Date: Mon, 04 Sep 2023 17:46:05 GMT
Etag: "3147526947"
Expires: Mon, 11 Sep 2023 17:46:05 GMT
Last-Modified: Thu, 17 Oct 2019 07:18:26 GMT
Server: ECS (oxr/830D)
Vary: Accept-Encoding
X-Cache: HIT
Content-Length: 1256

<!doctype html>
<html>
<head>
  <title>Example Domain</title>

  <meta charset="utf-8" />
  <meta http-equiv="Content-type" content="text/html; charset=utf-8" />
  <meta name="viewport" content="width=device-width, initial-scale=1" />
  <style type="text/css">
body {
  background-color: #f0f0f2;
  margin: 0;
```

bhojan — bash — 68x27
BadduBlueMountain:~ bhojan\$ telnet arivu.comp.nus.edu.sg 80
Trying 137.132.83.224...
Connected to arivu.d2.comp.nus.edu.sg.
Escape character is '^].
GET /index.html HTTP/1.1
Host: arivu.comp.nus.edu.sg

HTTP/1.1 200 OK
Date: Wed, 21 Jun 2017 07:03:14 GMT
Server: Apache/2.4.7 (Ubuntu)
Last-Modified: Thu, 22 Mar 2012 12:25:22 GMT
ETag: "b8-4bbd40082d39a"
Accept-Ranges: bytes
Content-Length: 184
Vary: Accept-Encoding
Content-Type: text/html

<html>
<head>
<title>ARIVU</title>
</head>
<body><h1>Hi, You are in ARIVU Server @ SoC</h1>
<p>It runs multiple services</p>
<p>Pls point to the appropriate service</p>
</body></html>
Connection closed by foreign host.
BadduBlueMountain:~ bhojan\$

HTTP demo

curl -v -I https://varlabs.comp.nus.edu.sg/tools/index.html

```
● ● ● bhojan — zsh — 104x55
(base) bhojan@adminsoc-MacBook-Pro-9 ~ % curl -v -L https://varlabs.comp.nus.edu.sg/tools/index.html
*   Trying 137.132.84.180:443...
* Connected to varlabs.comp.nus.edu.sg (137.132.84.180) port 443 (#0)
* ALPN, offering http/1.1
* successfully set certificate verify locations:
*   CAfile: /opt/anaconda3/ssl/cacert.pem
  CApth: none
* TLSv1.3 (OUT), TLS handshake, Client hello (1):
* TLSv1.3 (IN), TLS handshake, Server hello (2):
* TLSv1.3 (OUT), TLS change cipher, Change cipher spec (1):
* TLSv1.3 (OUT), TLS handshake, Client hello (1):
* TLSv1.3 (IN), TLS handshake, Server hello (2):
* TLSv1.3 (IN), TLS handshake, Encrypted Extensions (8):
* TLSv1.3 (IN), TLS handshake, Certificate (11):
* TLSv1.3 (IN), TLS handshake, CERT verify (15):
* TLSv1.3 (IN), TLS handshake, Finished (20):
* TLSv1.3 (OUT), TLS handshake, Finished (20):
* SSL connection using TLSv1.3 / TLS_AES_256_GCM_SHA384
* ALPN, server accepted to use http/1.1
* Server certificate:
*   subject: CN=*.comp.nus.edu.sg
*   start date: Nov 21 00:00:00 2022 GMT
*   expire date: Dec 22 23:59:59 2023 GMT
*   subjectAltName: host "varlabs.comp.nus.edu.sg" matched cert's "*.comp.nus.edu.sg"
*   issuer: C=GB; ST=Greater Manchester; L=Salford; O=Sectigo Limited; CN=Sectigo RSA Domain Validation Secure Server CA
*   SSL certificate verify ok.
> GET /tools/index.html HTTP/1.1
> Host: varlabs.comp.nus.edu.sg
> User-Agent: curl/7.71.1
> Accept: */*
>
* Mark bundle as not supporting multiuse
< HTTP/1.1 200 OK
< Server: nginx/1.10.2
< Date: Mon, 04 Sep 2023 17:32:17 GMT
< Content-Type: text/html
< Content-Length: 185
< Last-Modified: Sun, 17 Sep 2017 14:12:55 GMT
```

Alternatively, you can
use 'wget'

-v : verbose
-I : follow
redirects

HTTP/2

Key goal: decreased delay in multi-object HTTP requests

HTTP 1.1: introduced **multiple, pipelined GETs** over single TCP connection

- server responds *in-order* (FCFS: first-come-first-served scheduling) to GET requests
- with FCFS, small object may have to wait for transmission (**head-of-line (HOL) blocking**) behind large object(s)
- loss recovery (retransmitting lost TCP segments) stalls object transmission



HTTP/2

Key goal: decreased delay in multi-object HTTP requests

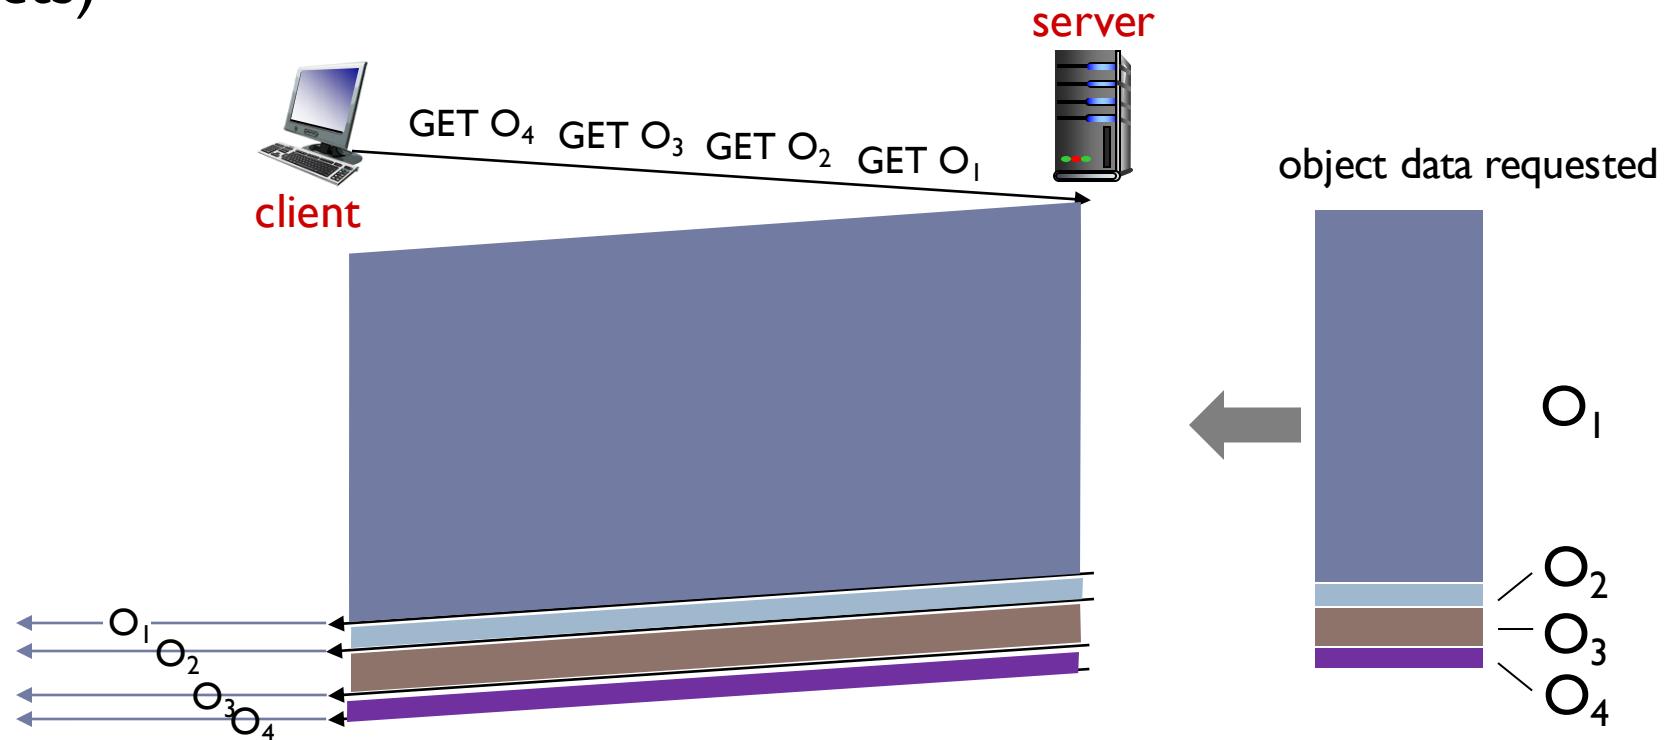
HTTP/2: [RFC 7540, 2015] increased flexibility at server in sending objects to client:

- methods, status codes, most header fields unchanged from HTTP 1.1
- transmission order of requested objects based on **client-specified object priority** (not necessarily FCFS)
- **push** unrequested objects to client
- divide objects into frames, schedule frames to mitigate HOL blocking



HTTP/2: mitigating HOL blocking

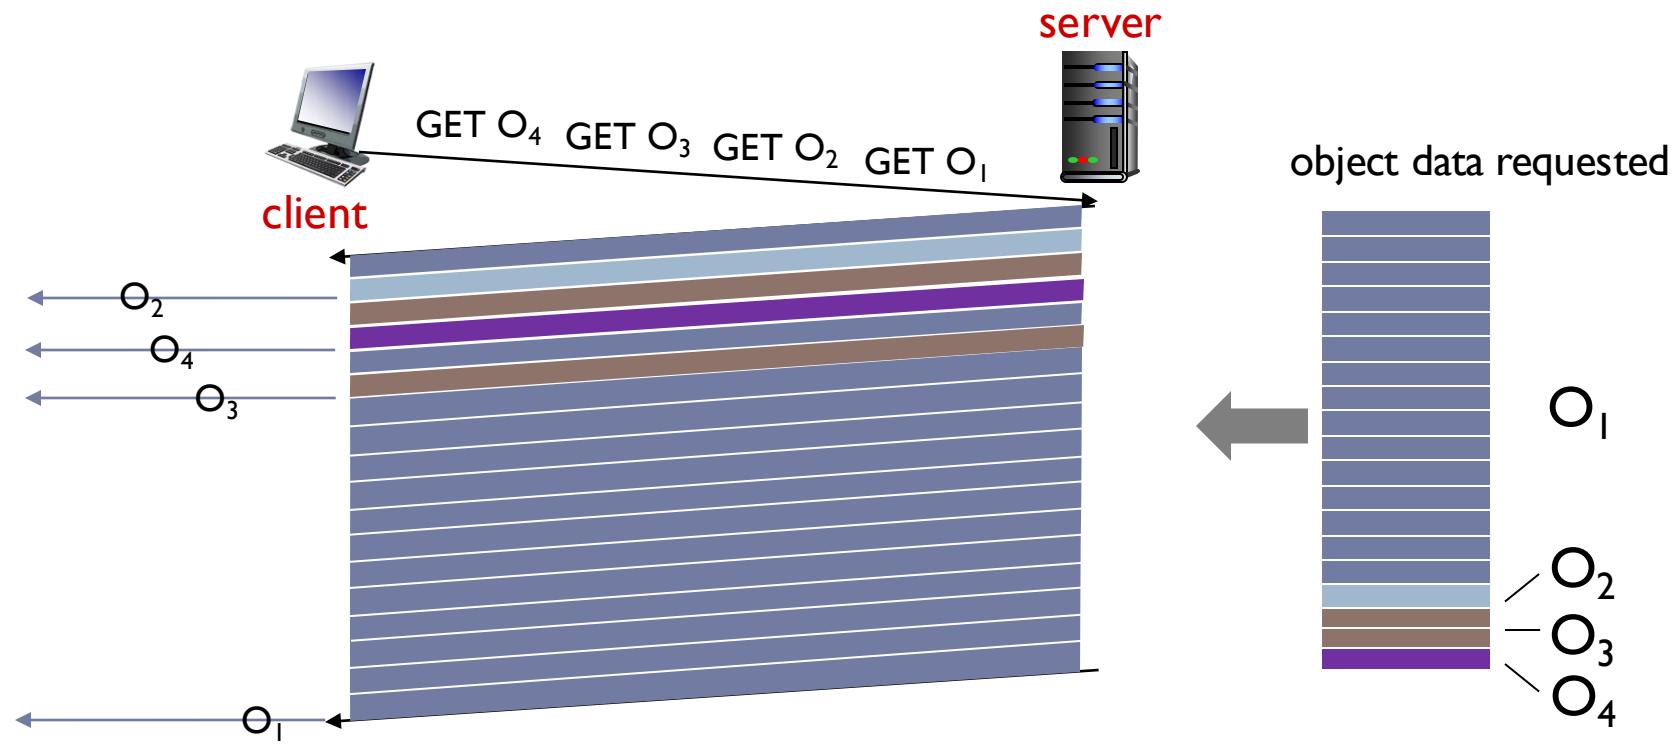
HTTP 1.1: client requests 1 large object (e.g., video file, and 3 smaller objects)



objects delivered in order requested: O₂, O₃, O₄ wait behind O₁

HTTP/2: mitigating HOL blocking

HTTP/2: objects divided into frames, frame transmission interleaved



O_2, O_3, O_4 delivered quickly, O_1 slightly delayed

HTTP/2 to HTTP/3

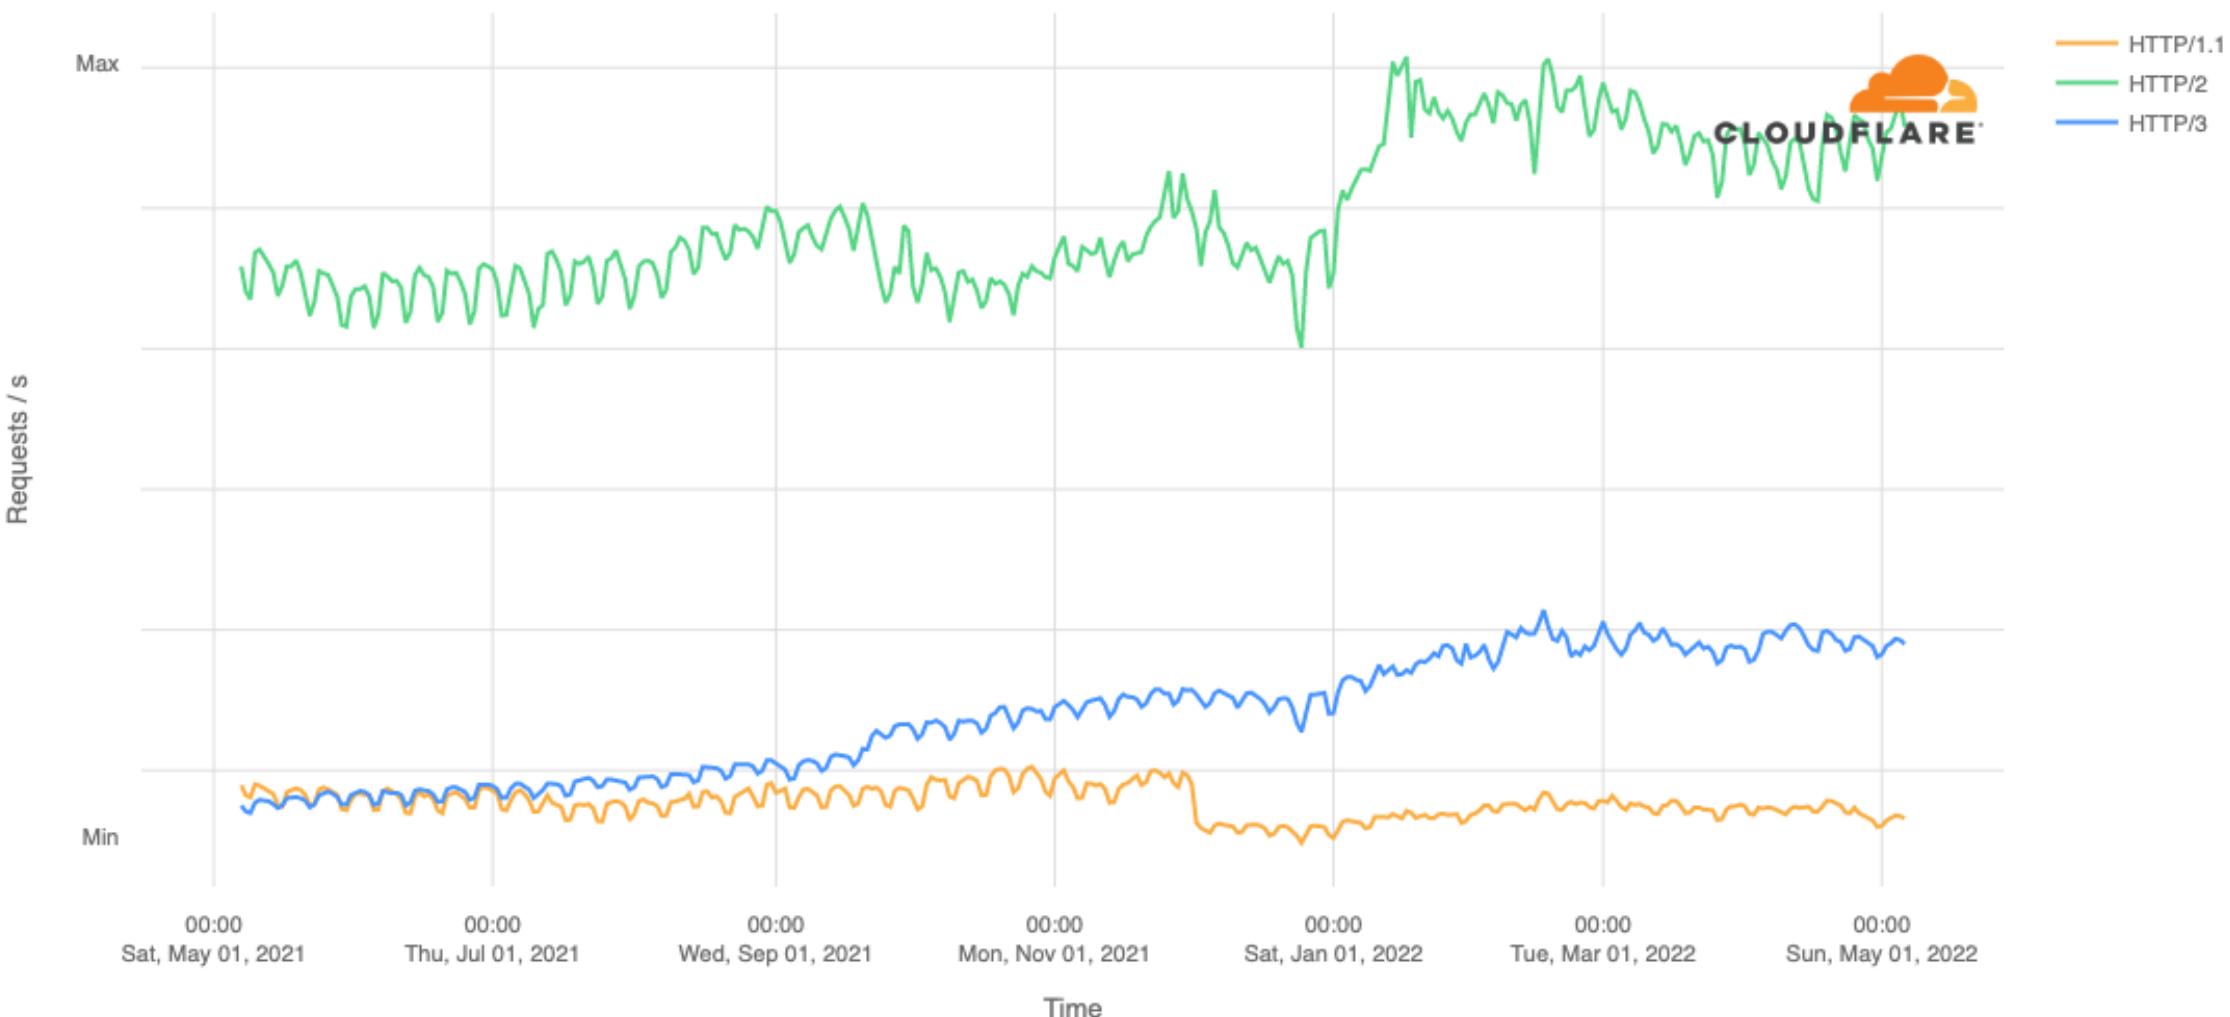
Key goal: decreased delay in multi-object HTTP requests

HTTP/2 over **single TCP connection** means:

- recovery from packet loss still stalls all object transmissions
 - as in HTTP 1.1, browsers have incentive to open multiple parallel TCP connections to reduce stalling, increase overall throughput
- no security over vanilla TCP connection
- **HTTP/3:** adds security , per object error- and congestion-control (more pipelining) **over UDP.** [uses, **QUIC: Quick UDP Internet Connections, A UDP-Based Multiplexed and Secure Transport**]



Number of requests by HTTP version (Worldwide)



Source: <https://blog.cloudflare.com/cloudflare-view-http3-usage/>



CS3103: Computer Networks Practice

Network Applications & Demo

Principles of Network Applications

HTTP & WWW

E-mail Protocols (SMTP, POP, IMAP)

P2P Applications

Dr. Anand Bhojan

COM3-02-49, School of Computing

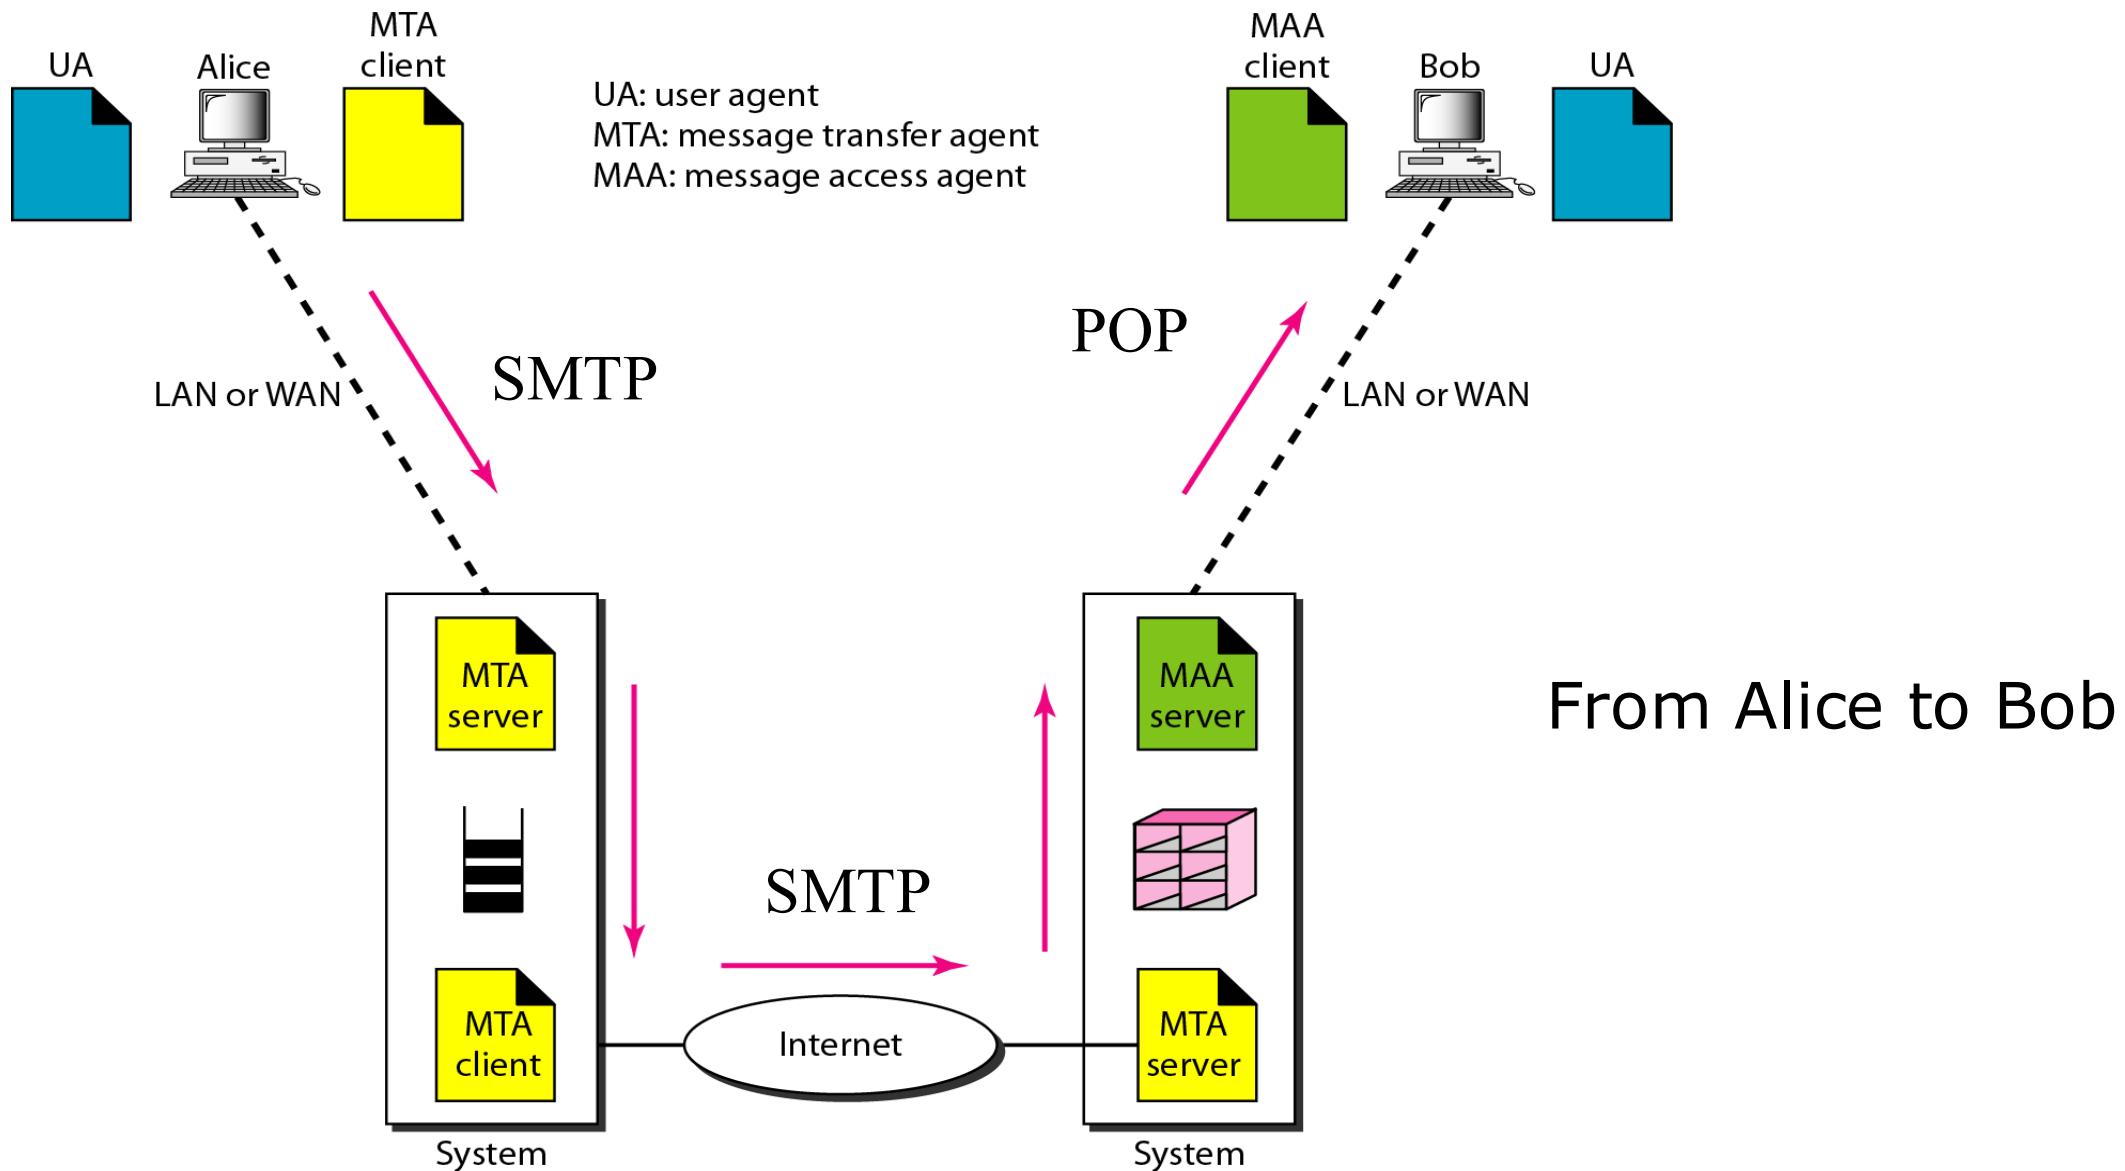
banand@comp.nus.edu.sg ph: 651-67351

E-Mail : Three major components

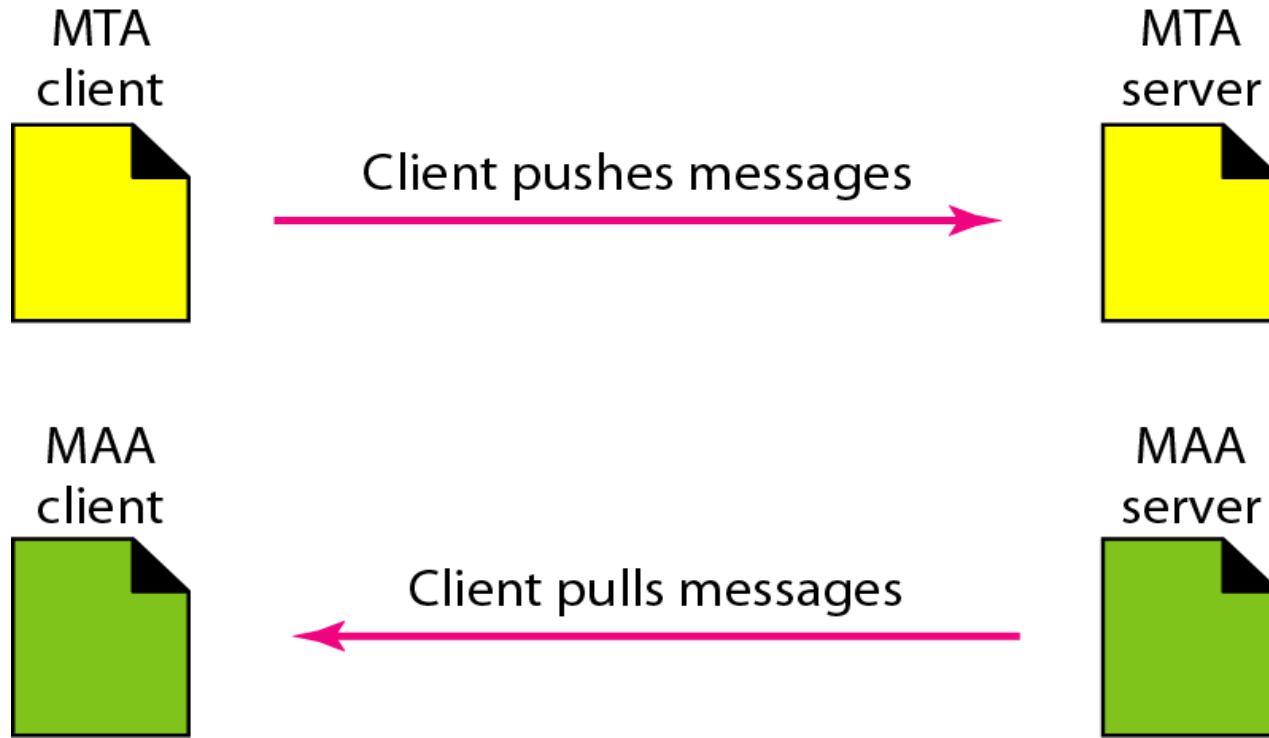
- ▶ **User Agents/Mail Readers** (eg. Eudora, Outlook, elm, Mozilla Thunderbird, pine)
- ▶ **Message Transfer Agents** (SMTP - port 25)
 - SMTP : Simple Mail Transfer Protocol
- ▶ **Message Access Agents** (POP3 – port 110 and IMAP4 - port 143)
 - POP3 – Post Office Protocol v3
 - IMAP4 – Internet Mail Access Protocol v4

E-mail

wats gg on here

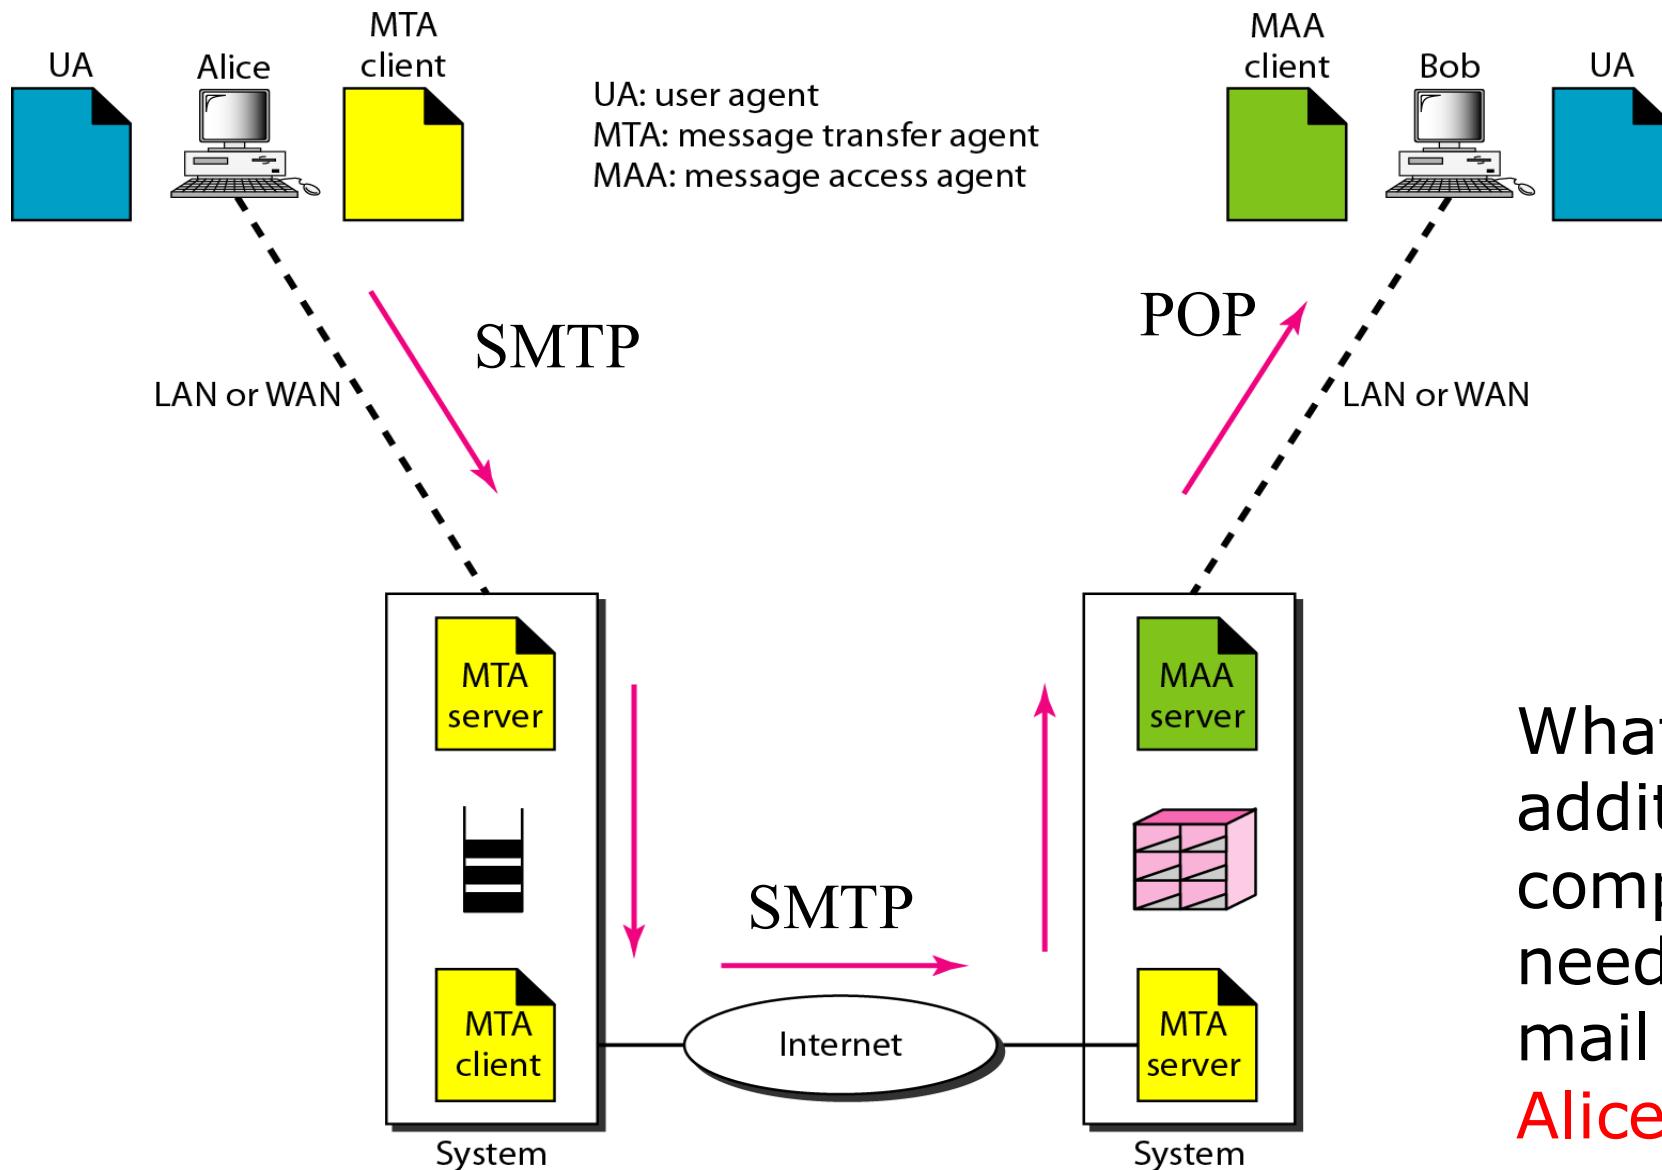


E-mail : PUSH vs PULL



- ▶ POP- PULL protocol
- ▶ SMTP – PUSH protocol
- ▶ Then, what is HTTP/1.1? (PUSH or PULL)

E-mail



What are the additional components needed to send mail from **Bob** to **Alice**?

Email - Format



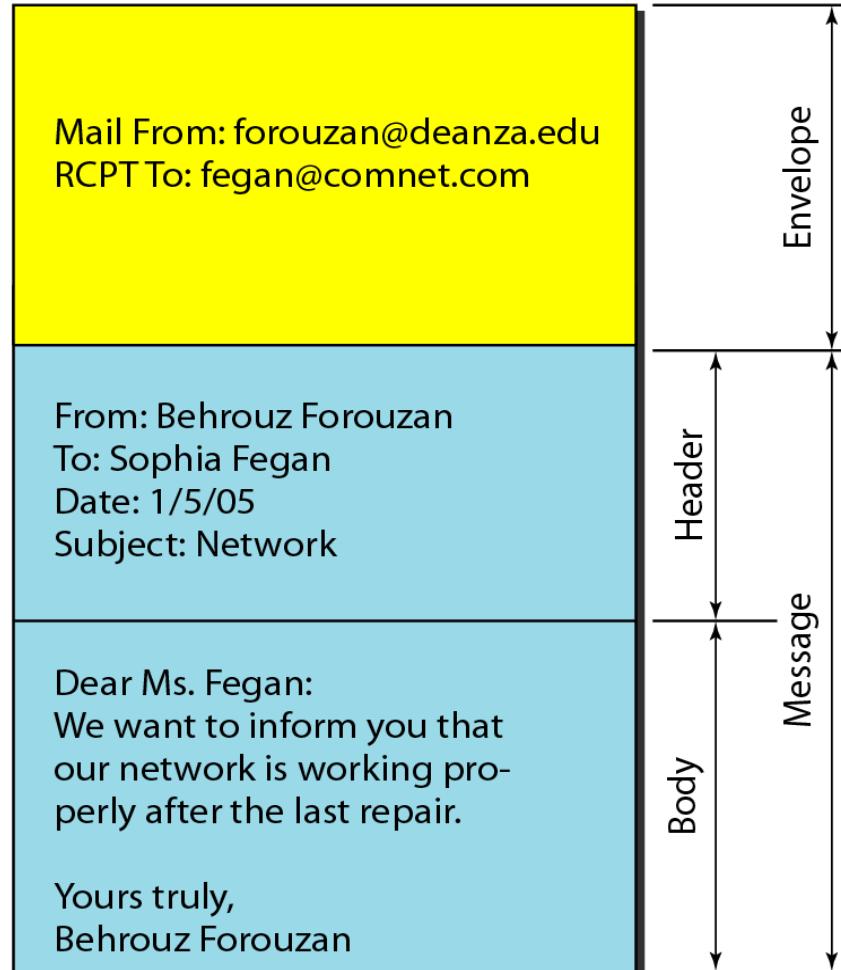
Sophia Fegan
Com-Net
Cupertino, CA 95014
Jan. 5, 2005

Subject: Network

Dear Ms. Fegan:
We want to inform you that
our network is working pro-
perly after the last repair.

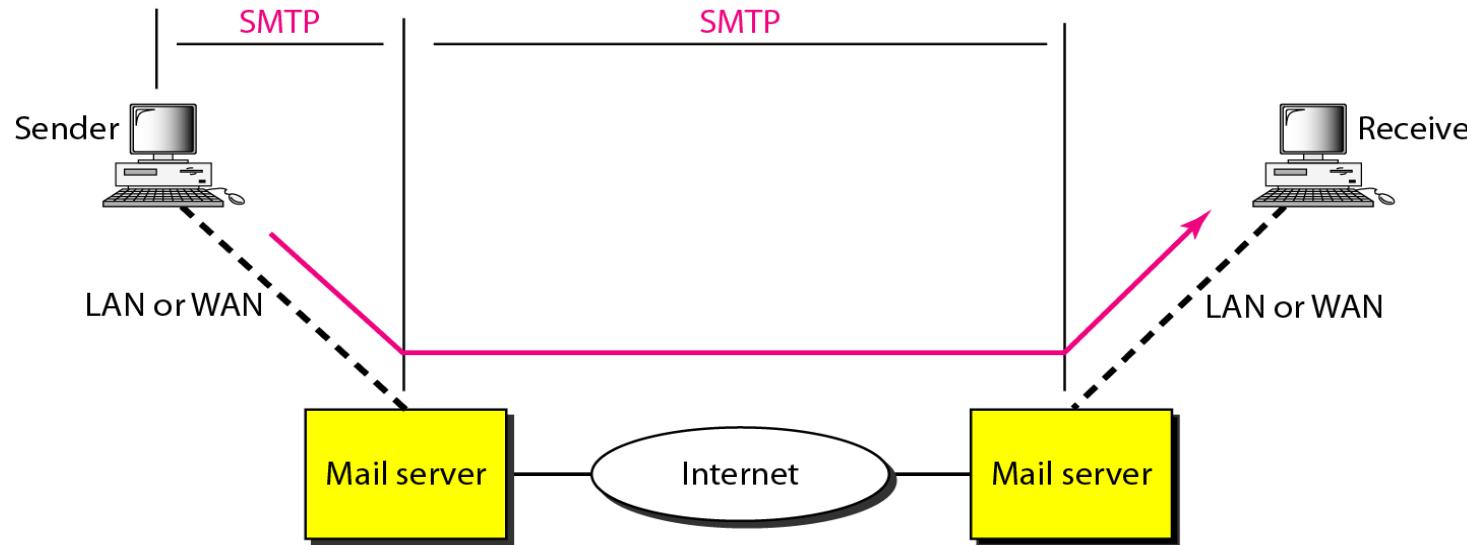
Yours truly,
Behrouz Forouzan

a. Postal mail



b. Electronic mail

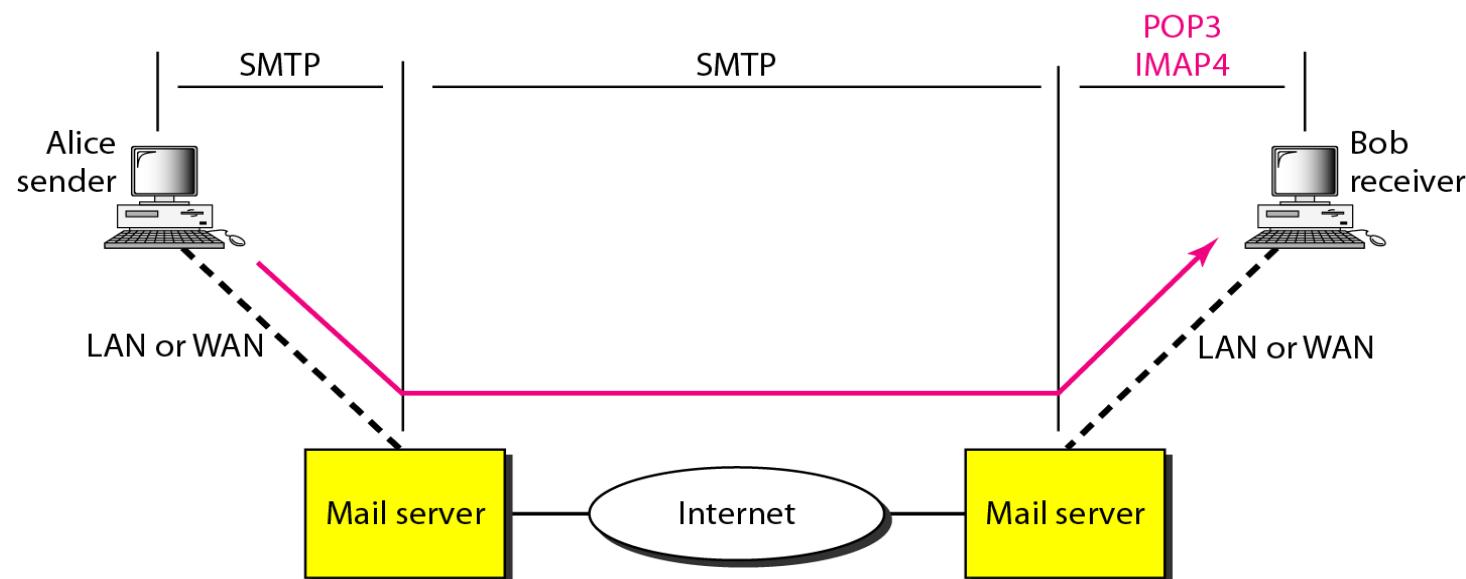
E-Mail:- MTA (eg. SMTP), MAA (eg. POP/IMAP)



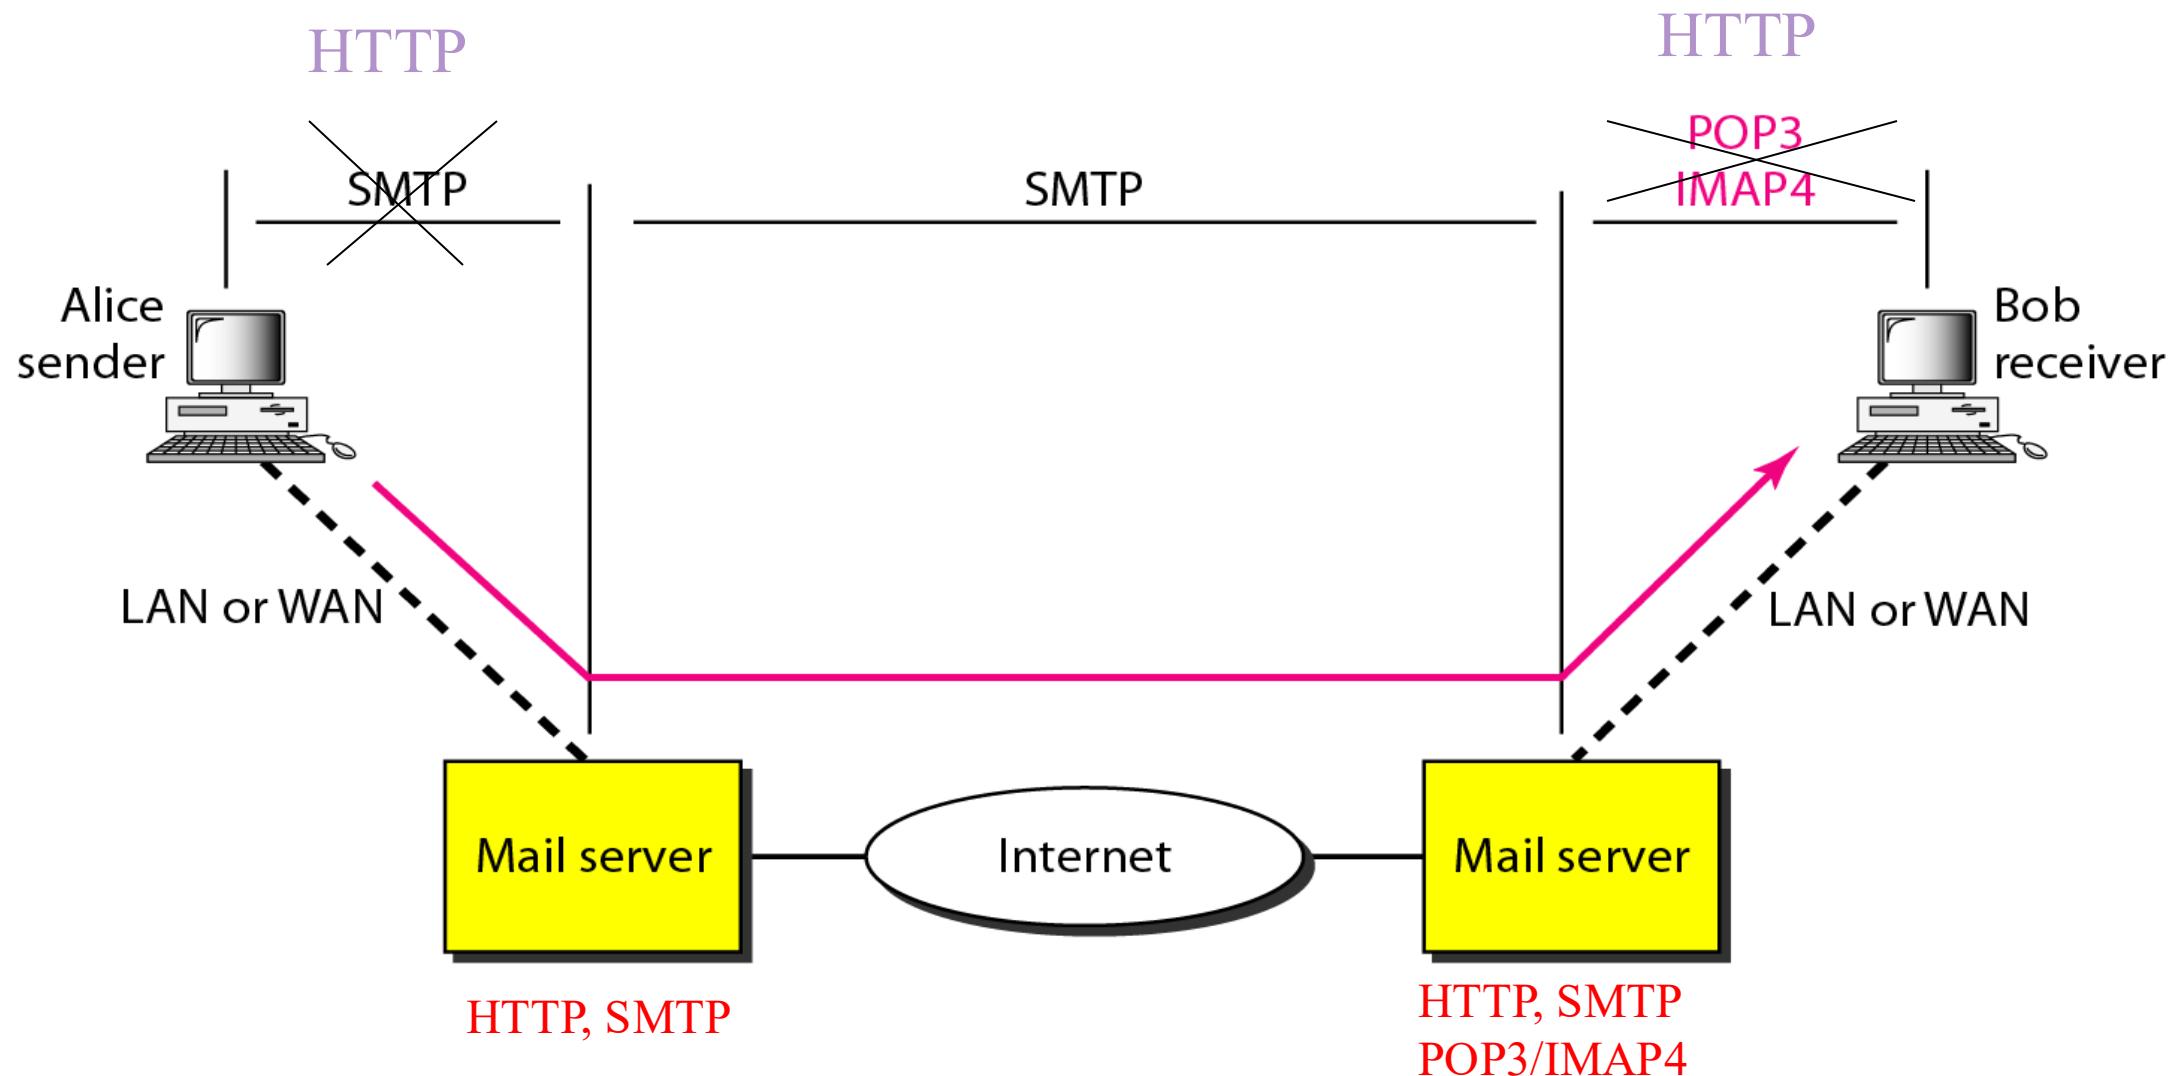
SMTP - Simple Mail Transfer Protocol

POP3-Post Office Protocol (v3) [RFC 1939]

IMAP4-Internet Mail Access Protocol (v4) [RFC 1730]



E-mail- Web based



SMTP Example

- ▶ \$openssl s_client -connect smtp.gmail.com:465 -crlf -ign_eof
- ▶ CONNECTED(00000005)
- ▶ EHLO gmail.com
- ▶ 250-mx.google.com at your service,
[137.132.84.20]....
- ▶ **AUTH LOGIN yourBASE64encodedUserID**
- ▶ 334 UAPLGGFdskzcmQ6
- ▶ **yourBase64encodedPassword**
- ▶ 235 2.7.0 Accepted
- ▶ **MAIL FROM: <yourmailID@gmail.com>**
- ▶ 250 2.1.0 OK ys4sm22528541pb.9 - gsmt
- ▶ **RCPT TO: <yourMailID@nus.edu.sg>**
- ▶ 250 2.1.5 OK ys4sm22528541pb.9 - gsmt
- ▶ **DATA**
- ▶ 354 Go ahead ys4sm22528541pb.9 - gsmt
- ▶ **FROM:me@gmail.com**
- ▶ **TO:you@nus.edu.sg**
- ▶ **Subject: It works**
- ▶ *The last line should have a*
- ▶ **DOT.**
- ▶ **My first email with openssl!**
- ▶ .
- ▶ 250 2.0.0 OK 1378814309 ys4sm22528541pb.9 - gsmt
- ▶ **QUIT**
- ▶ 221 2.0.0 closing connection
ys4sm22528541pb.9 – gsmt
- ▶ \$

- SMTP uses persistent connections
- SMTP requires message (header & body) to be in 7-bit ASCII
- SMTP: multiple objects sent in multipart msg

POP Example

- ▶ **telnet pop.nus.edu.sg 110**
- ▶ *+OK The Microsoft Exchange POP3 service is ready.*
- ▶ **USER myUsername**
- ▶ **+OK**
- ▶ **PASS myPassword**
- ▶ *+OK User successfully logged on.*
- ▶ **LIST**
- ▶ *1 1051*
- ▶ *2 908*
- ▶ *3 10884*
- ▶ *.*
- ▶ **RETR 1**
- ▶ **+OK**
- ▶ *From: "somesender" <somesender@nus.edu.sg>*
- ▶ *Content-Class: urn:kvsplc-com:cc:vault:shortcut*
- ▶ *Date: Fri, 17 Sep 2010 17:21:03 +0800*
- ▶ *Subject:Title of the mail*
- ▶ *...*
- ▶ **QUIT**
- ▶ *+OK Microsoft Exchange Server 2007 POP3 server signing off.*

Use OpenSSL for Secured POP access:

```
$openssl s_client -connect pop.nus.edu.sg:995 -crlf -ign_eof  
                                <or>  
$openssl s_client -crlf -connect pop.nus.edu.sg:110 -starttls pop3
```

POP3 (more) and IMAP

more about POP3

- ▶ previous example uses POP3 “download and delete” mode
 - ▶ Bob cannot re-read e-mail if he changes client
- ▶ POP3 “download-and-keep”: copies of messages on different clients
- ▶ POP3 is stateless across sessions

IMAP

- ▶ keeps all messages in one place: at server
- ▶ allows user to organize messages in folders
- ▶ keeps user state across sessions:
 - ▶ names of folders and mappings between message IDs and folder name

CS3103: Computer Networks Practice

P2P Applications

Principles of Network Applications

HTTP & WWW

E-mail Protocols (SMTP, POP, IMAP)

P2P Applications

Dr. Anand Bhojan

COM3-02-49, School of Computing

banand@comp.nus.edu.sg ph: 651-67351

P2P architecture

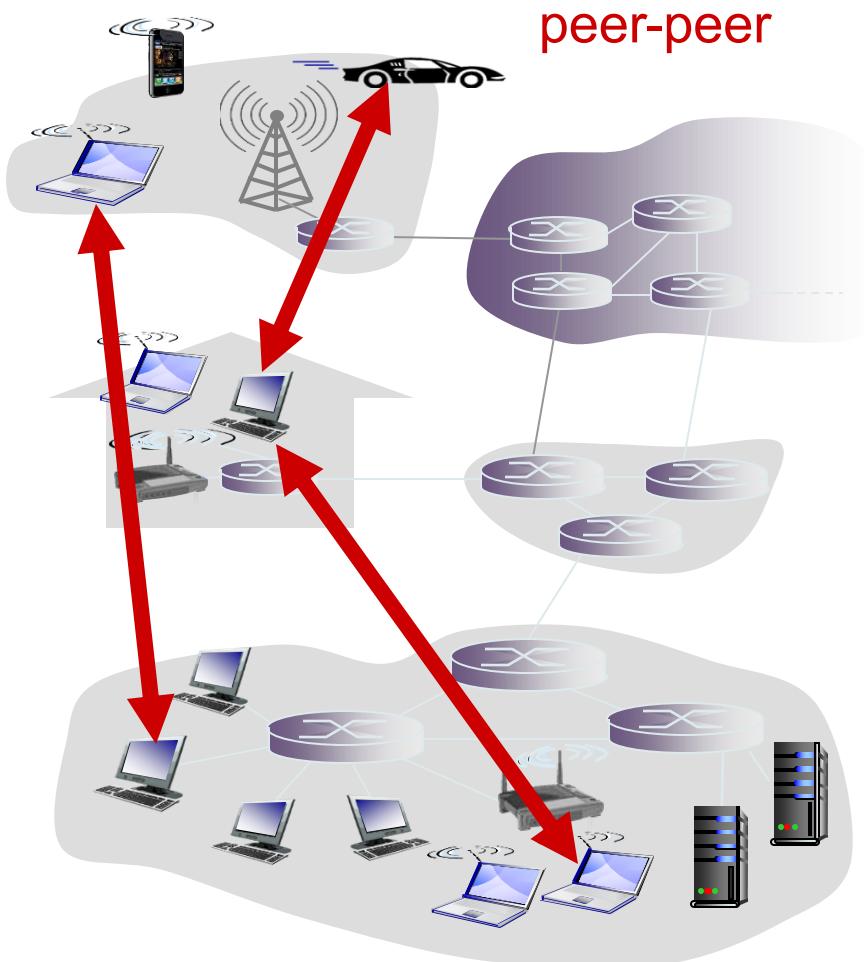
- ▶ no always-on server
- ▶ arbitrary end systems directly communicate and serve each other.
 - ▶ Each host runs both server and client processes

examples:

- ▶ file distribution (BitTorrent)
- ▶ Streaming (KanKan)
- ▶ VoIP (Skype)
- ▶ Multiplayer Games
- ▶ Blockchain (bitcoin)

Q1: What is the advantage of P2P?
- Scalability (explained next)

Q2: How the hosts find the IP addresses of each other?

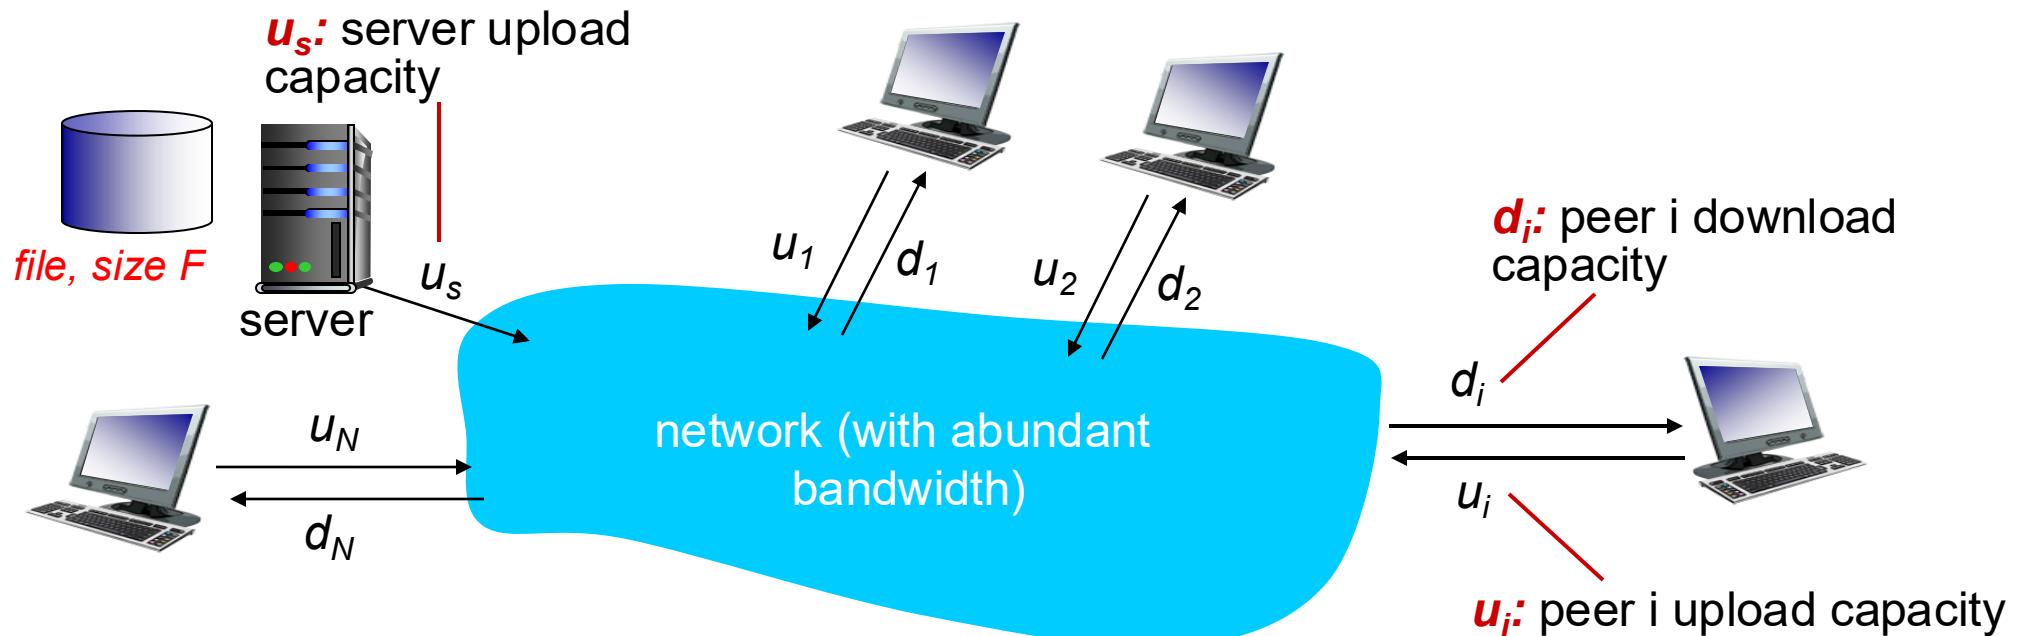


An interesting article to read: <https://www.cse.wustl.edu/~jain/cse570-21/ftp/p2p/index.html>

File distribution: client-server vs P2P

Question: how much time to distribute file (size F) from one server to N peers?

- peer upload/download capacity is limited resource



File distribution time: client-server

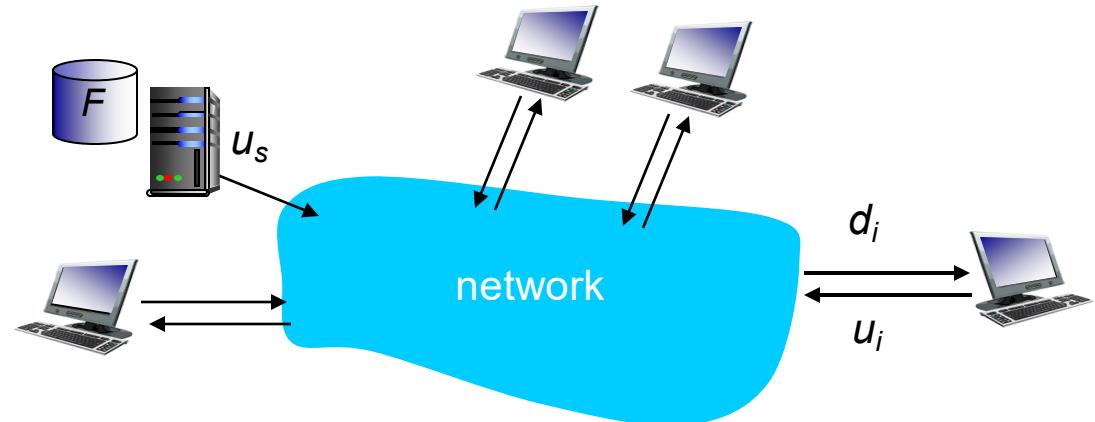
❖ **server transmission:** must sequentially send (upload) N file copies:

- ❖ time to send one copy: F/u_s
- ❖ time to send N copies: NF/u_s

❖ **client:** each client must download file

copy

- d_{\min} = min client download rate
- min client download time: F/d_{\min}



time to distribute F
to N clients using
client-server approach

$$D_{c-s} \geq \max\{NF/u_s, F/d_{\min}\}$$

increases linearly in N

File distribution time: P2P

Question: how much time to distribute file (size F) from one server to N peers?

- ❖ **server transmission:** must upload at least one copy

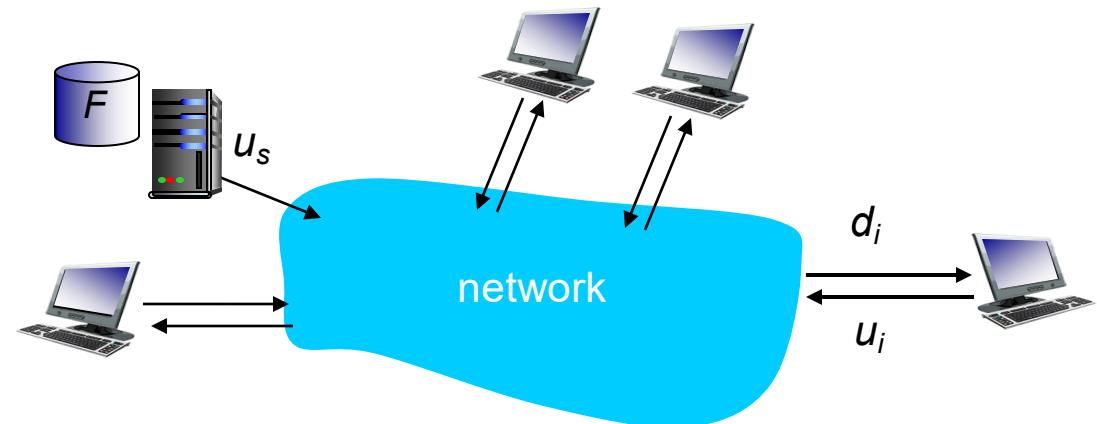
- ❖ time to send one copy: F/u_s

- ❖ **client:** each client must download file copy

- min client download time: F/d_{min}

- ❖ **clients:** as aggregate must download NF bits

- max upload rate (limiting max download rate) is $u_s + \sum u_i$



time to distribute F to N clients using P2P approach

$$D_{P2P} \geq \max\{F/u_s, F/d_{min}, NF/(u_s + \sum u_i)\}$$

increases linearly in N ...

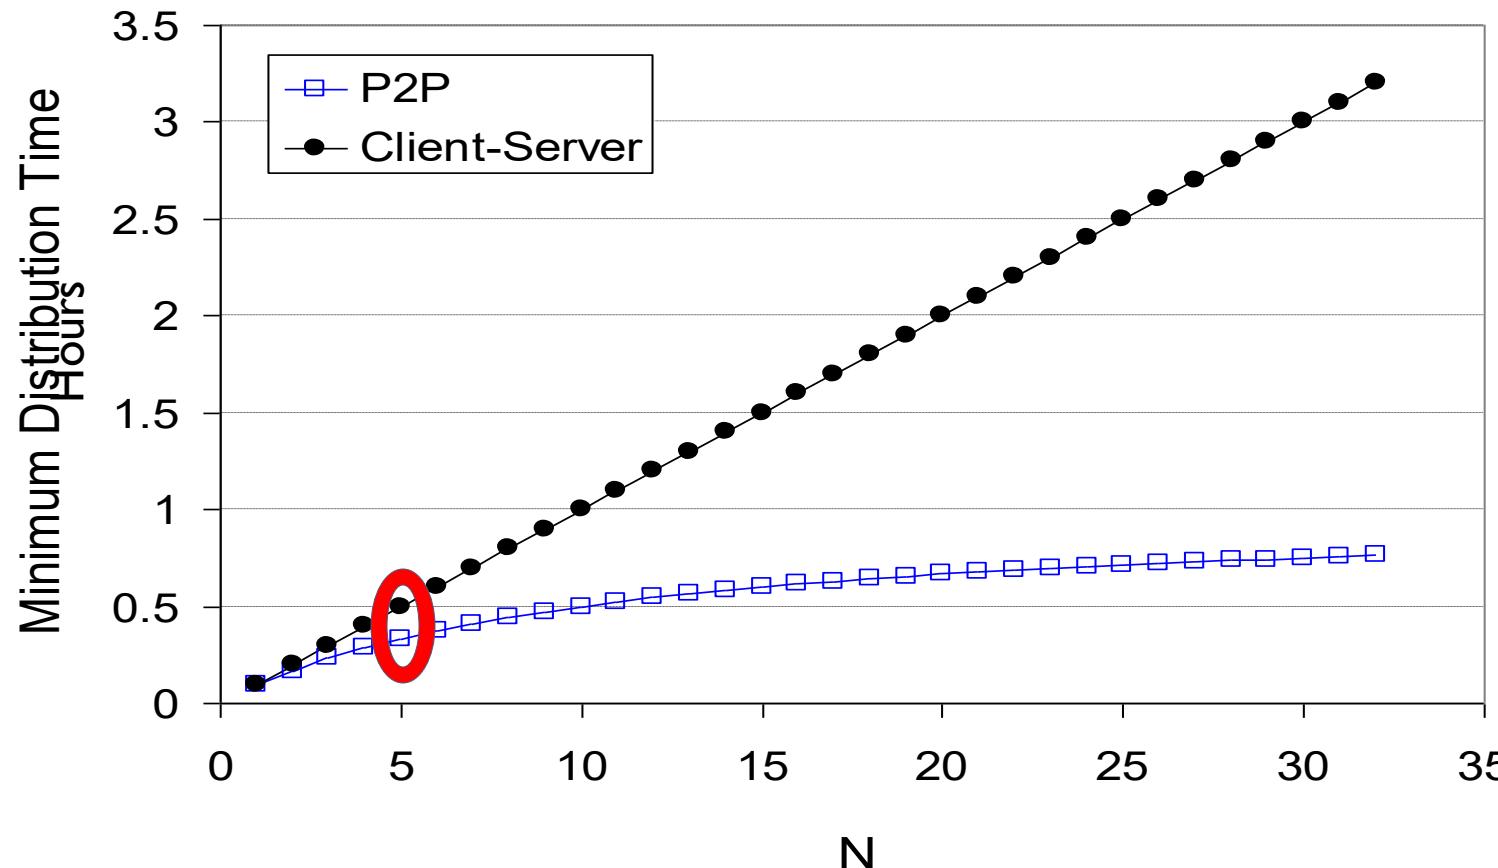
... but so does this, as each peer brings service capacity

Client-server vs. P2P: example (Just based only on upload rate)

- Client upload rate = u
- Server upload rate $u_s = 10u$

To keep it simple, Assume client download rate $d_{min} \geq u_s$
[So, **the bottleneck is upload rate in C-S mode!**]

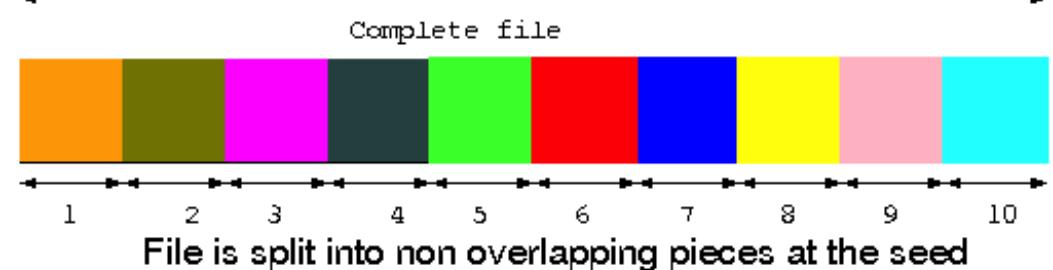
- File size $F = 60$ GBits
- Assume, client upload time $F/u = 1$ hour
 - => Server upload time $F/10u = 60/10 = 6$ mins
- So, to transmit (upload) to 5 clients in seq, it takes **5 x 6 mins**



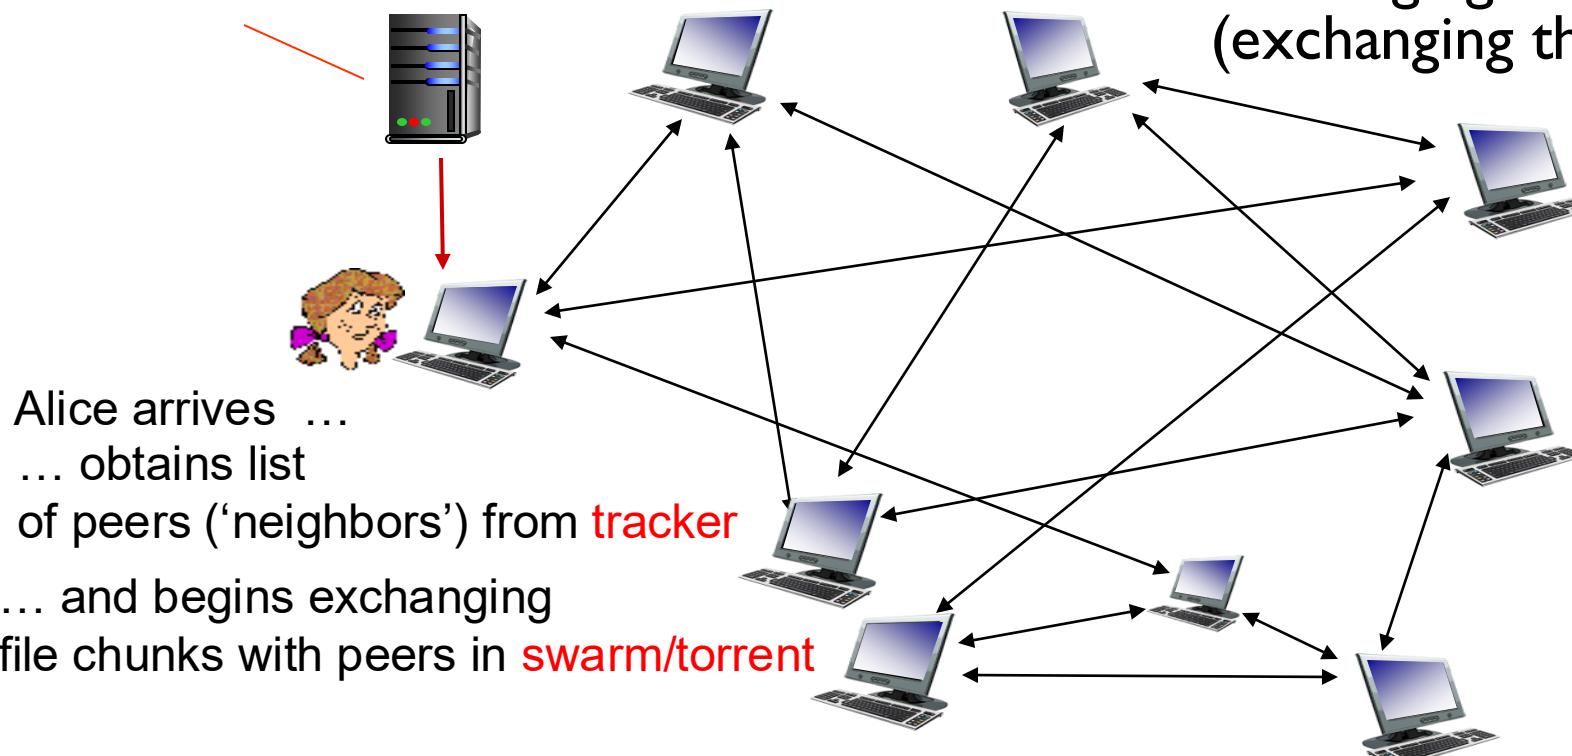
- In p2p transmission, it takes $N*60/(10u+5u)$
- Which is, $5 * 60/15u = 5 * 4$ mins

P2P file distribution: BitTorrent (BT)

- ❖ file divided into 256Kb chunks
- ❖ peers in torrent (aka, swarm) send/receive file chunks



tracker: tracks peers participating in a torrent



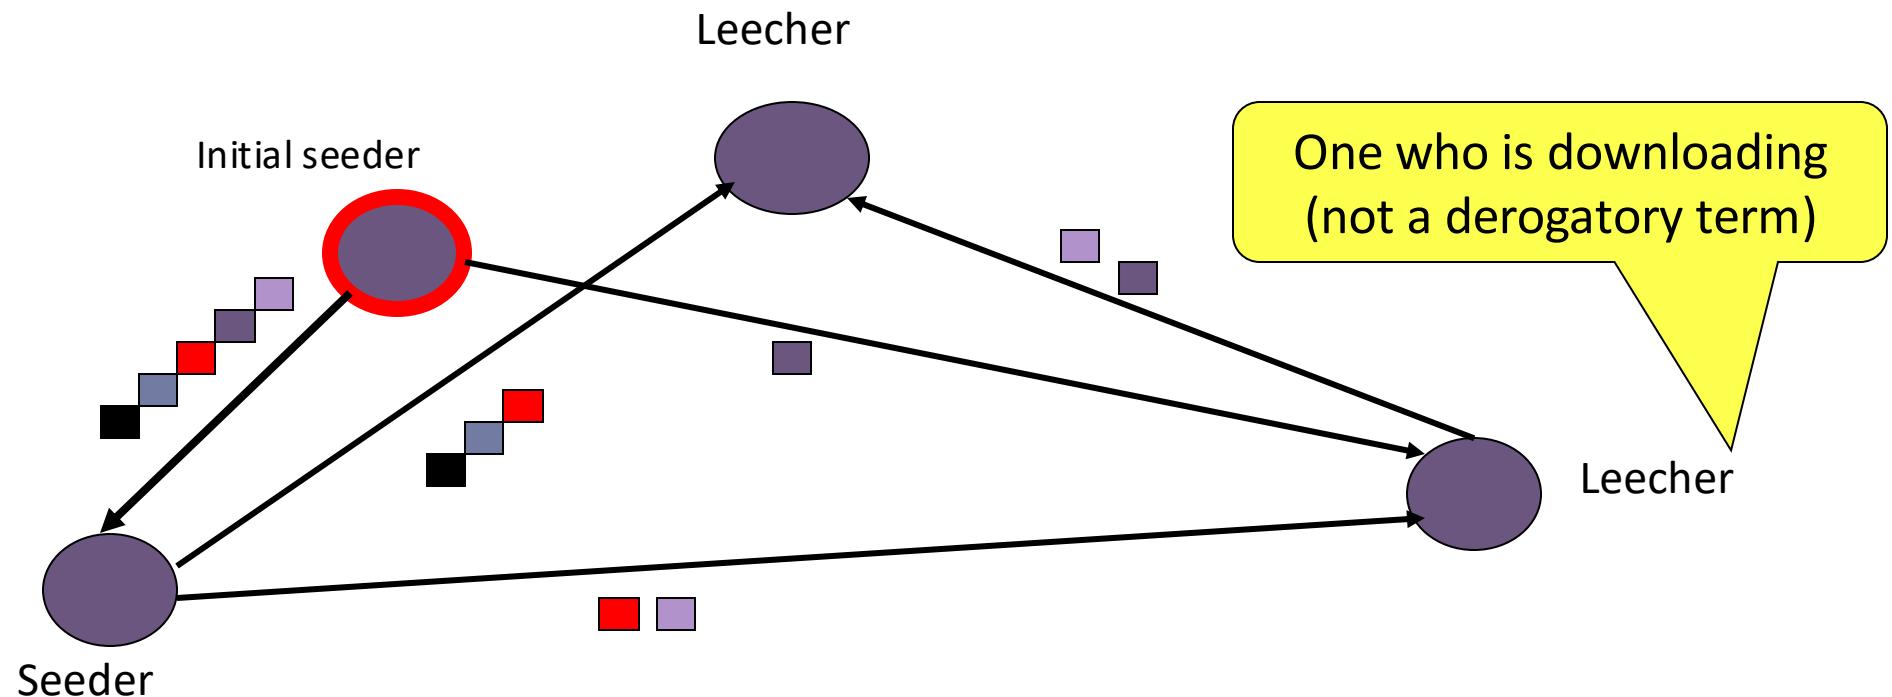
Torrent/swarm: group of peers exchanging chunks of a file.
(exchanging the same file)

churn: peers may come and go.
once peer has entire file, it may (selfishly)
leave or (altruistically)
remain in the torrent as **seeder**

BitTorrent Lingo

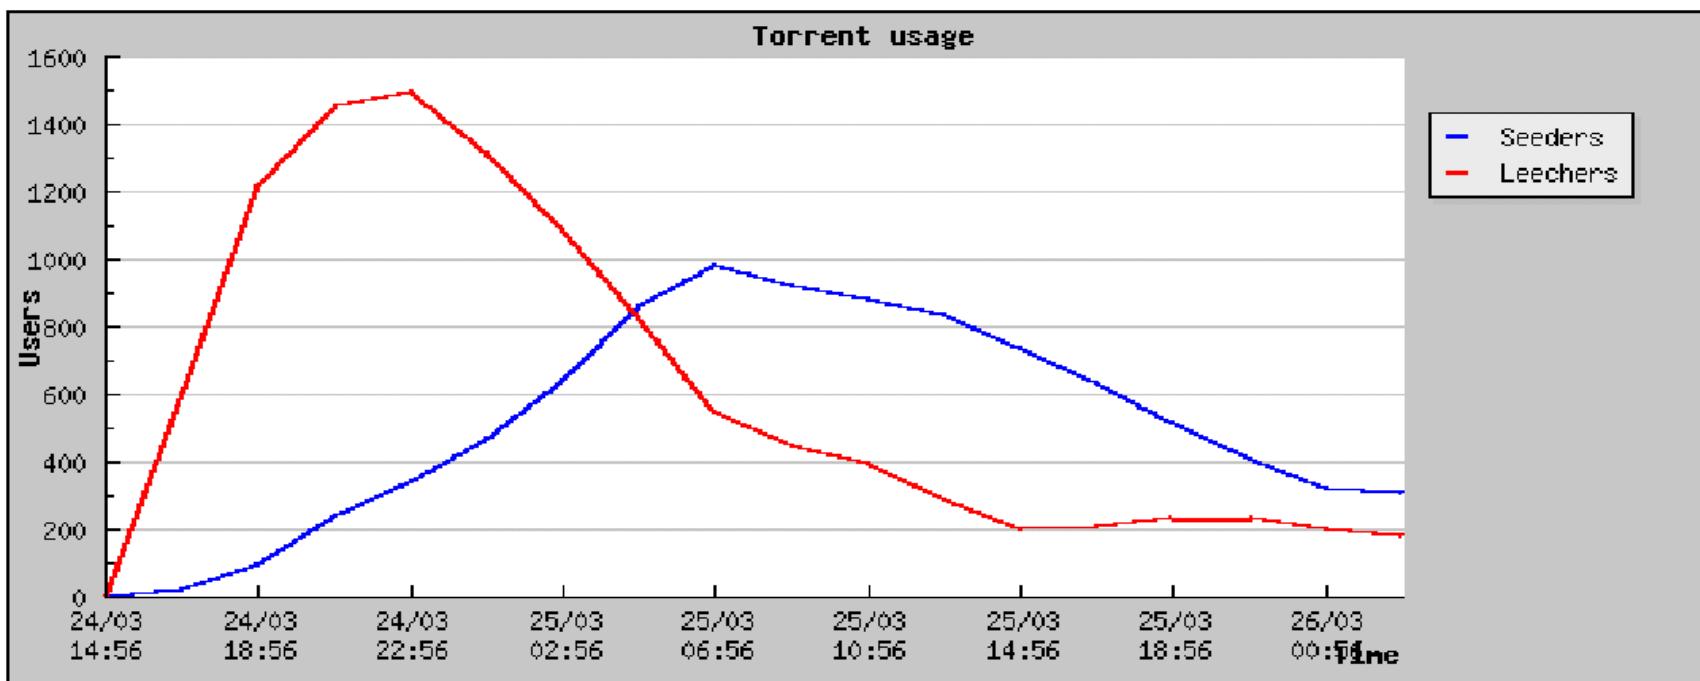
Seeder = a peer that provides the complete file.

Initial seeder = a peer that provides the initial copy.



Basic Idea

- ◆ As a leecher downloads pieces of the file, replicas of the pieces are created. ***More downloads mean more replicas available.***
- ◆ As soon as a leecher has a complete piece, it can potentially share it with other downloaders. Eventually each leecher becomes a seeder by obtaining all the pieces, and assembles the file. Verifies the checksum.



BitTorrent: requesting, sending file chunks

requesting chunks:

- ❖ at any given time, different peers have different subsets of file chunks
- ❖ periodically, Alice asks each peer for list of chunks that they have
- ❖ Alice requests missing chunks from peers.
- ❖ **Requires - Piece Selection Policies:**
 - ▶ The **order** in which pieces are selected by different peers is critical for good performance
 - ▶ If an inefficient policy is used, then peers may end up in a **situation where each has all identical set of easily available pieces**, and **none of the missing ones**.
 - ▶ If the original seed is prematurely taken down, then the file cannot be completely downloaded! What are “**good policies?**”

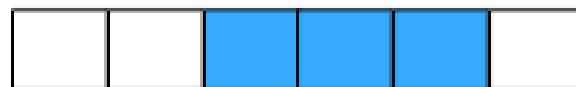
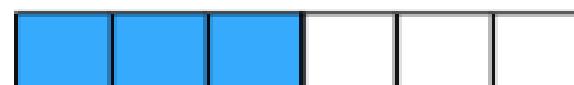
Piece Selection



peer 1



peer 2



peer 3



peer 4



Small overlap is good

Large overlap is bad
-- wastes bandwidth

Piece selection

- ▶ **Strict Priority**
 - ▶ Not common
- ▶ **Random First Piece**
 - ▶ Special case, at the beginning
- ▶ **Rarest First**
 - ▶ General rule
- ▶ **Endgame Mode**
 - ▶ Special case, at the end

Random First Piece

- ▶ Initially, a peer has nothing to trade
- ▶ Important to get a complete piece ASAP
- ▶ Select a random piece of the file and download it

Rarest Piece First

- ▶ Determine the pieces that are **most rare** among your peers, and download those first.
- ▶ **Logically**, the most commonly available pieces are left till the end to download.

Endgame Mode

- ◆ Near the end, missing pieces are requested from every peer containing them.
- ◆ This ensures that a download is not prevented from completion due to a single peer with a slow transfer rate.
- ◆ Some bandwidth is wasted, but in practice, this is not too much.

How to discourage free riding?

- ▶ Built-in **incentive** mechanism (where all the magic happens):
 - ▶ Choking Algorithm
 - ▶ Optimistic Unchoking

Choking

- ▶ **Choking** is a *temporary refusal* to upload to free riders. It is one of BT's most powerful idea to deal with **free riders (those who only download but never upload)**.
- ▶ ***Tit-for-tat strategy*** is based on game-theoretic concepts.

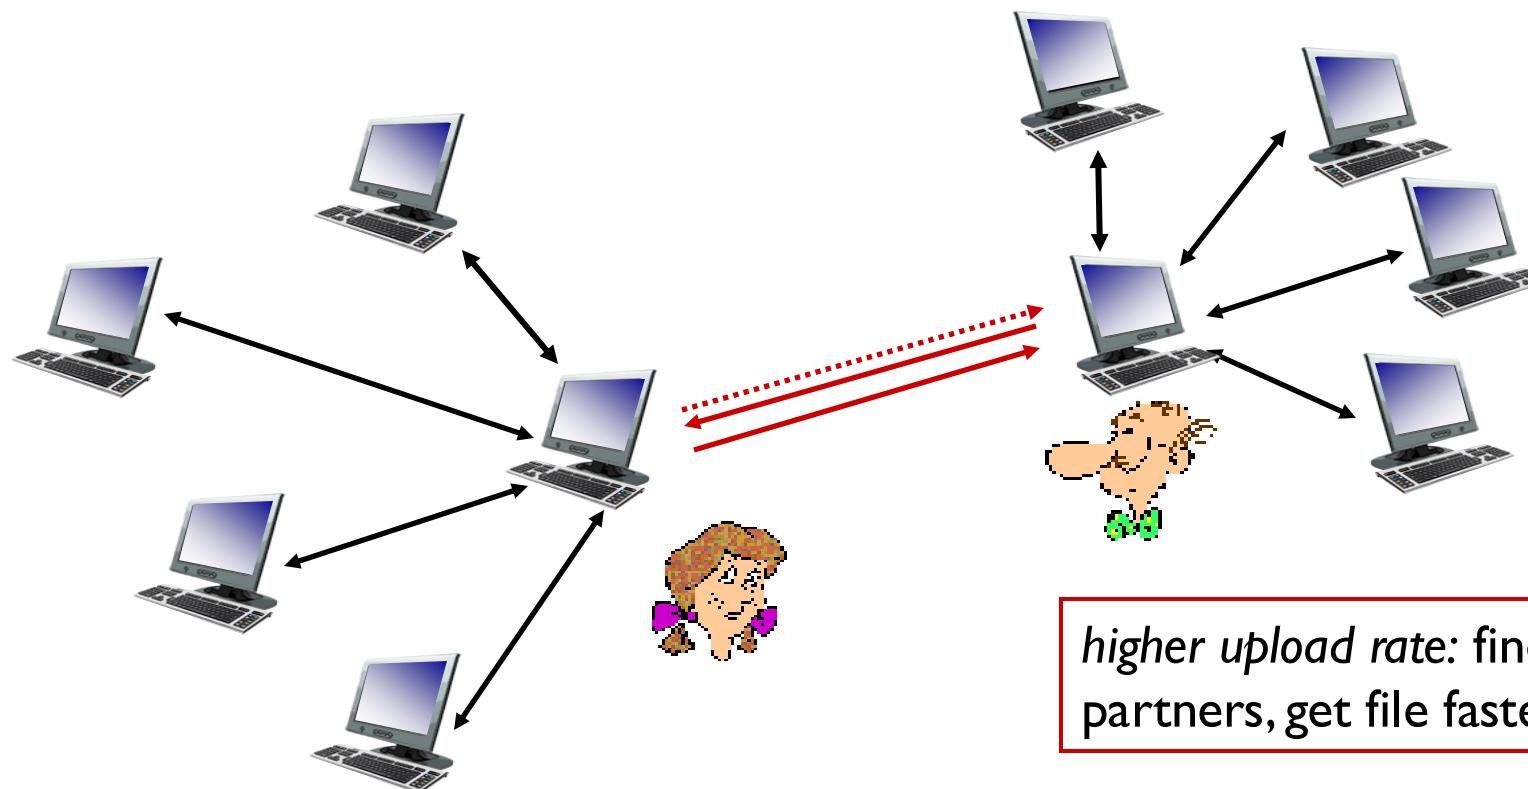
BitTorrent: requesting, sending file chunks

sending chunks: tit-for-tat

- ❖ Alice sends chunks to **four peers** currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate **top 4 every 10 secs**
- ❖ every 30 secs: randomly select another peer, starts sending chunks
 - “**optimistically unchoke**” this peer
 - newly chosen peer may join top 4 [See next slide]

BitTorrent: tit-for-tat

- (1) Alice “optimistically unchoke” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Upload-Only mode

- ▶ Once download is complete, a peer can **only upload**. The question is, which nodes to upload to?
- ▶ **Policy:** Upload to those with the best upload rate. This ensures that pieces get replicated faster, and new seeders are created fast

What is Trackerless torrents

- ▶ BitTorrent also supports "trackerless" torrents, featuring a DHT implementation that allows the client to download torrents that have been created without using a BitTorrent tracker.
- ▶ **Question to ponder:** How and where to maintain the peer database, when there is no centralised *tracker* ?

Questions about BT to discuss...

- ▶ What is the effect of bandwidth constraints?
- ▶ Is the Rarest First policy really necessary?
- ▶ Must nodes perform seeding after downloading is complete?
- ▶ How serious is the Last Piece Problem?
- ▶ Does the incentive mechanism affect the performance much?
- ▶ What if the peers are behind NAT (using private IP)?

Some Decentralised Systems/Apps (p2p)

For your reading pleasure (FRYOP)...



Enterprises

Devs

Community

Explorer

Wallet

About

Docs



Reinventing DePIN: resource sharing economy

The Decentralized Cloud for
AI, Media & Entertainment



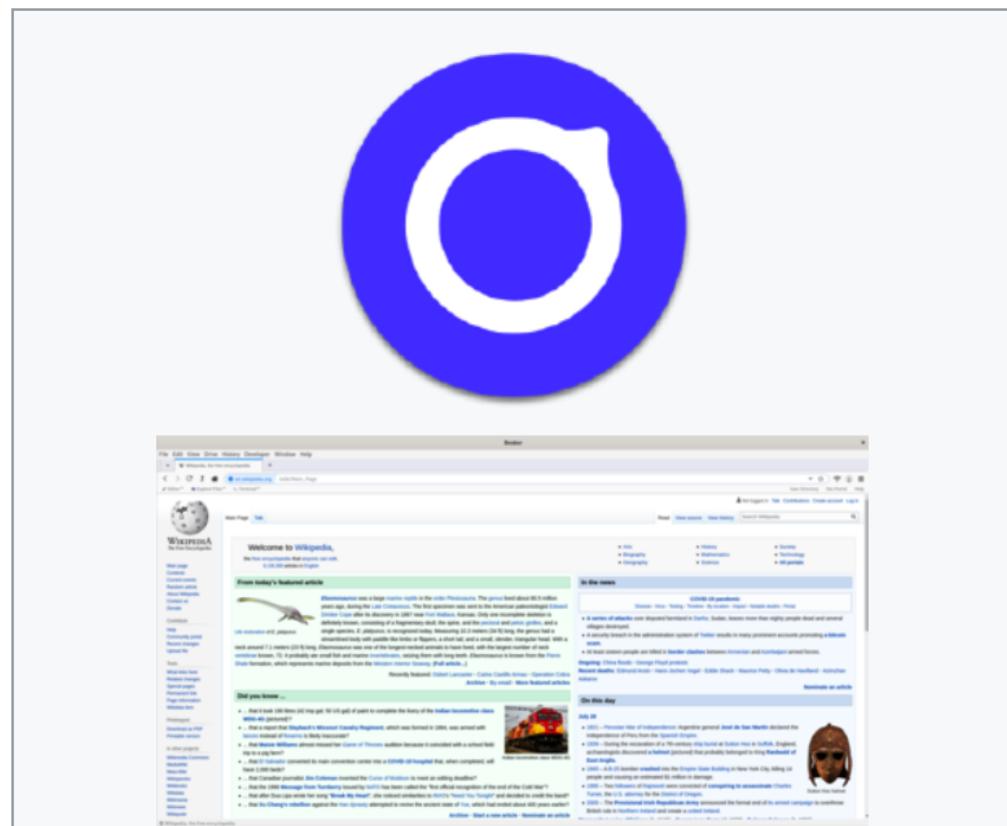
THETA

Beaker is a discontinued^[4] free and open-source web browser^[5] developed by Blue Link Labs.^{[6][7]}

Beaker Browser peer-to-peer technology allows users to self-publish websites and web apps^[8] directly from the browser, without the need to set up and administrate a separate **web server** or host their content on a third-party server. All files and websites are transferred using **Dat**, a hypermedia peer-to-peer protocol, which allows files to be shared and hosted by several users.^[9]

The browser also supports the **HTTP protocol** to connect to traditional servers.^[5]

Beaker



Screenshot of the Beaker browser on Fedora
31

Developer(s) Blue Link Labs.^[1]

Initial release 1 August 2017; 7 years ago

Stable release(s)

1.1.0 (December 8, 2020; 3 years ago)^[2] [+] [↑]

[Browser](#)[Brave Search](#)[Why Brave](#)[Search API](#)[Private advertising](#)[Get Brave](#)[Search The Web](#)

The browser that puts you first

Block ads. Save data. And get way faster webpages. Just by switching your browser.

[Get Brave](#)

Private search

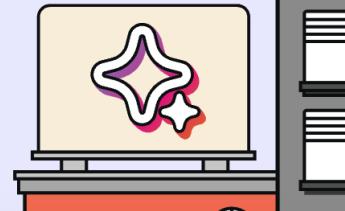
Search the web privately...



Better results, less SEO spam, and zero profiling.

Built-in AI assistant

Get answers, generate content, & more. Right in the browser.



Powerful VPN

Protect every app, on your entire device.



Speak Freely

Say "hello" to a different messaging experience. An unexpected focus on privacy, combined with all of the features you expect.

[Get Signal](#)

← → ⌂ filecoin.io

Deploy smart contracts on Filecoin's Virtual Machine →

EN 中文

f

Store Provide Build Blog

Explore the Network



Filecoin is a decentralized storage network designed to store humanity's most important information.

▶ Explore the Filecoin vision

Deploy smart contracts on Filecoin's Virtual Machine →

- ▶ Uses IPFS IPFS (InterPlanetary File System) Protocol

- ▶ P2P transactions for lending, borrowing, and trading assets without intermediaries.



THE END

- ▶ Principles of Network Applications
- ▶ HTTP & WWW
- ▶ E-mail Protocols (SMTP, POP, IMAP)
- ▶ P2P Applications

Activities:

LAB2 goes virtual with mini-net. It is take-home Lab.

You will have post lab Quiz

Assignment 2 (NW tools) will be released tomorrow.

Q: How bitcoin's p2p network works?

Q: What are the problems with current c/s social apps? Can that be solved using p2p?



Attendance

<https://inetapps.nus.edu.sg/ctr/>

