

CS3103 Computer Networks Practice

IP-options, ICMP, NAT-Hole Punching

IPv4 – Options Processing
Internet Control Message Protocol (ICMP)
Ping and Traceroute
NAT Traversal (Port Forwarding and Hole Punching)

(Things to know for Assignment 3b and OSPF lab)

Dr Anand Bhojan

COM3-02-49, School of Computing

banand@comp.nus.edu.sg ph: 651-67351

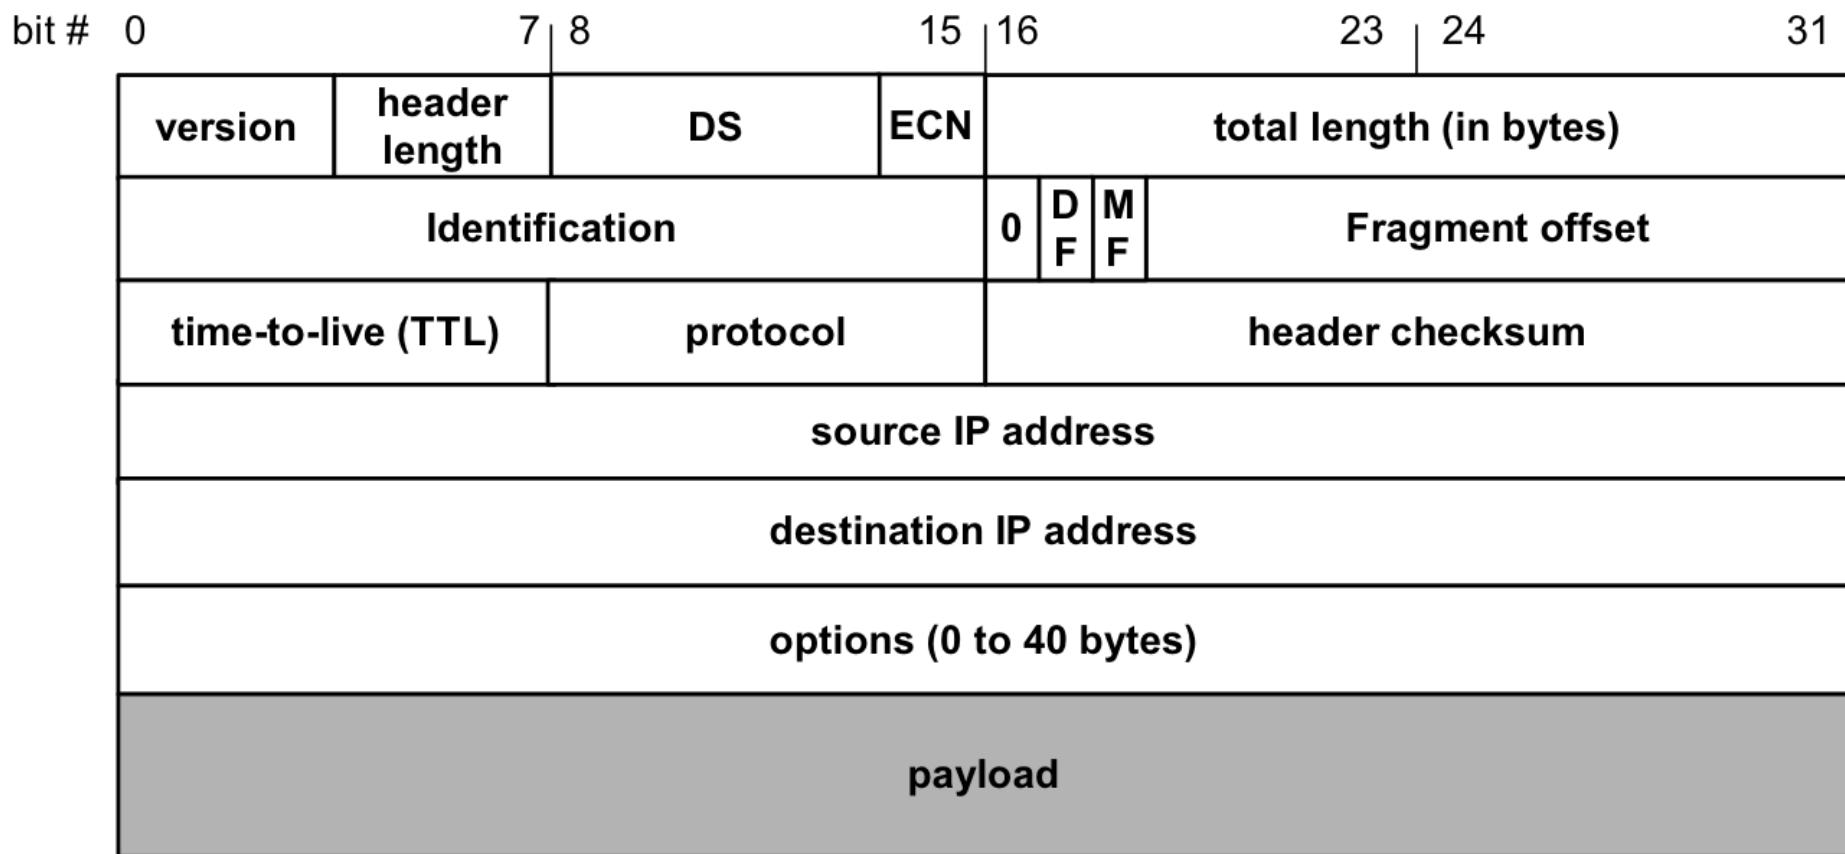
Learning Objectives

- ▶ Bridge-In (Recap) – Recap on IP Datagram Format & IP Fragmentation
- ▶ Understand IPv4 Options Processing – Record route and Source route
- ▶ Understand ICMP protocol and its applications
- ▶ Understand NAT traversal methods

Related Assessments: Assignment, Labs

FRORP – For Your Own Reading Pleasure. May not be discussed in class.

IP Datagram Format. - For your reference.



- ▶ $20 \text{ bytes (min header)} \leq \text{Header Size} < 2^4 \times 4 \text{ bytes} = 64 \text{ bytes}$
- ▶ $20 \text{ bytes (min header)} \leq \text{Total Length} < 2^{16} \text{ bytes} = 65536 \text{ bytes}$

- ▶ **Question:** In which order are the bytes of an IP datagram transmitted?
- ▶ **Answer:**
 - ▶ Transmission is row by row
 - ▶ For each row:
 1. First transmit bits 0-7
 2. Then transmit bits 8-15
 3. Then transmit bits 16-23
 4. Then transmit bits 24-31
- ▶ This is called **network byte order** or **big endian** byte ordering.
- ▶ **Note:** Many computers (incl. Intel processors, Apple M1 processors) store 32-bit words in little endian format..

Big endian vs. small endian

FRORP

- Conventions to store a multi-byte word

- Example: a 4 byte Long Integer

Byte3 Byte2 Byte1 Byte0

LittleEndian

- ▶ Stores the low-order byte at the lowest address and the highest order byte in the highest address.

Base Address+0 Byte0

Base Address+1 Byte1

Base Address+2 Byte2

Base Address+3 Byte3

- ▶ Intel processors use this order

BigEndian

- ▶ Stores the high-order byte at the lowest address, and the low-order byte at the highest address.

Base Address+0 Byte3

Base Address+1 Byte2

Base Address+2 Byte1

Base Address+3 Byte0

- LEON, Motorola processors use big endian.



Differentiated Services

▶ DS/ECN field (1 byte)

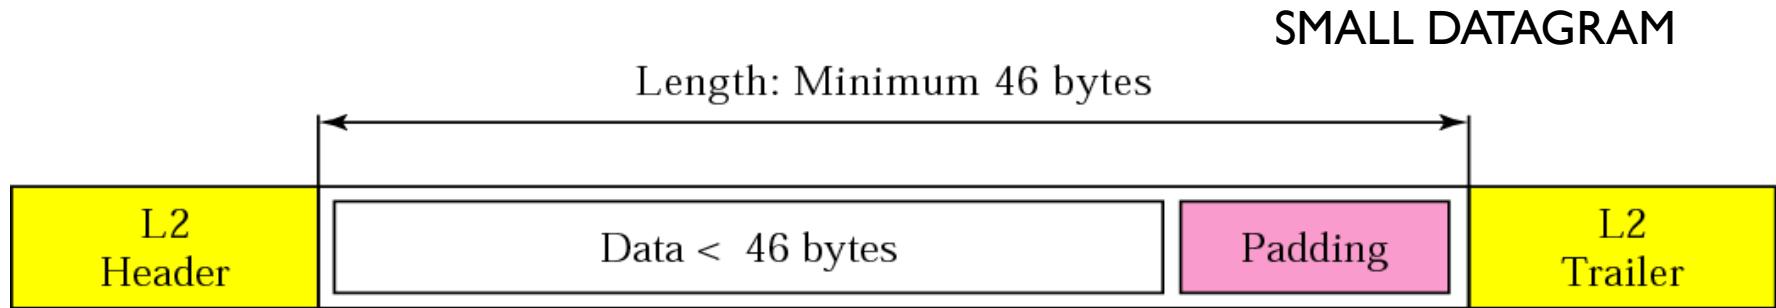
- ▶ This field was previously called as Type-of-Service (TOS) field. The role of this field has been re-defined, but is “backwards compatible” to TOS interpretation
- ▶ Differentiated Service (DS) (6 bits):
 - ▶ Used to specify service level (currently not supported in the Internet). Under DiffServ, **DSCP (Differentiated Services Code Point)** is defined.
- ▶ Explicit Congestion Notification (ECN) (2 bits):
 - ▶ New feedback mechanism used by TCP (Most server support ECN today)



ECN

- 00 – Non ECN-Capable Transport, Non-ECT
- 10 – ECN Capable Transport, ECT(0)
- 01 – ECN Capable Transport, ECT(1)
- 11 – Congestion Encountered, CE.

Encapsulation of a small datagram in an Ethernet frame



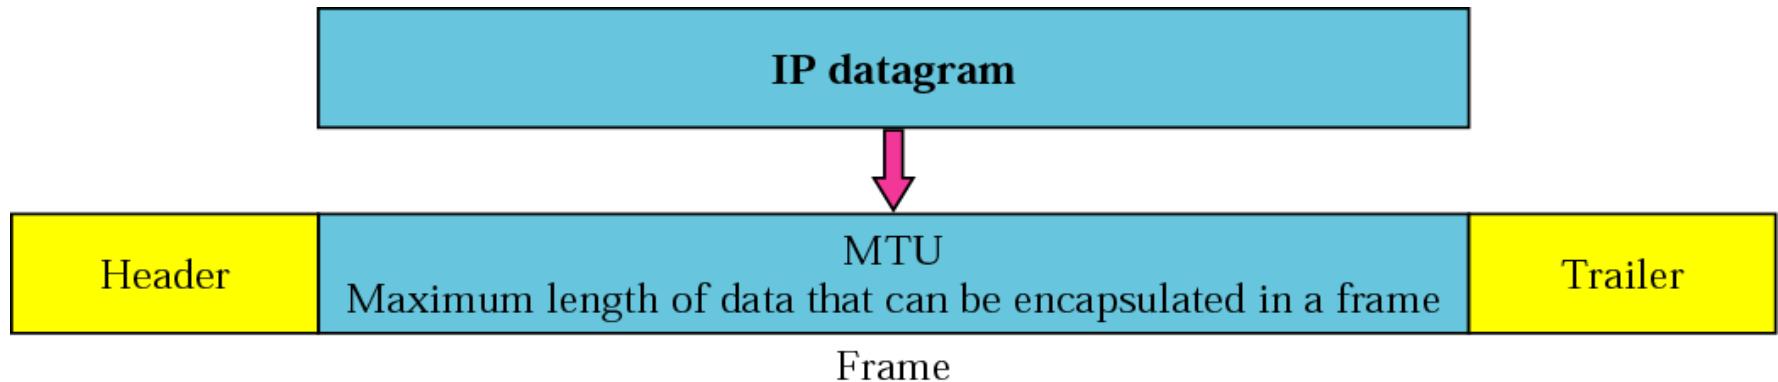
Q: Why there is a minimum length requirement?

Make sure you **LOGIN** using your NUSNET ID.

Encapsulation of a big datagram in an Ethernet frame

MTU – Maximum Transmission Unit

BIG DATAGRAM

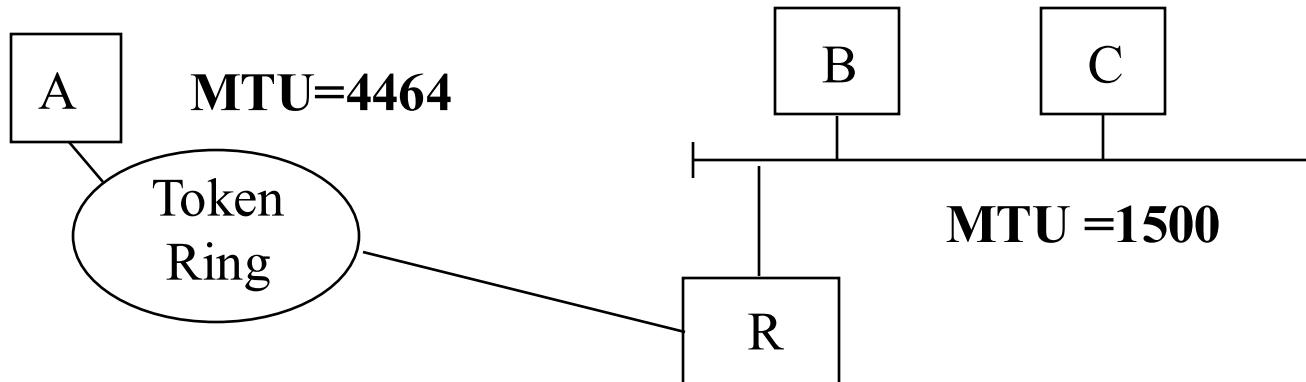


**Q: Why there
is a maximum
length
requirement
for Ethernet?**

<i>Protocol</i>	<i>MTU</i>
Hyperchannel	65,535
Token Ring (16 Mbps)	17,914
Token Ring (4 Mbps)	4,464
FDDI	4,352
Ethernet	1,500
X.25	576
PPP	296

IP Fragmentation

fragmentation total length = 16 bits, offset = 13 bits the 3 bits is for 0, reserved for future use, bit 1 for dont frag, bit 2 - more frag means that there are more frag-packets



- IP Fragmentation takes place because of physical level limitation - MTU.
- Fragmentation takes place at a **Router** or **Original Sending Host**.
- Fragments are reassembled at the **destination host**.

Segmentation and Fragmentation

UDP

TCP

← 2000 bytes →

Application data

UDP Application data

IP UDP Application data

IP UDP

1480 bytes

540 bytes

← 2000 bytes →

Application data

TCP

TCP

IP TCP

IP TCP

1480 bytes

560 bytes

TCP Segmentation at TCP layer

IP Fragmentation at IP layer

Example

- An IP datagram with 636 bytes of data (IP payload) arrives at a router and is destined to a network with a MTU of 256 bytes. Assume 20 bytes IP header.
- IP Header (recap) - 3-bit Flag:-
 - D: Do not fragment
 - M: More fragments



Example

636 bytes of data to be transferred. MTU is 256 bytes.

- Maximum possible data length in each fragment
 $256 - 20 = 236$ bytes
- Min number of fragments required = $636/236 = 3$
- But length must be divisible by 8, except for last fragment
- Revised length of each fragment is 232, 232, 172

Question to ponder: Why the length must be divisible by 8?



Example (cont)

- First Fragment:

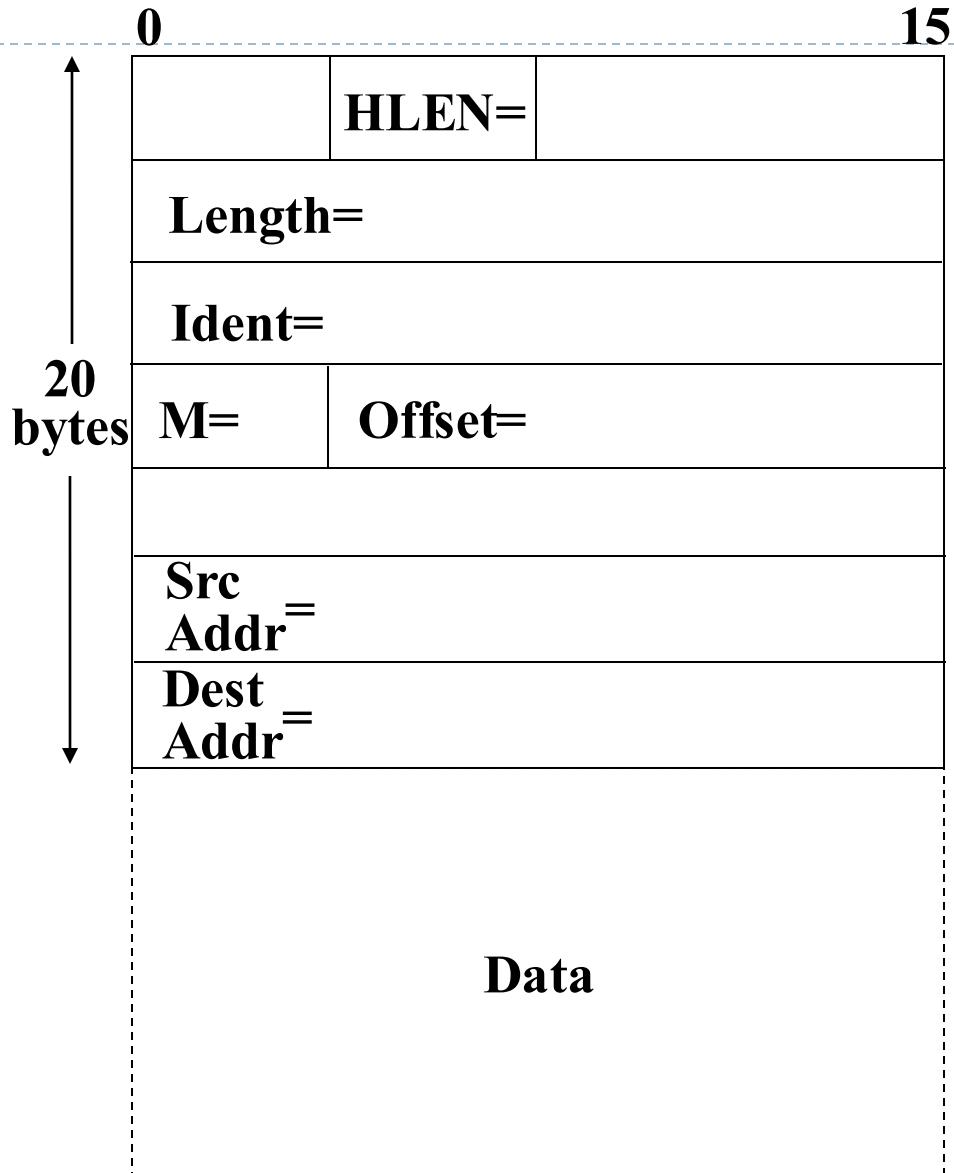
HLEN=5;

Fragment length = 232

Total Length =
 =
 =
 =252

Fragment Offset=0

More =1



Example (cont)

- Second Fragment:

HLEN=5;

Fragment length = 232

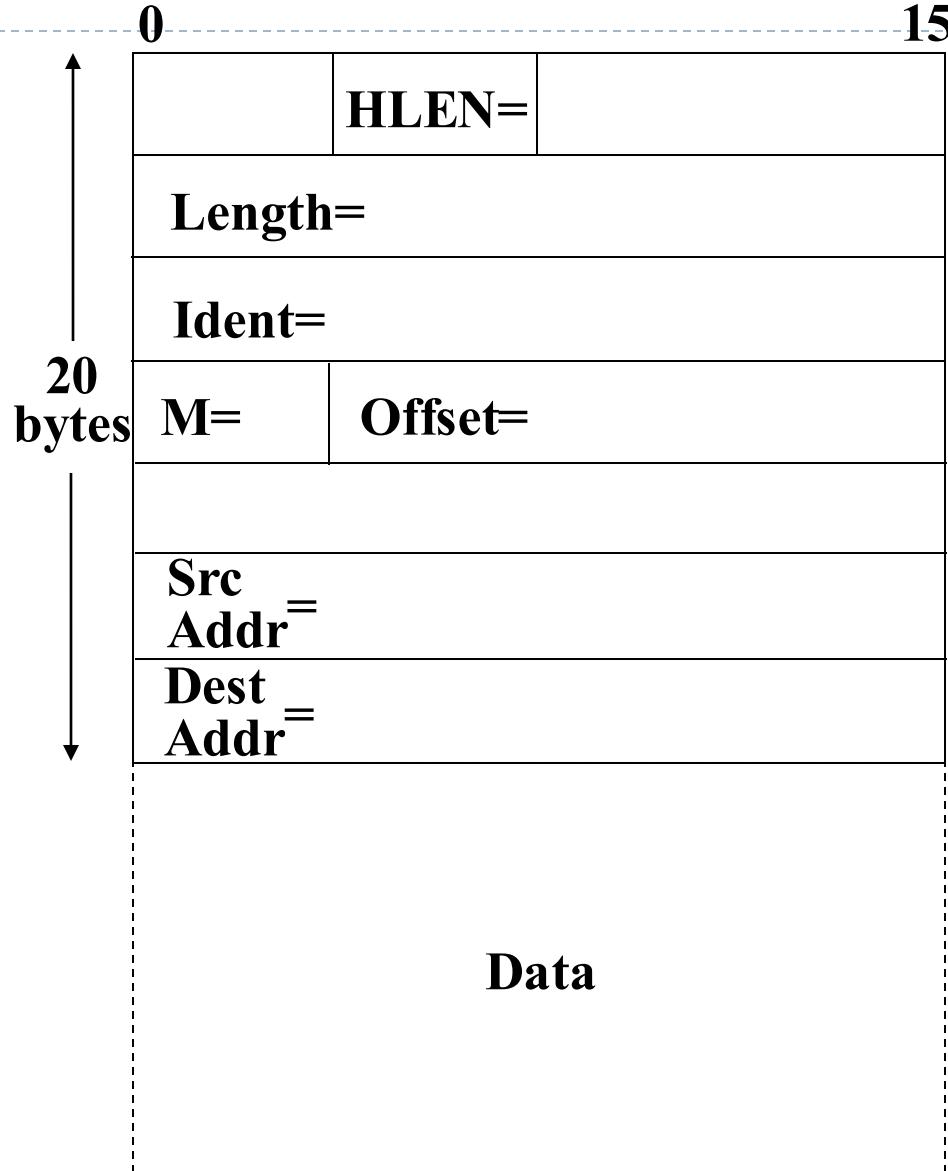
Total Length=252

Fragment Offset=29

More=1

Question to ponder: Why fragment offset is 29 instead of 232?

Why can't we have it as 232? save space?



Example (cont)

- Third Fragment

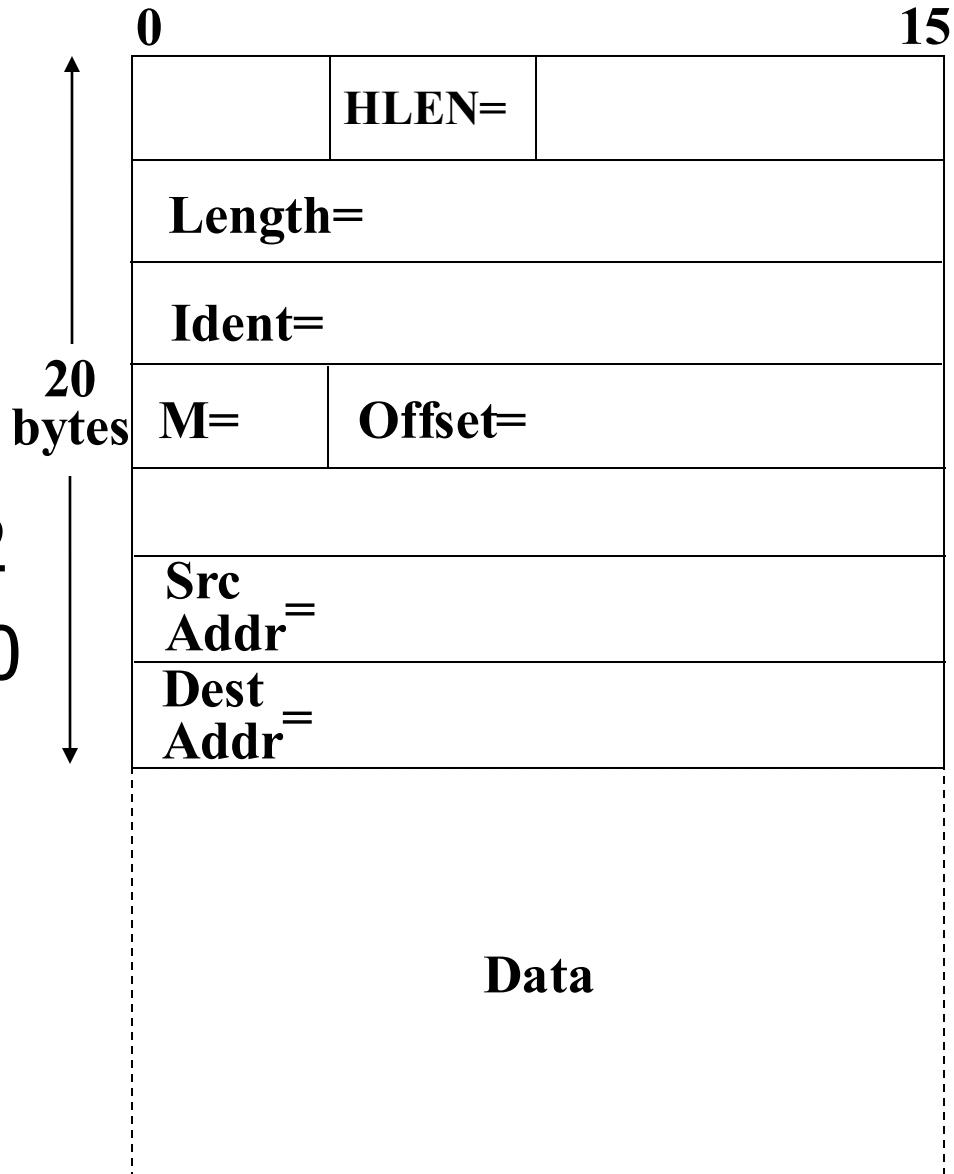
HLEN=5;

Fragment length=172

$$\begin{aligned}\text{Total Length} &= 172 + 20 \\ &= 192\end{aligned}$$

Fragment offset=58

More=0



EXAMPLES

- 1) *A packet has arrived with an M bit value of 1 and a fragmentation offset value of zero. Is this the first fragment, the last fragment, or a middle fragment?.*
- 2) *A packet has arrived in which the offset value is 100, the value of HLEN is 5 and the value of the total length field is 100. What is the number of the first byte and the last byte of the IP payload?*



IP-options, ICMP, NAT-Hole Punching

IPv4 – Options Processing

Internet Control Message Protocol (ICMP)

Ping and Traceroute

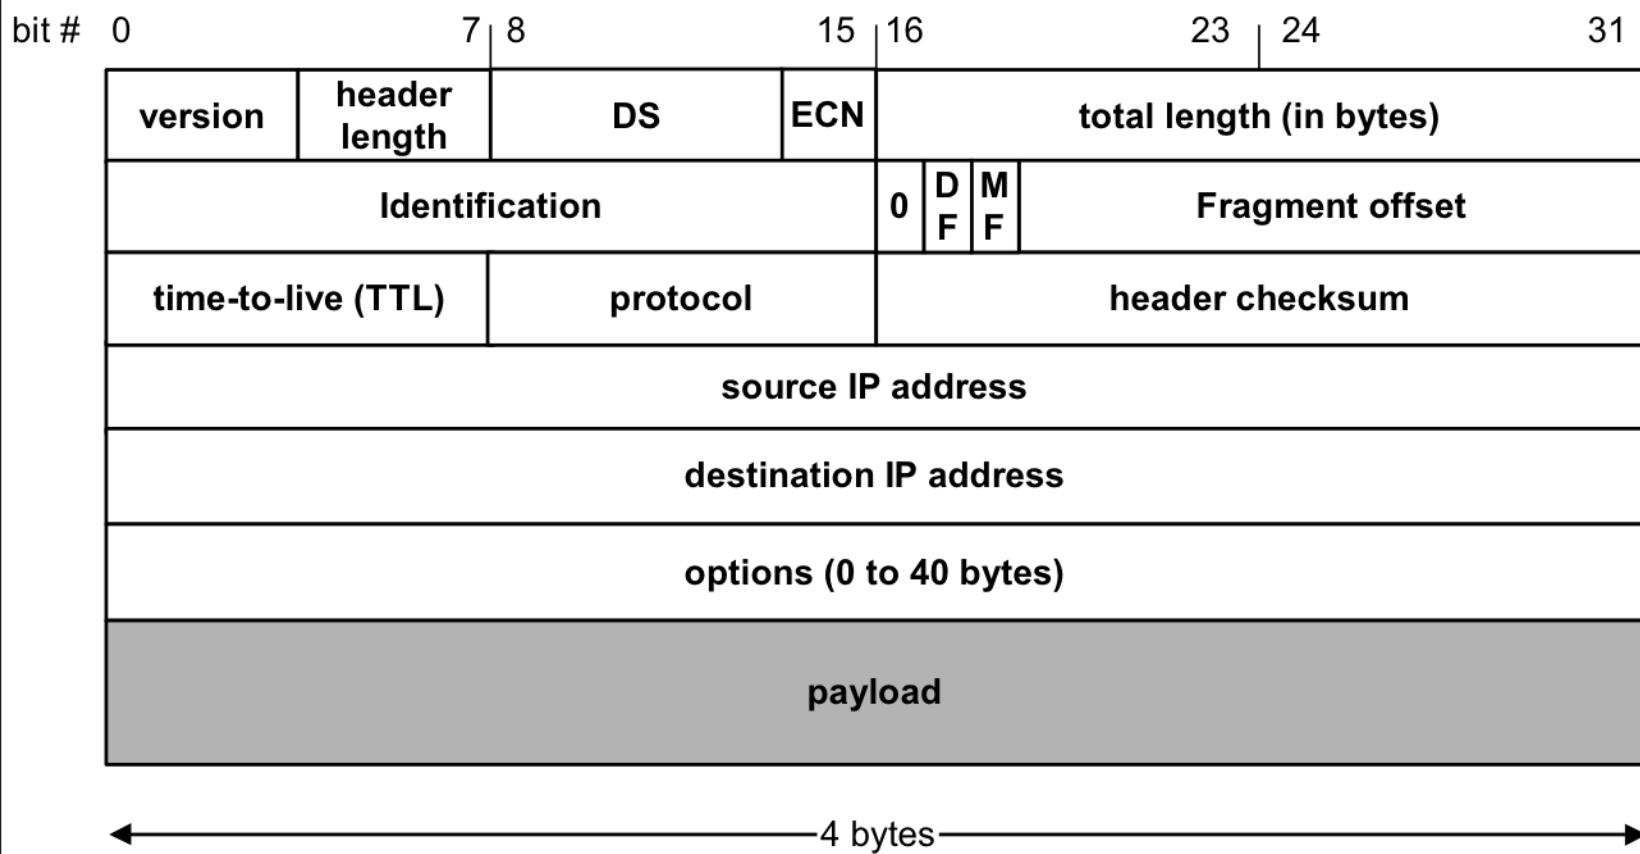
NAT Traversal (Port Forwarding and Hole Punching)

Dr Anand Bhojan

COM3-02-49, School of Computing

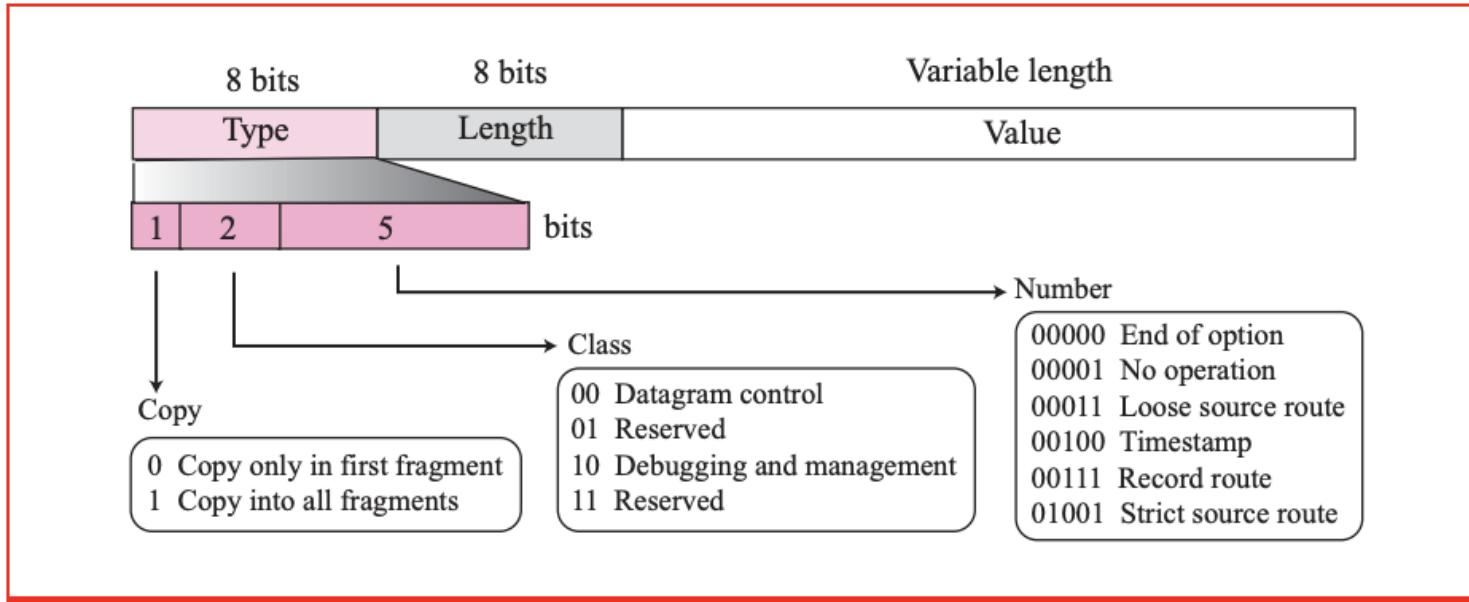
banand@comp.nus.edu.sg ph: 651-67351

OPTIONS



- Q: Given header-length is only 4 bits, how the header size is 60 bytes?**
- Q: Given, these 40 bytes are optional, do every routers and end-hosts need to process the options?**

Option format



- First bit – 0 (Copy the option only in first fragment); 1 (Copy in all fragments)
- Second and Third Bits – General purpose of the option (option category). 00 (Datagram Control), 01 (Reserved), 10 (Debugging and Management), 11(Reserved)
- Remaining 5 bits gives 32 options. We have only 6 options defined so far.

[Length and Value are optional]

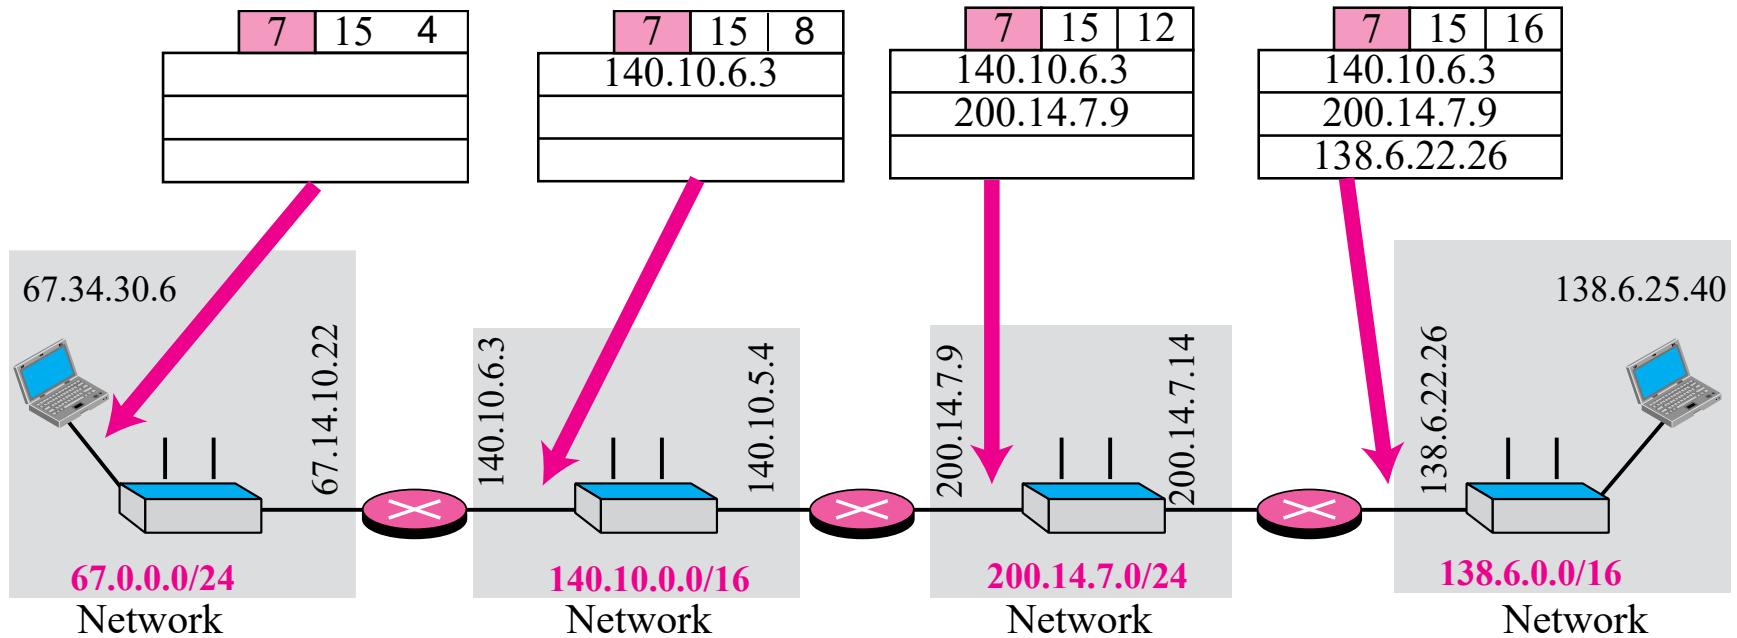
Length is total length including the Type, and Length field itself.

Record-route option

Only 9 addresses
can be listed.

Type: 7 00000111	Length (Total length)	Pointer
	First IP address (Empty when started)	
	Second IP address (Empty when started)	
	⋮	
	Last IP address (Empty when started)	

Record-route concept

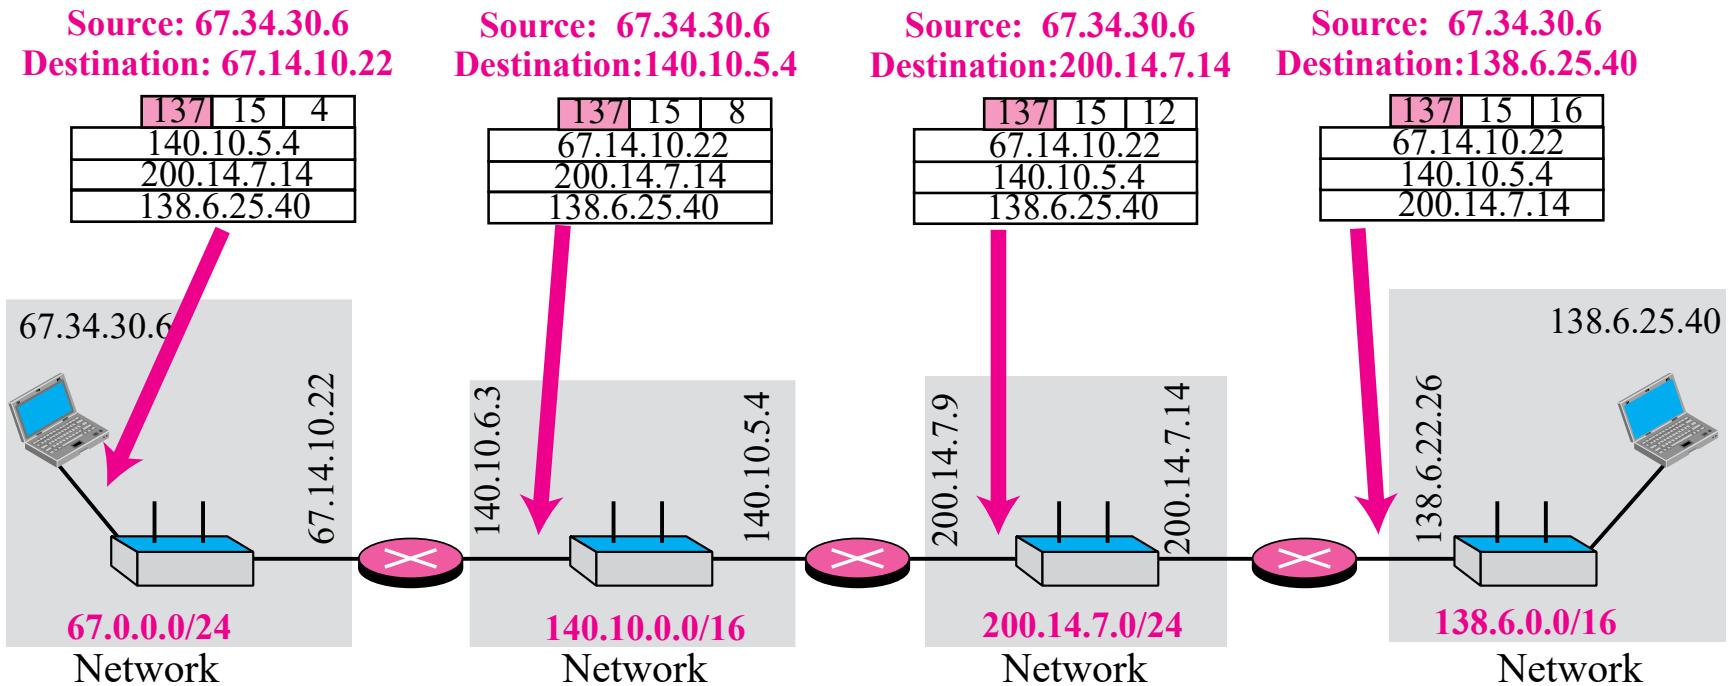


Strict-source-route option

Only 9 addresses
can be listed.

Type: 137 10001001	Length (Total length)	Pointer
	First IP address (Filled when started)	
	Second IP address (Filled when started)	
	• • •	
	Last IP address (Filled when started)	

Strict-source-route option



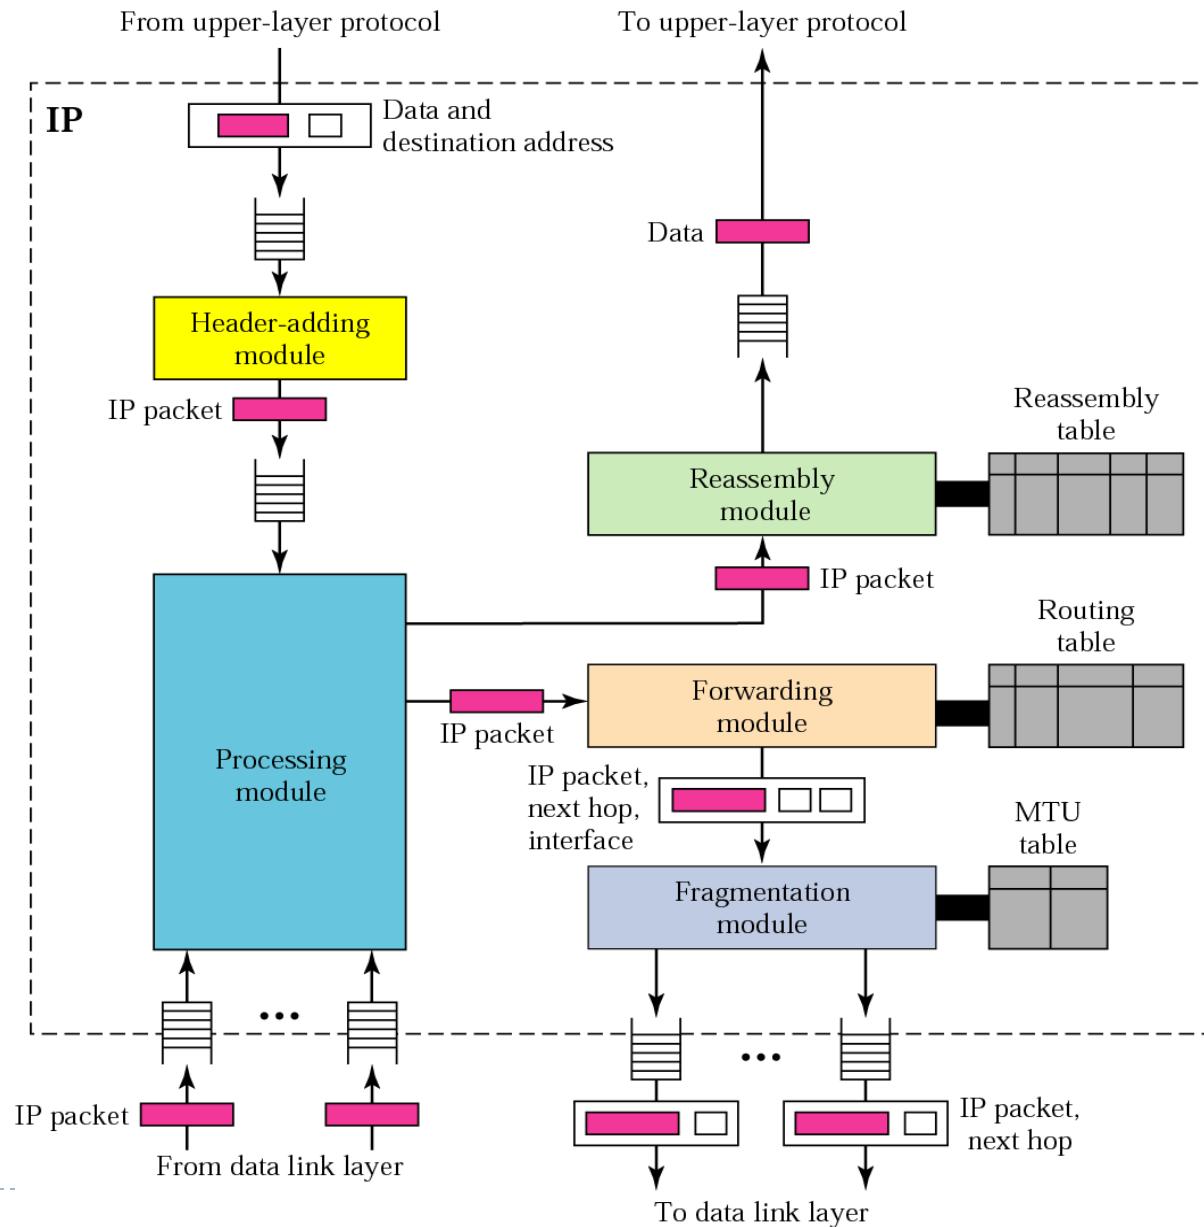
Loose-source-route option

Only 9 addresses
can be listed.

Type: 131 10000011	Length (Total length)	Pointer
	First IP address (Filled when started)	
	Second IP address (Filled when started)	
	• • •	
	Last IP address (Filled when started)	



IP components



CS3103 Computer Networks Practice

IP-options, ICMP, NAT-Hole Punching

IPv4 – Options Processing

Internet Control Message Protocol (ICMP)
Ping and Traceroute

NAT Traversal (Port Forwarding and Hole Punching)

Dr Anand Bhojan

COM3-02-49, School of Computing

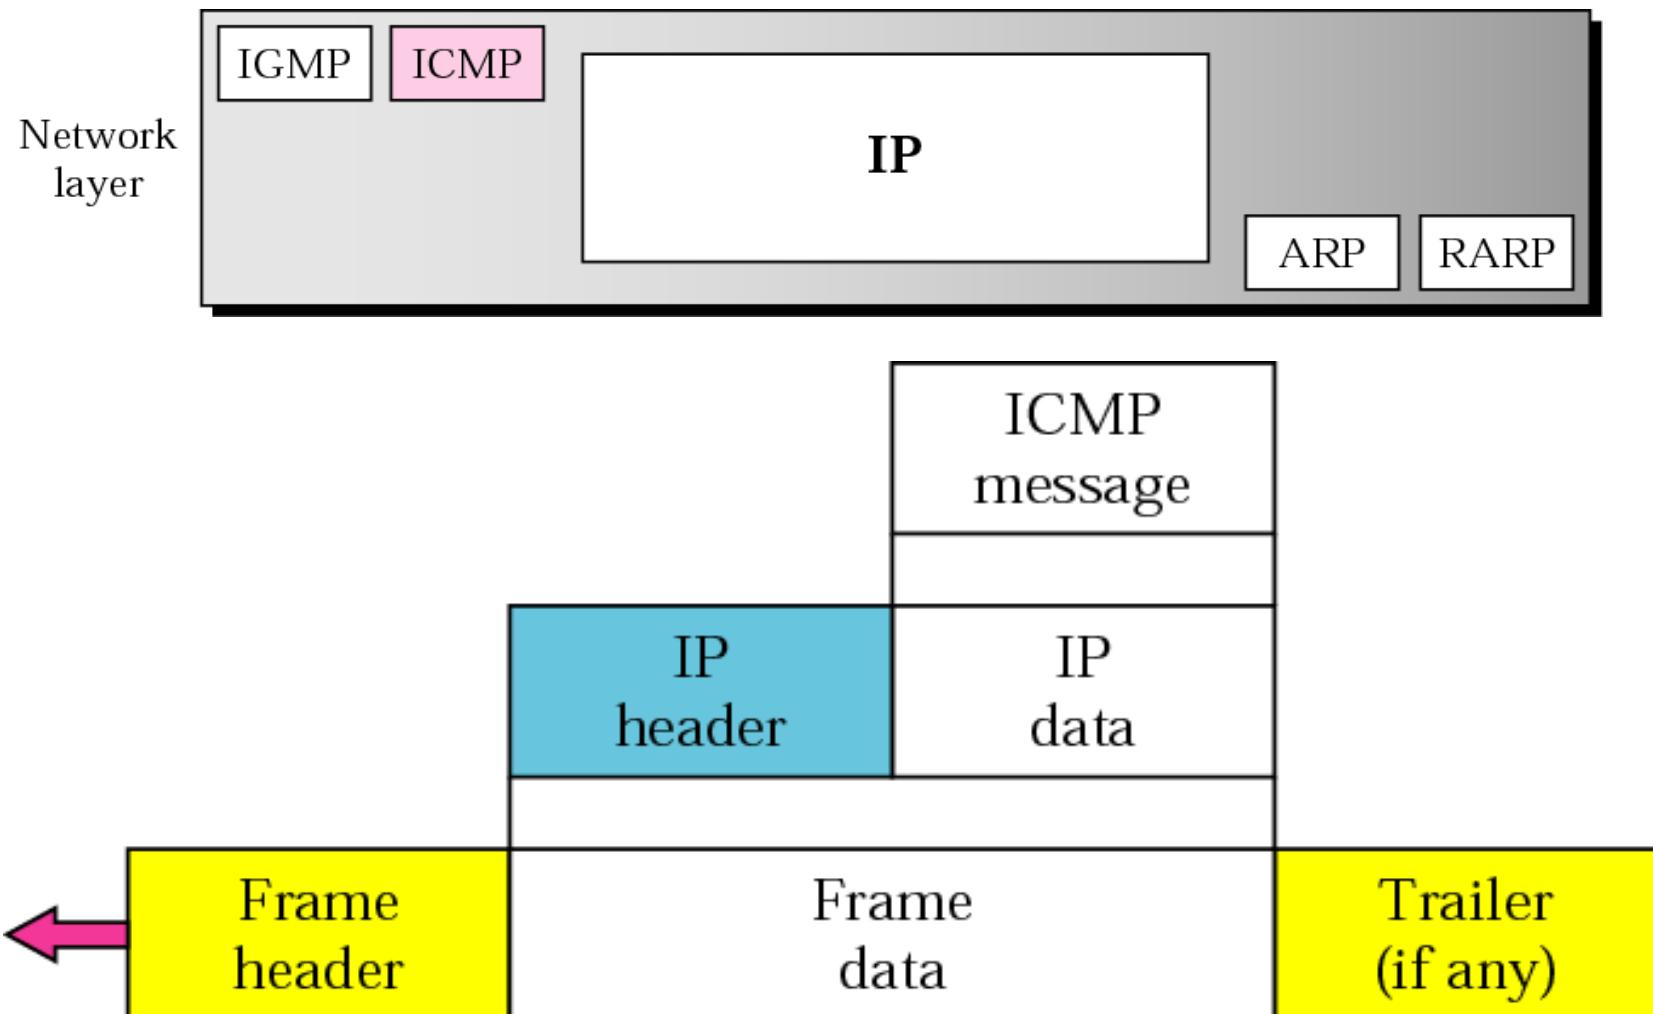
banand@comp.nus.edu.sg ph: 651-67351

Why ICMP?

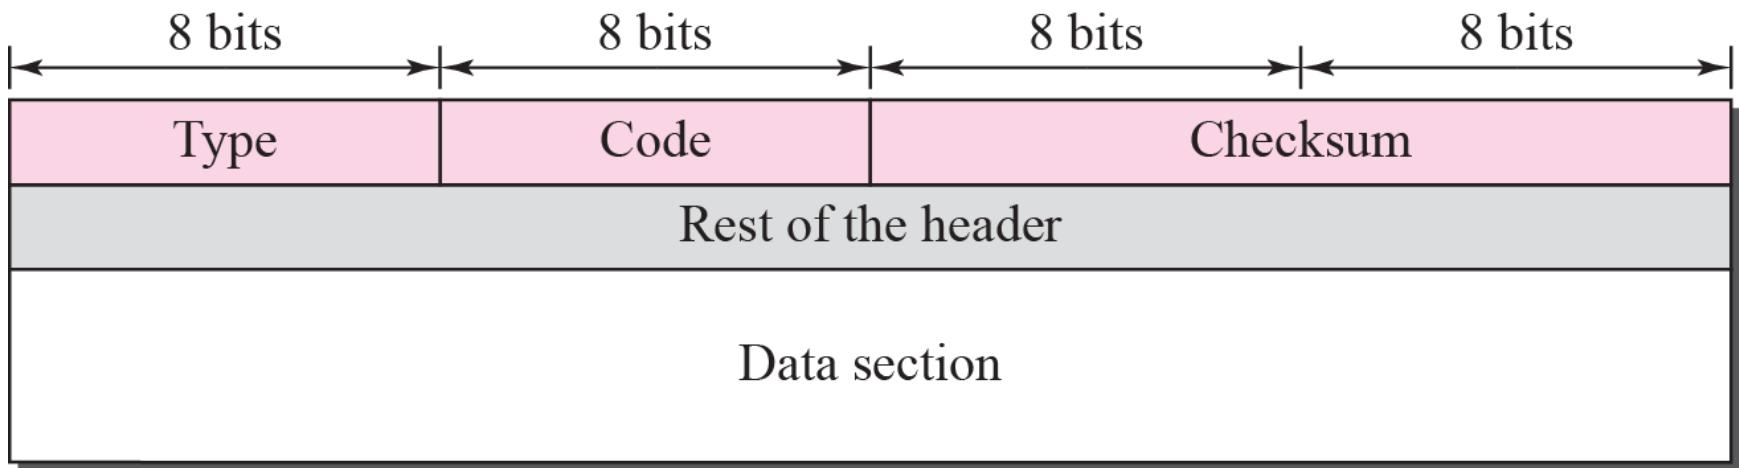
- ▶ *Q:What if a router cannot route or deliver a datagram?*
- ▶ *Q:What if a router experiences a congestion?*
- ▶ *Q:What if the TTL expires?*
- ...
- ▶ Router needs to inform source to take action to avoid or correct the problem
- ▶ ICMP is an error reporting mechanism, and can only report error condition **back to the original source**
- ▶ In addition, ICMP allows **routers and hosts** to send **information or control messages** to other **routers or hosts**.
- ▶ ICMP is specified in **RFC 792**



Position of ICMP in the network layer & Encapsulation



General format of ICMP messages



- 8 byte header, variable size data section
- format for first 4 bytes of header is common to all ICMP packets
- Type – ICMP message type
- Code – reason for the message type generated

ICMP MESSAGES

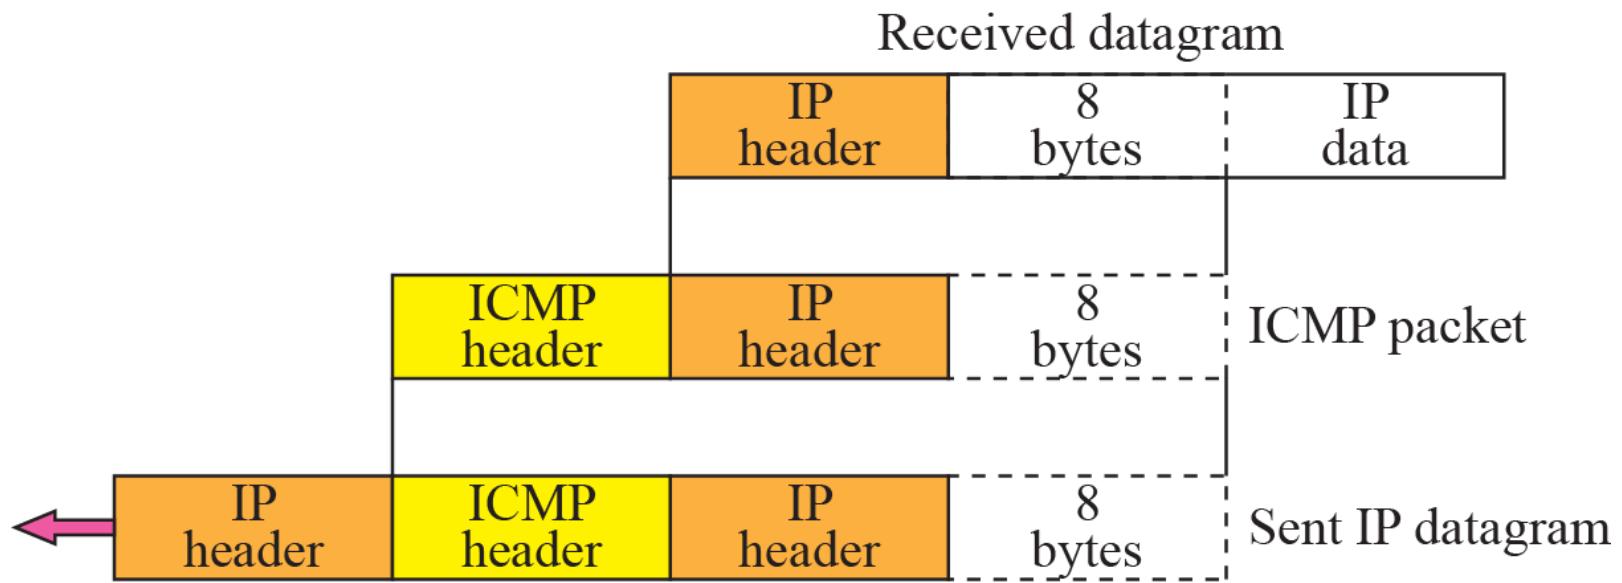
<i>Category</i>	<i>Type</i>	<i>Message</i>
Error-reporting messages	3	Destination unreachable
	4	Source quench
	11	Time exceeded
	12	Parameter problem
	5	Redirection
Query messages	8 or 0	Echo request or reply
	13 or 14	Timestamp request or reply

The **error-reporting messages** report problems that a router or a host (destination) may encounter when it processes an IP packet.

The **query messages**, which occur in pairs, help a host or a network manager get specific information from a router or another host.



Contents of data field for the error message



Q: Why the IP header of the received datagram is included in the Error message?

Q: Why the first 8-bytes of the received datagram is included in the Reply?

ERROR MESSAGES

ICMP error messages are used by routers and hosts to tell a device that sent a datagram about problems encountered in delivering it.

- ▶ **Q1:** What if a datagram carrying ICMP error message causes another error?
- ▶ **Q2:** Do we need ICMP error message for each fragment of a fragmented datagram that causes the error?
- ▶ **Q3:** ICMP error messages will not be generated for a datagram having a multicast or broadcast address in the destination address field, Why?
- ▶ **Q4:** ICMP error messages will not be for a datagram whose source address is not a single host, 0.0.0.0, 127.x.x.x, broadcast or multicast address, Why?

ICMP error messages - Examples

Type: 3	Code: 0 to 15	Checksum
Unused (All 0s)		
Part of the received IP datagram including IP header plus the first 8 bytes of datagram data		

(type 3) Destination-unreachable format

Codes:

0: network unreachable

1: host unreachable

2: protocol unreachable

3: port unreachable

4: need fragmentation but DF flag is set

.....

Q: Who generates these ICMP error messages, hosts or routers?



ICMP error messages - Examples

Type: 4	Code: 0	Checksum
Unused (All 0s)		
Part of the received IP datagram including IP header plus the first 8 bytes of datagram data		

Source-quench format

- *A source-quench message informs the source that a datagram has been discarded due to congestion in a router or the destination host.*
- *The source must slow down the sending of datagrams until the congestion is relieved.*
- *One message sent for each datagram discarded.*



ICMP error messages - Examples

Type: 11	Code: 0 or 1	Checksum
Unused (All 0s)		
Part of the received IP datagram including IP header plus the first 8 bytes of datagram data		

Time-exceeded message format

Two possible cases:

- ***Whenever a router decrements a datagram with a time-to-live value to zero, it discards the datagram and sends a time-exceeded message to the original source.***
- ***When the final destination does not receive all of the fragments in a set time, it discards the received fragments and sends a time-exceeded message to the original source.***



ICMP error messages - Examples

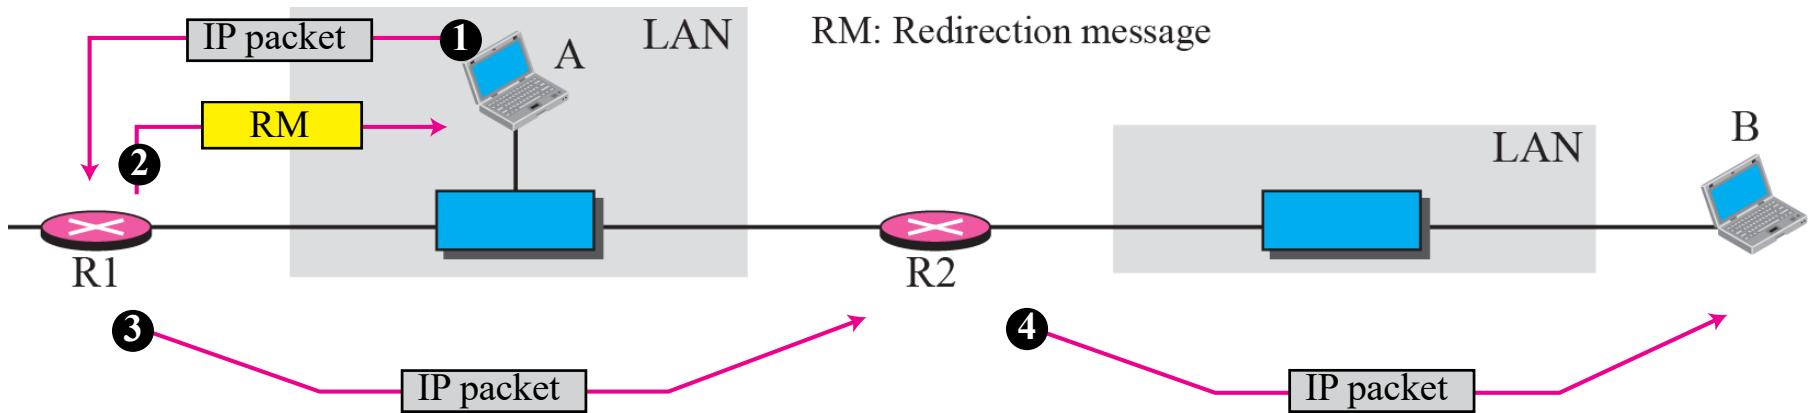


Note:

A host usually starts with a small (statically defined) routing table that is gradually augmented and updated. One of the tools to accomplish this is the ICMP redirection message.



ICMP error messages - Examples



Redirection concept

Type: 5

Code: 0 to 3

Checksum

IP address of the target router

Part of the received IP datagram including IP header
plus the first 8 bytes of datagram data



QUERY

ICMP can also diagnose some network problems through the query messages, a group of four different pairs of messages. In this type of ICMP message, a node sends a message that is answered in a specific format by the destination node.

ICMP query messages - Examples

The timestamp-request and timestamp-reply messages can be used to synchronize two clocks in two machines if the exact one-way time duration is known.

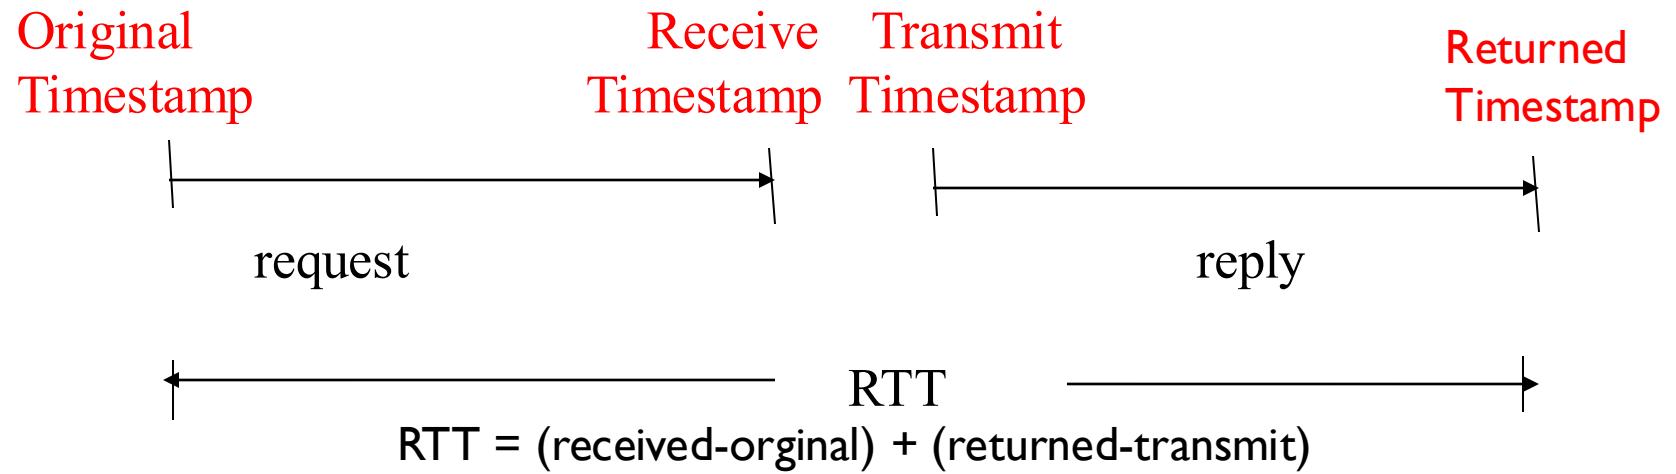
assignment

Type: 13 or 14	Code: 0	Checksum
Identifier		Sequence number
Original timestamp		
Receive timestamp		
Transmit timestamp		

13: request
14: reply

Timestamp-request and timestamp-reply message format

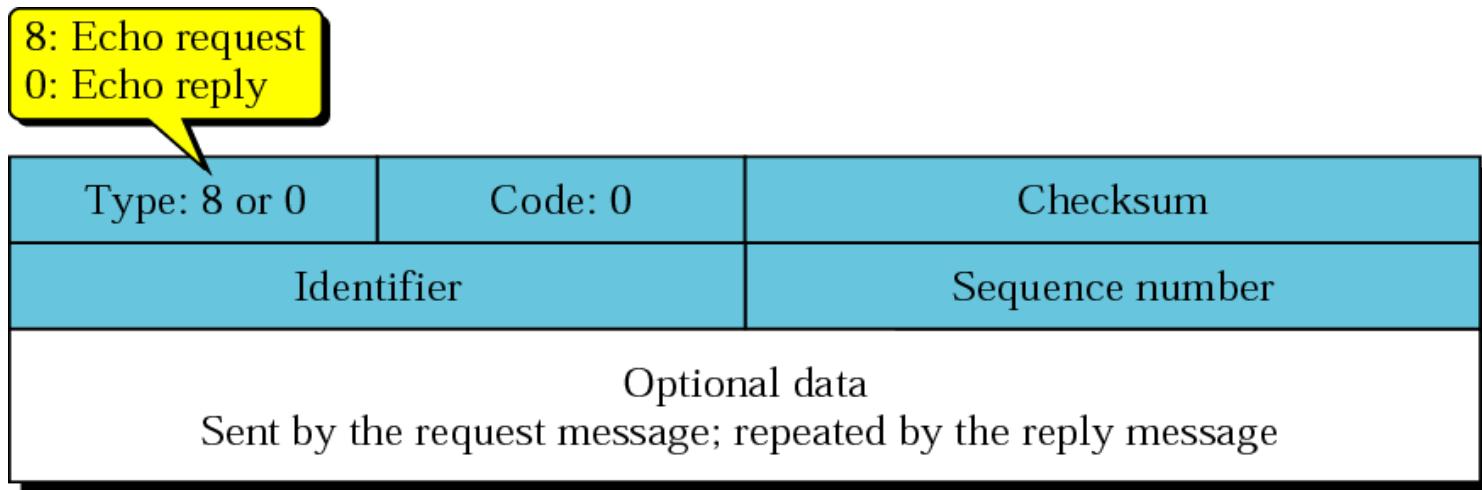




Q: If the two clocks are NOT synchronised, will the RTT be correct?



Echo-request and echo-reply messages



Echo-request and echo-reply messages can test the reachability of a host. This is usually done by invoking the **ping** command.



EXAMPLE 1

We use the ping program to test the server fhda.edu. The result is shown below:

```
$ ping fhda.edu
```

```
PING fhda.edu (153.18.8.1) 56 (84) bytes of data.
```

```
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=0 ttl=62 time=1.91 ms
```

```
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=1 ttl=62 time=2.04 ms
```

```
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=2 ttl=62 time=1.90 ms
```

```
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=3 ttl=62 time=1.97 ms
```

```
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=4 ttl=62 time=1.93 ms
```

See Next Slide



EXAMPLE (CONTINUED)

```
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=5 ttl=62 time=2.00 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=6 ttl=62 time=1.94 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=7 ttl=62 time=1.94 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=8 ttl=62 time=1.97 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=9 ttl=62 time=1.89 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=10 ttl=62 time=1.98 ms
```

--- fhda.edu ping statistics ---

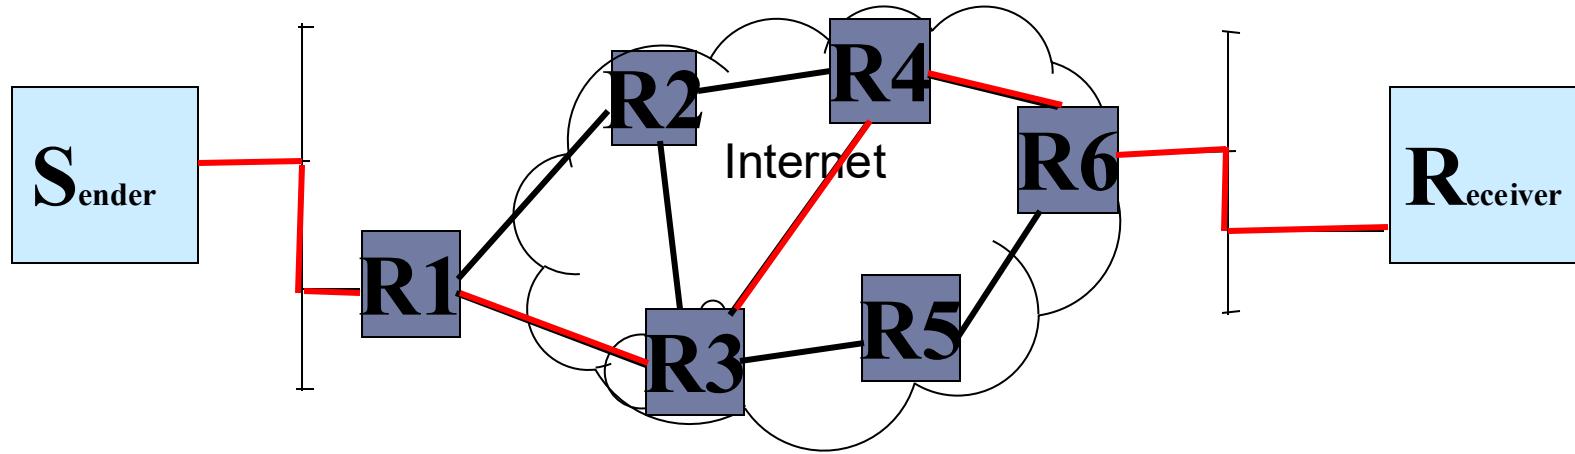
11 packets transmitted, 11 received, 0% packet loss, time 10103ms

rtt min/avg/max = 1.899/1.955/2.041 ms

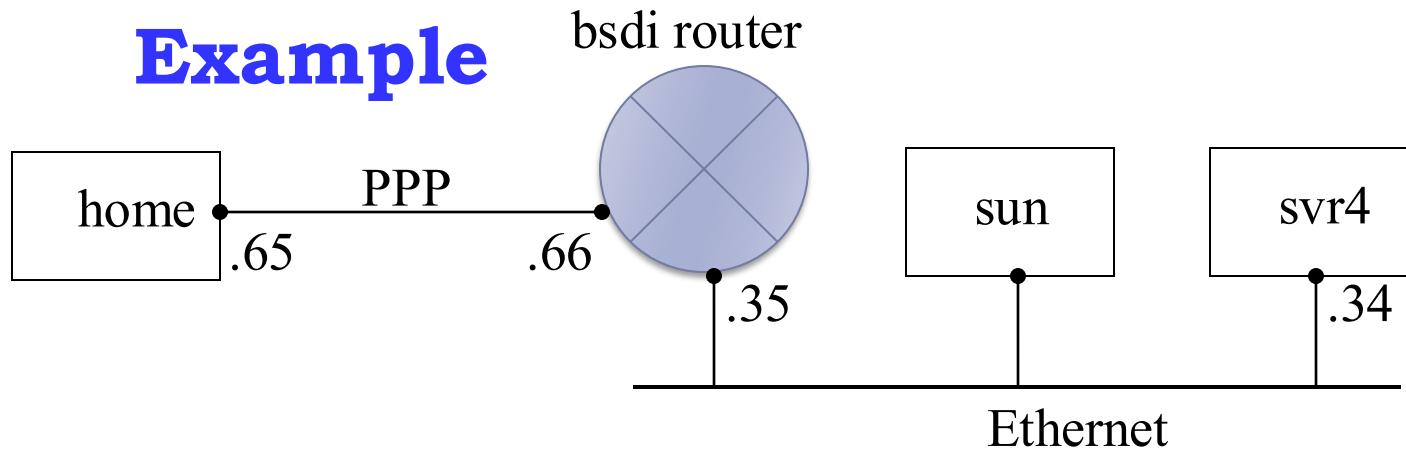


Traceroute Program (*cont*)

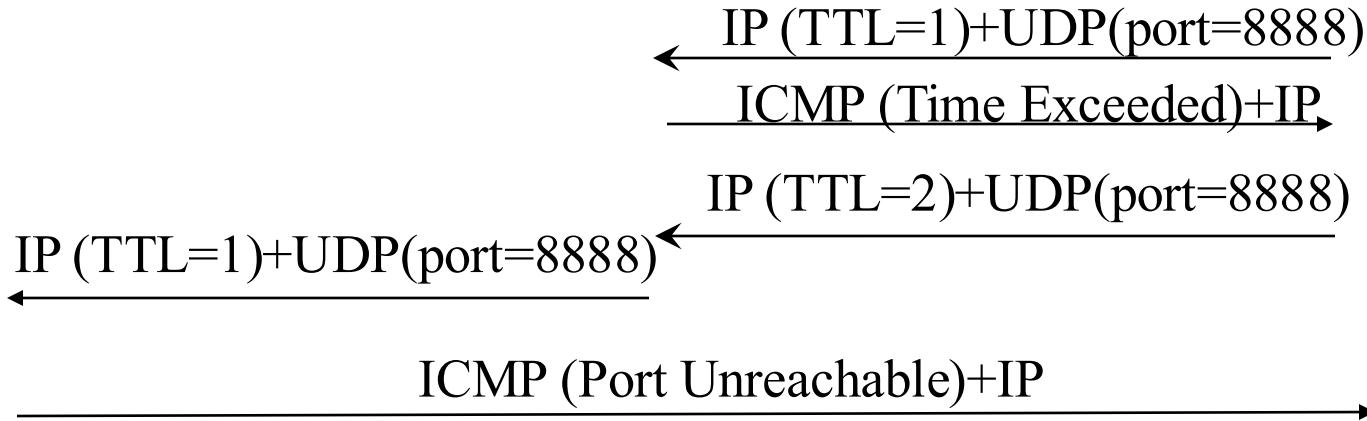
Allows us to see the route that IP datagrams follow from one host to another.



Example



assignment



svr4 % traceroute home

traceroute to home (140.252.13.65), 30 hops max, 40 byte packets

1	bsdi	(140.252.13.35)	20 ms	10 ms	10 ms
2	home	(140.252.13.65)	120 ms	120 ms	120 ms

Traceroute Operation

- ▶ Sends a UDP datagram to destination with TTL field in IP header set to value 1
 - ▶ datagram = 40bytes
(8 bytes UDP + 20 bytes IP + 12 bytes data)
 - ▶ 12 bytes data = seq no, time sent
- ▶ Causes router to generate “**time exceeded**” ICMP error
- ▶ Increment TTL progressively until final destination is reached
- ▶ How to terminate?
 - ▶ UDP port chosen is a non-existence port
 - ▶ Destination sends ICMP “**port unreachable**” error
- ▶ For each TTL value, 3 datagrams (probes) are sent. **WHY?**

Q: Some firewalls disable UDP messages. Is there any alternate method?

Q: Is there any difference between the route info shown by Traceroute and IP-record route?



CS3103 Computer Networks Practice

IP-options, ICMP, NAT-Hole Punching

IPv4 – Options Processing
Internet Control Message Protocol (ICMP)

Ping and Traceroute

**NAT Traversal (Port Forwarding and Hole
Punching)**

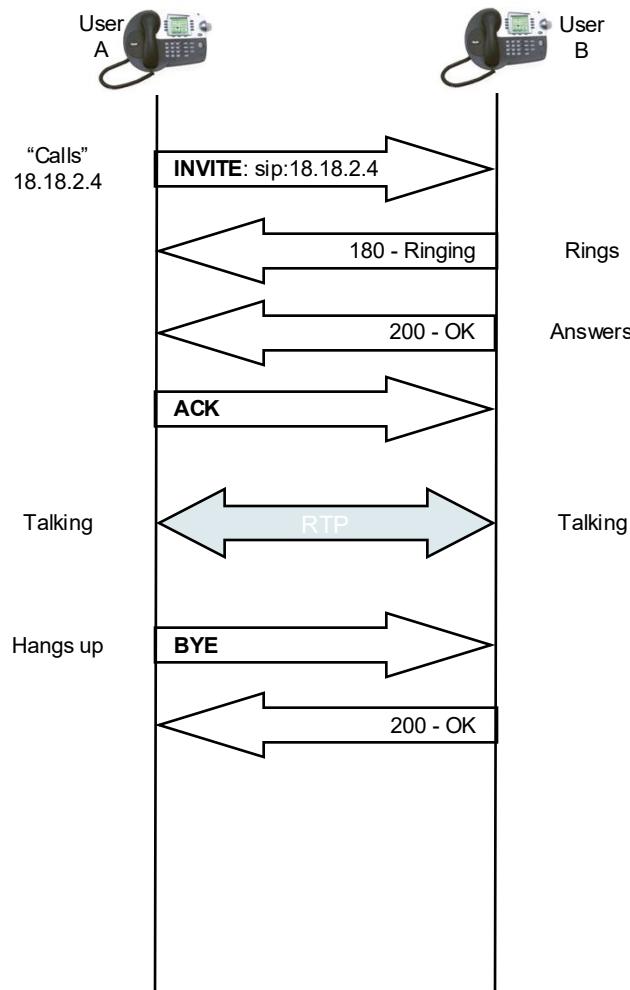
Dr Anand Bhojan

COM3-02-49, School of Computing

banand@comp.nus.edu.sg ph: 651-67351

A Simple SIP Flow

Session Initiation Protocol (SIP)



What if one or more peers are behind NAT (using private IP)?

NAT Traversal

NAT traversal & why do we need it?

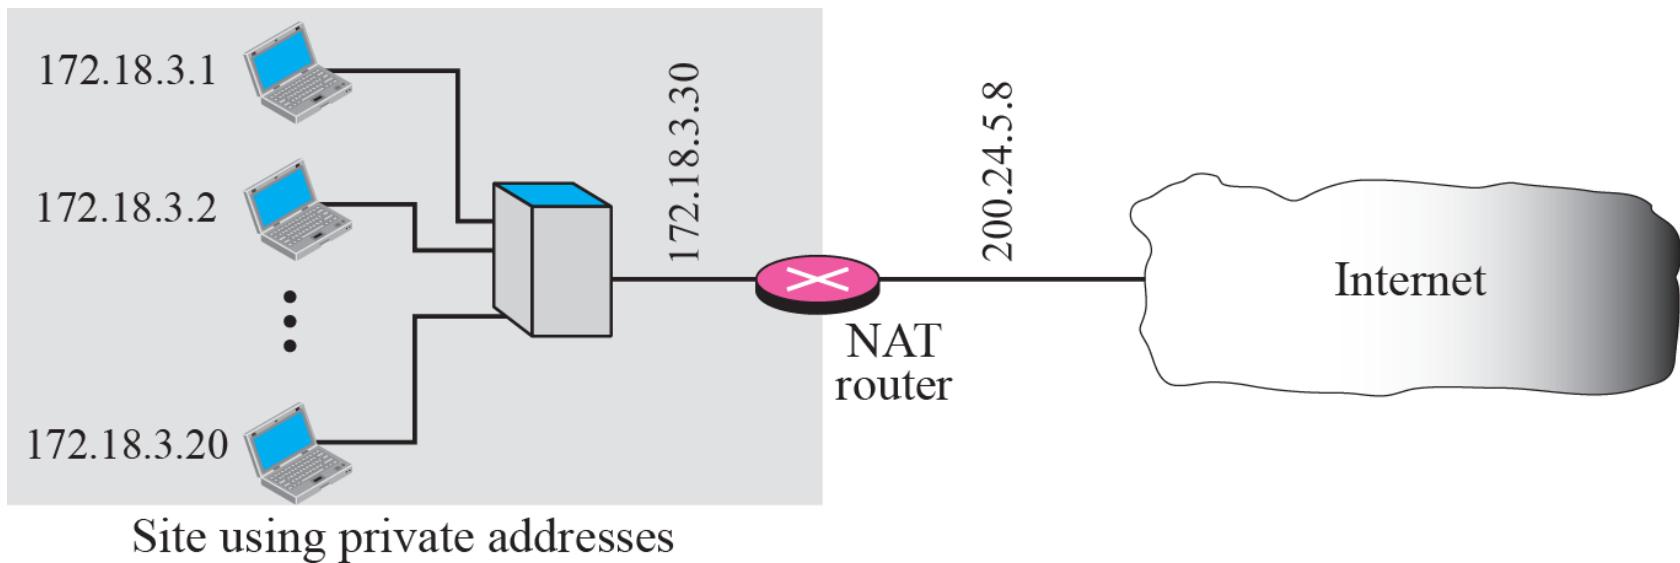
Network Address Translation (NAT)

◀ Remember NAT?

- Private IP addresses
- NAT gateway (usually on a gateway router)
- ← **Translates** between internal addresses/ports & external ones

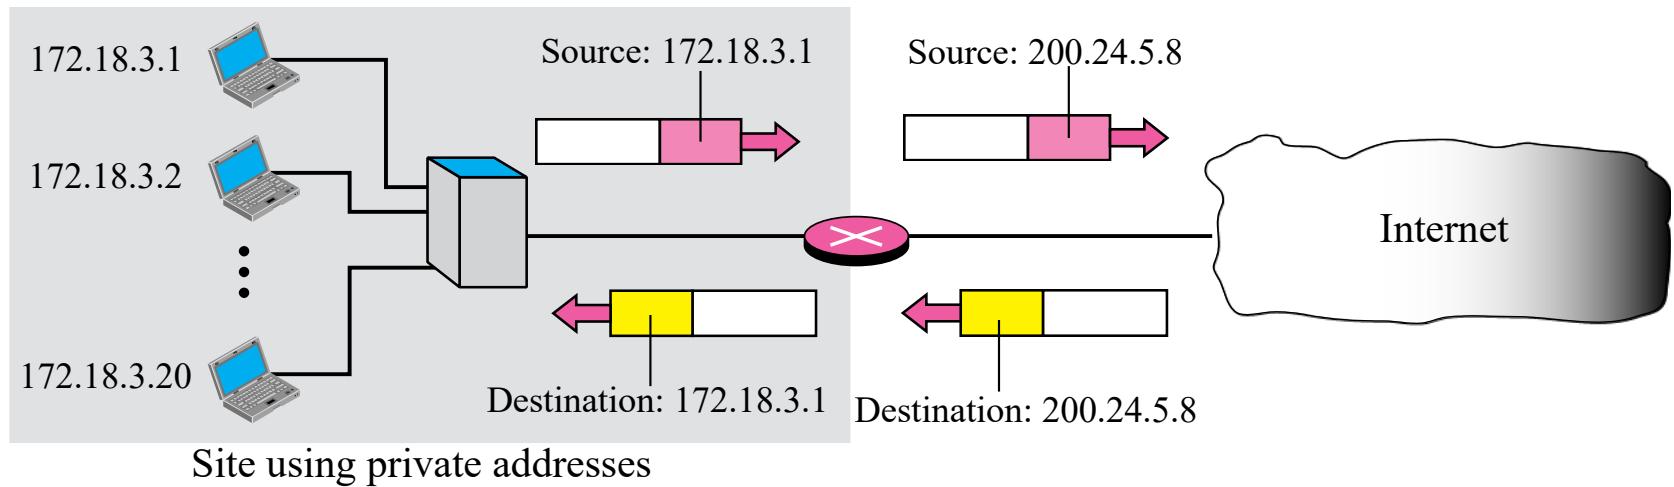
Network Address Translation (NAT)

Revision



Network Address Translation (NAT)

Revision

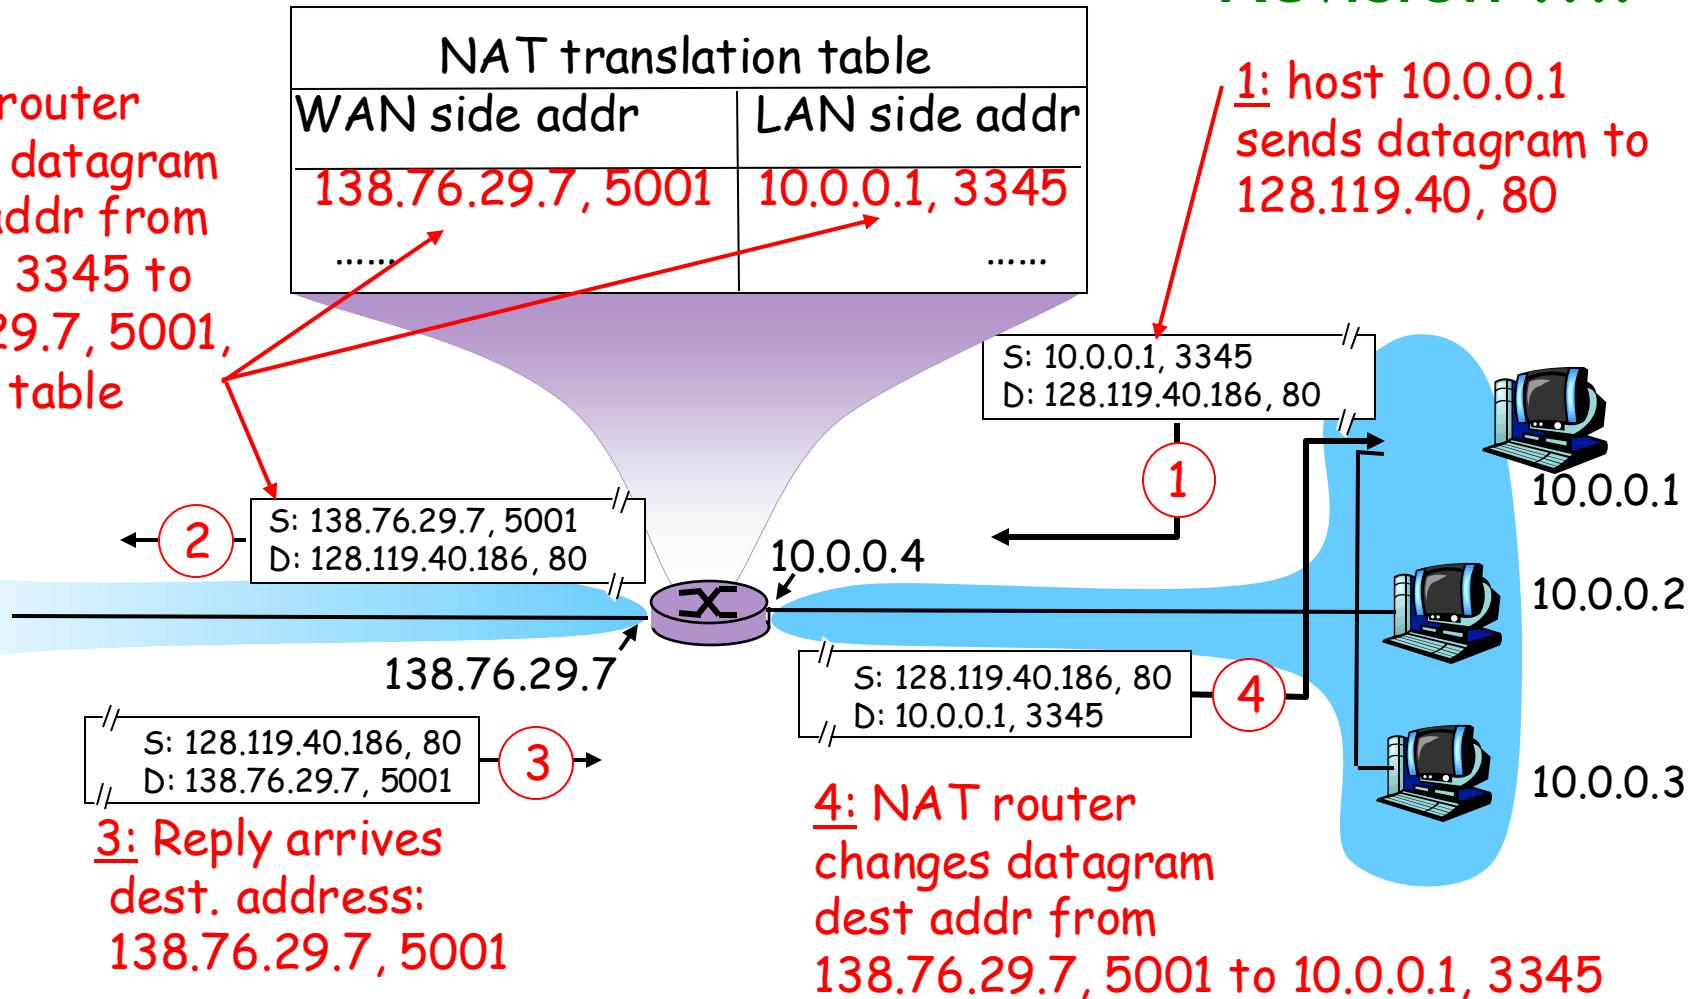


it simply maps the public ip to a private ip address alongside a port in the table.
so all the requests go through the same public ip address

NAT: Network Address Translation

Revision

2: NAT router changes datagram source addr from 10.0.0.1, 3345 to 138.76.29.7, 5001, updates table



1: host 10.0.0.1 sends datagram to 128.119.40, 80

4: NAT router changes datagram dest addr from 138.76.29.7, 5001 to 10.0.0.1, 3345

NAT Table

Revision

Five-column translation table

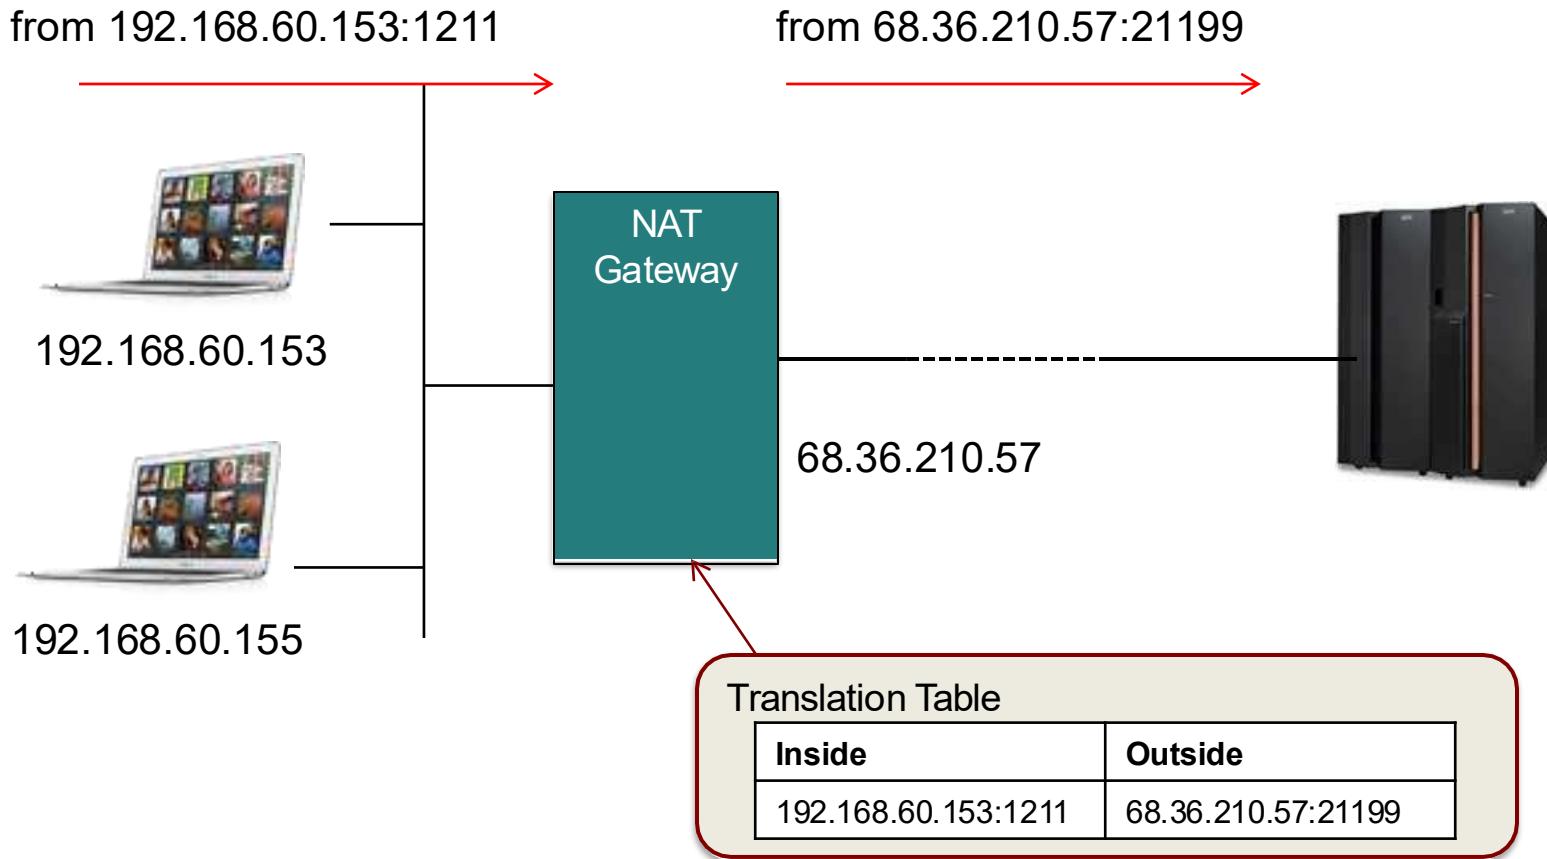
<i>Private Address</i>	<i>Private Port</i>	<i>External Address</i>	<i>External Port</i>	<i>Transport Protocol</i>
172.18.3.1	1400	25.8.3.2	80	TCP
172.18.3.2	1401	25.8.3.2	80	TCP
...

NAT traversal & why do we need it?

- NAT is awesome!
 - Maps multiple local private addresses to a single public IP address.
You don't need many public addresses!
 - **External hosts only see a single IP device** and do not have visibility into a network's internal/private structure
- **But** it breaks end-to-end connectivity!
 - Host behind NAT box can access services in the internet, but **hosts in the internet cannot (easily) reach hosts behind NAT box.**
 - Consider two **VoIP clients behind NAT** that want to communicate



NAT: Accessing external public IP server - Normal Operation - This is easy



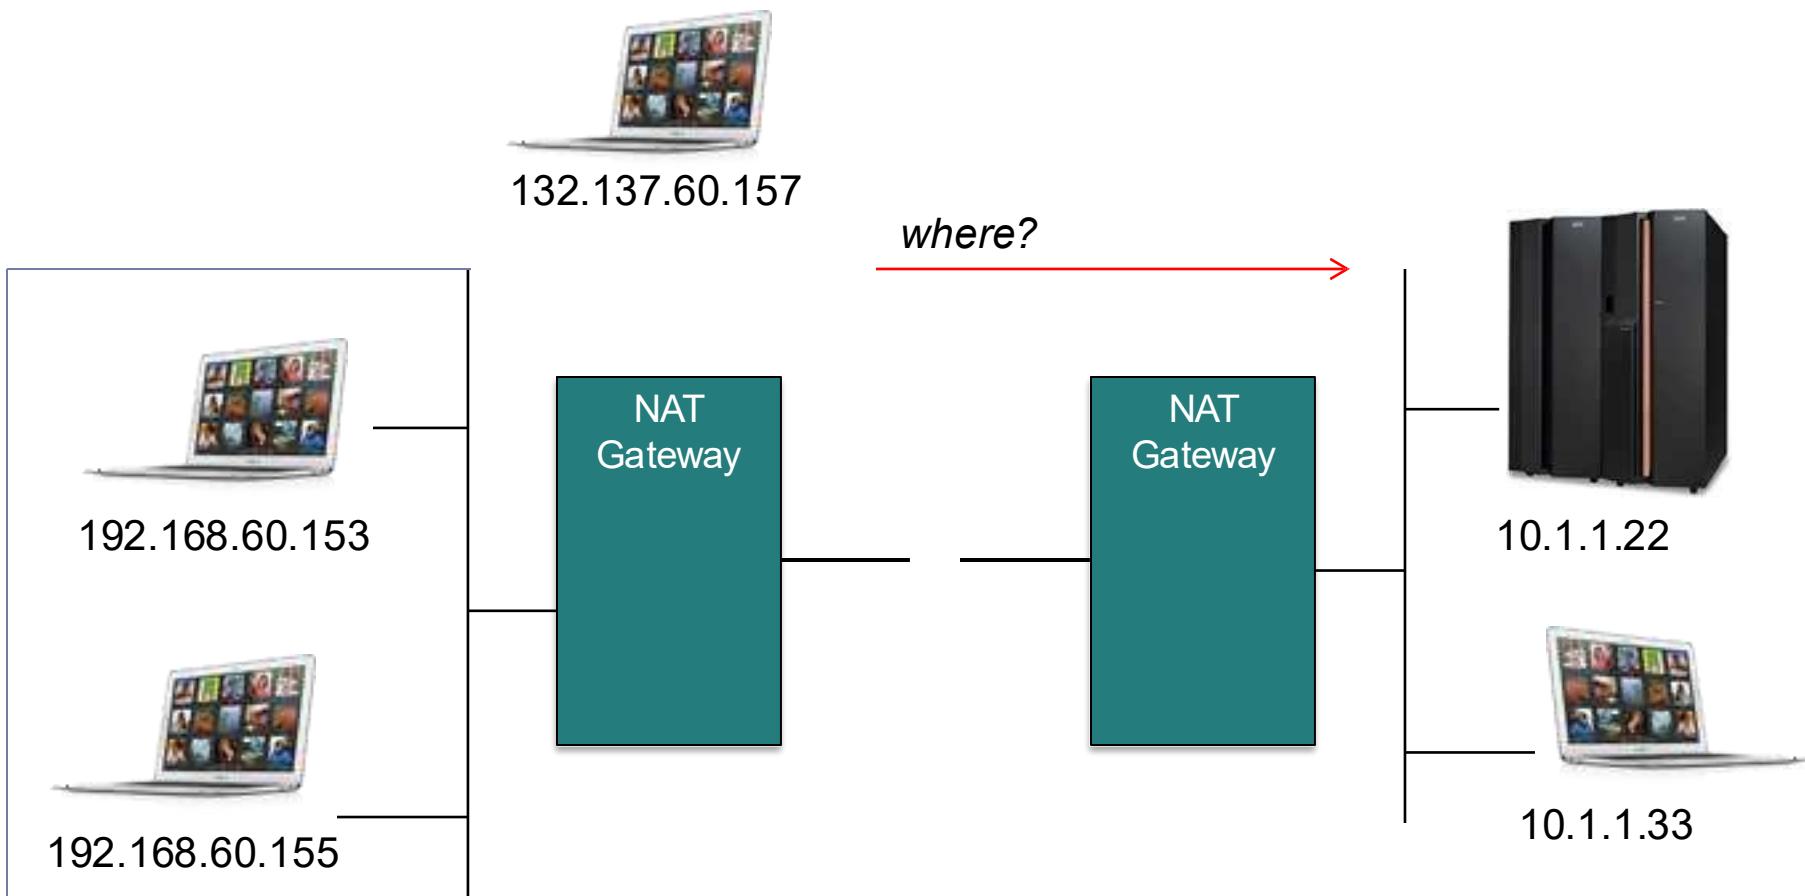
SIMPLE, on the fly **address mapping!**

“Normal” operations in NAT box

- ▶ Mapping is “reactive”,
 - ▶ X ($\text{SIP:SP}=\underline{192.168.31.3:5555}$, $\text{DIP:DP}=34.56.78.9:80$)
 - ▶ NAT entry is setup by the first packet from X to outside world
 - NAT ($\text{SIP:SP}=\underline{192.168.31.3:5555} > \text{132.137.3.13:2345}$,
 $\text{DIP:DP}=34.56.78.9:80$)
 - ▶ Expects a return packet with the source and destination information reversed.
 - NAT Entry ($\text{SIP:SP}=34.56.78.9:80$,
 $\text{DIP:DP}=\text{132.137.3.13:2345} > \underline{192.168.31.3:5555}$)

NAT: Accessing a server behind NAT- This is tricky

why is this so? isn't it just like sending from host to host? what's the issue?



NAT ISSUES....

so the issue is the inbound ?

- ▶ **Issue 1:** You need to access a server behind the NAT gateway from the Internet
 - ▶ Solution 1: Port Forwarding - Define a specific mapping for incoming packets (in advance) - **See lab 7**
 - ▶ Solution 2: Connection Reversal
- ▶ **Issue 2:** Both hosts are behind NAT boxes
 - ▶ Solution 1: Relay all with a Proxy
 - ▶ Solution 2: Hole Punching. - **See lab 7**

NAT Traversal Techniques

Port Forwarding (See lab 7)
Hole Punching (See lab 7)

Port Forwarding. - See lab 7

I. Define Mapping in advance (**Port Forwarding**):

- ▶ ~~X~~ (192.168.31.3) runs a TCP server listening at port 5555
 - ▶ Known to the outside world as **132.137.3.13:443**
 - ▶ We **set up this translation** in the NAT box **in advance.**
- ▶ NAT (SIP:SP=*, DIP:IP=**132.137.3.13:443** -> DIP:DP=192.168.31.3:5555)
- ▶ **What's next?** Anything to do with DNS?
 - ▶ Buy a domain and insert a 'A' record with **132.137.3.13:443**.

2. Another Solution: **Connection Reversal**

Connection reversal

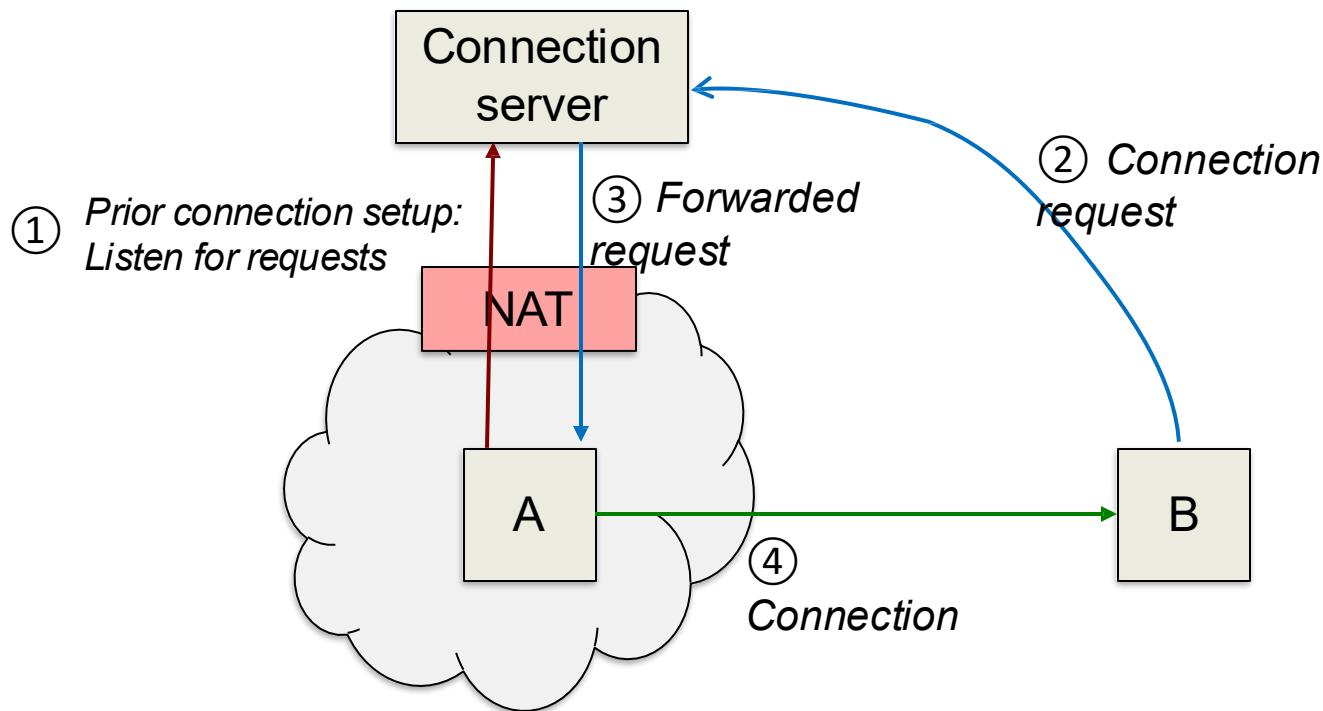
- ▶ B wants to connect to A
 - ▶ But A is behind a NAT
- ▶ Somehow get B to send a message to A,
 - ▶ Ask A to open a connection to B
- ▶ Two approaches
 - ▶ Relay the request via a server (but A must be connected to the server)
 - ▶ Assume an existing connection between A & B and ask for a new one

so does this mean for example, there is no current NAT entry on the router with mappings to A from the public IP, hence a connection server is used, once the server receives a connection request from B, it tells to A as it is listening, then A will send a request to B and map on the NAT table as well to receive the connection from B? So first, why doesn't A directly create the mapping on the NAT and instead requires a connection server. And what does B receive?



Connection reversal

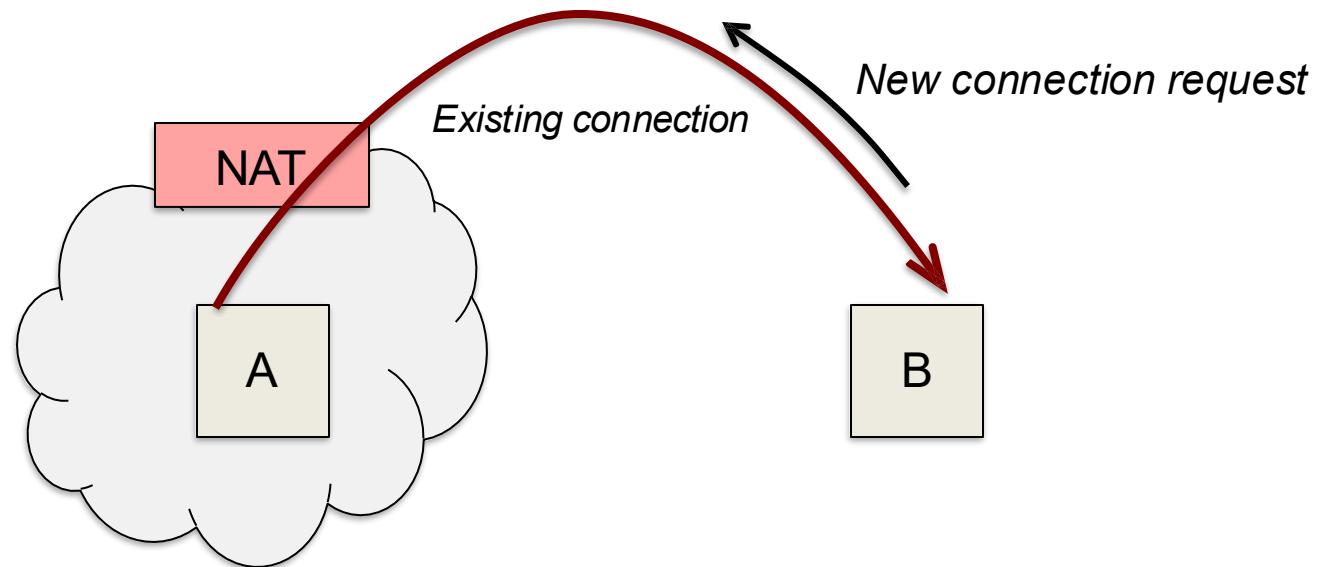
Use a server for sending only connection requests



Connection reversal

B wants to talk to A

There is an existing connection between A & B
(set up by B previously)



Hole Punching - See lab 8

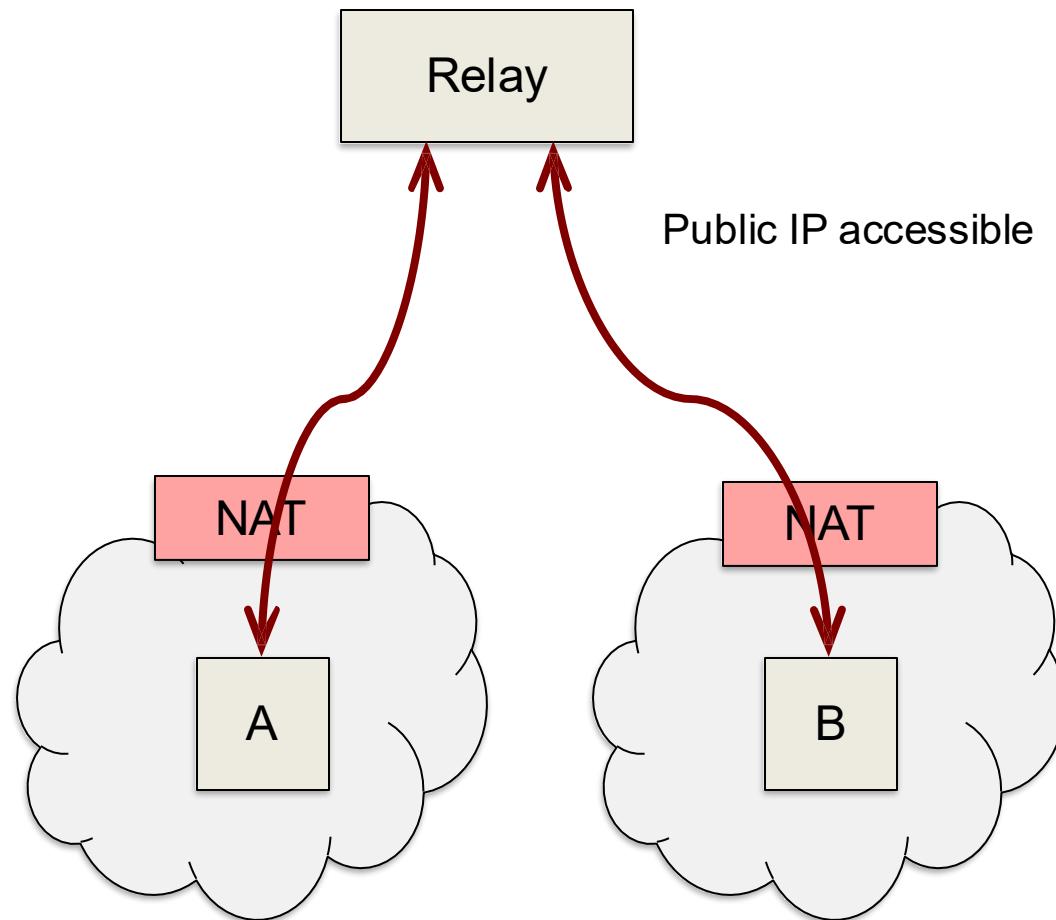
- ▶ Devices A and B are behind NAT gateways
 - ▶ Normal and port forwarding do not work!
1. Simple Solution - **Relay** : Use a proxy server to relay messages
 2. Better Solution - **Hole punching**: Use a rendezvous server (external controller) with known public **IPaddress:port** to inform both devices the NAT IP address and port to try

Relay all messages through a Proxy

- Hosts A & B want to communicate
- Have an Internet-accessible proxy, P
- A connects to P and waits for messages on the connection
- B talks to P; P relays messages to A
- Most reliable but not very efficient
 - Extra message relaying
 - Additional protocols needed (e.g., B needs to state what it wants)
 - Proxy can become a point of congestion (network links & CPU)



Relay all messages

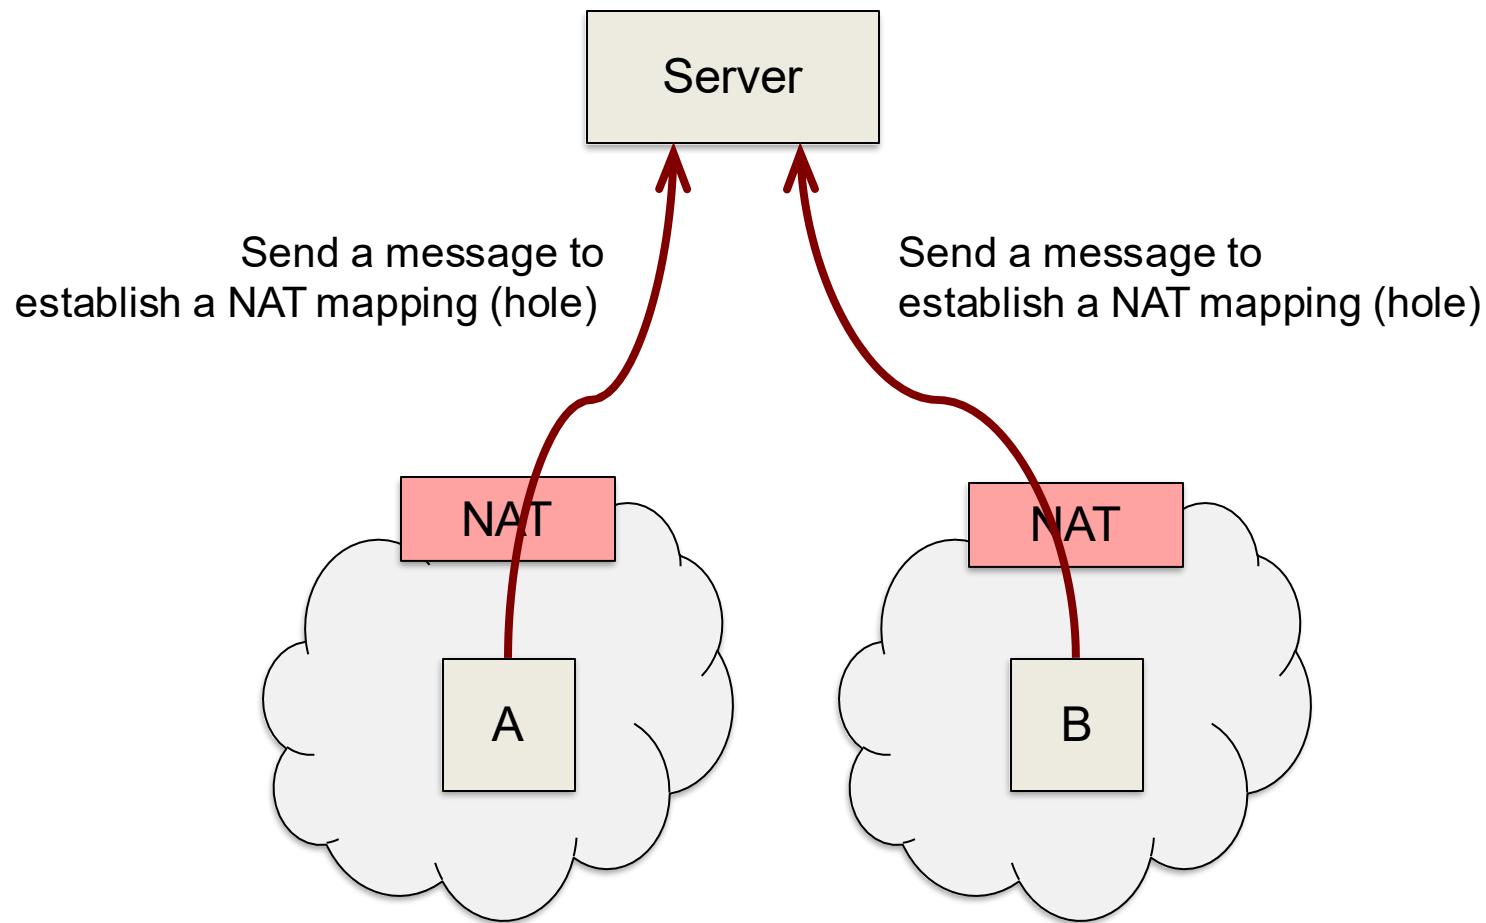


UDP hole punching

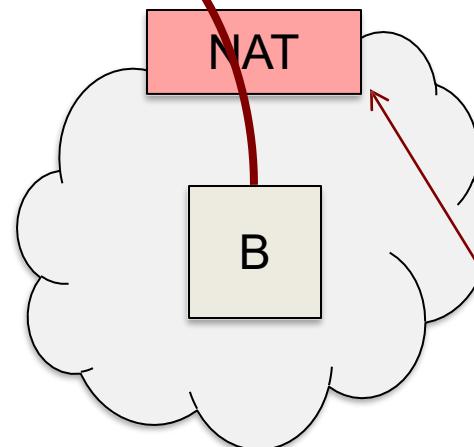
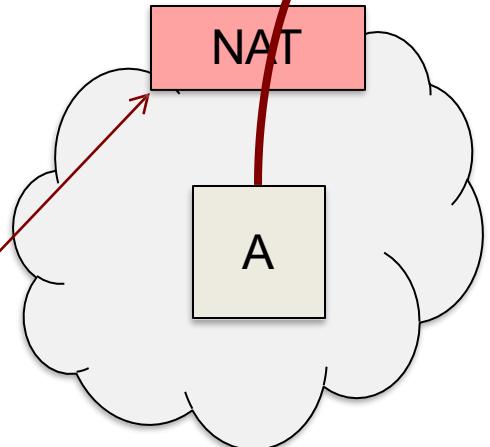
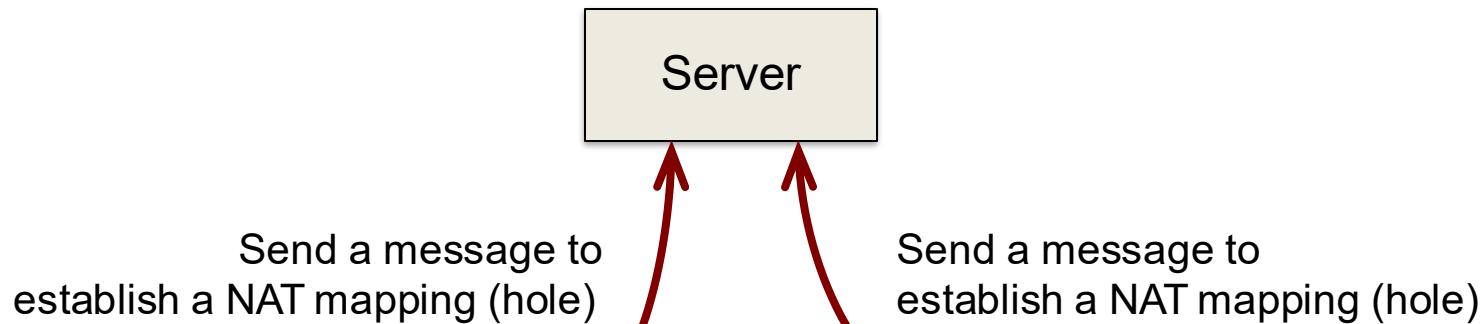
- Hosts A & B want to communicatee
- Have an Internet-accessible **rendezvous server** (external controller), S
- Host A sends a message to S
 - That sets up a NAT translation on A's NAT gateway
 - S now knows the external host & port
- Host B sends a message to S
 - That sets up a NAT translation on B's NAT gateway
 - S also knows the external host & port on B
- **S tells B:** *talk on A's IP address & port*
- **S tells A:** *talk to B's IP address & port*



UDP hole punching



UDP hole punching



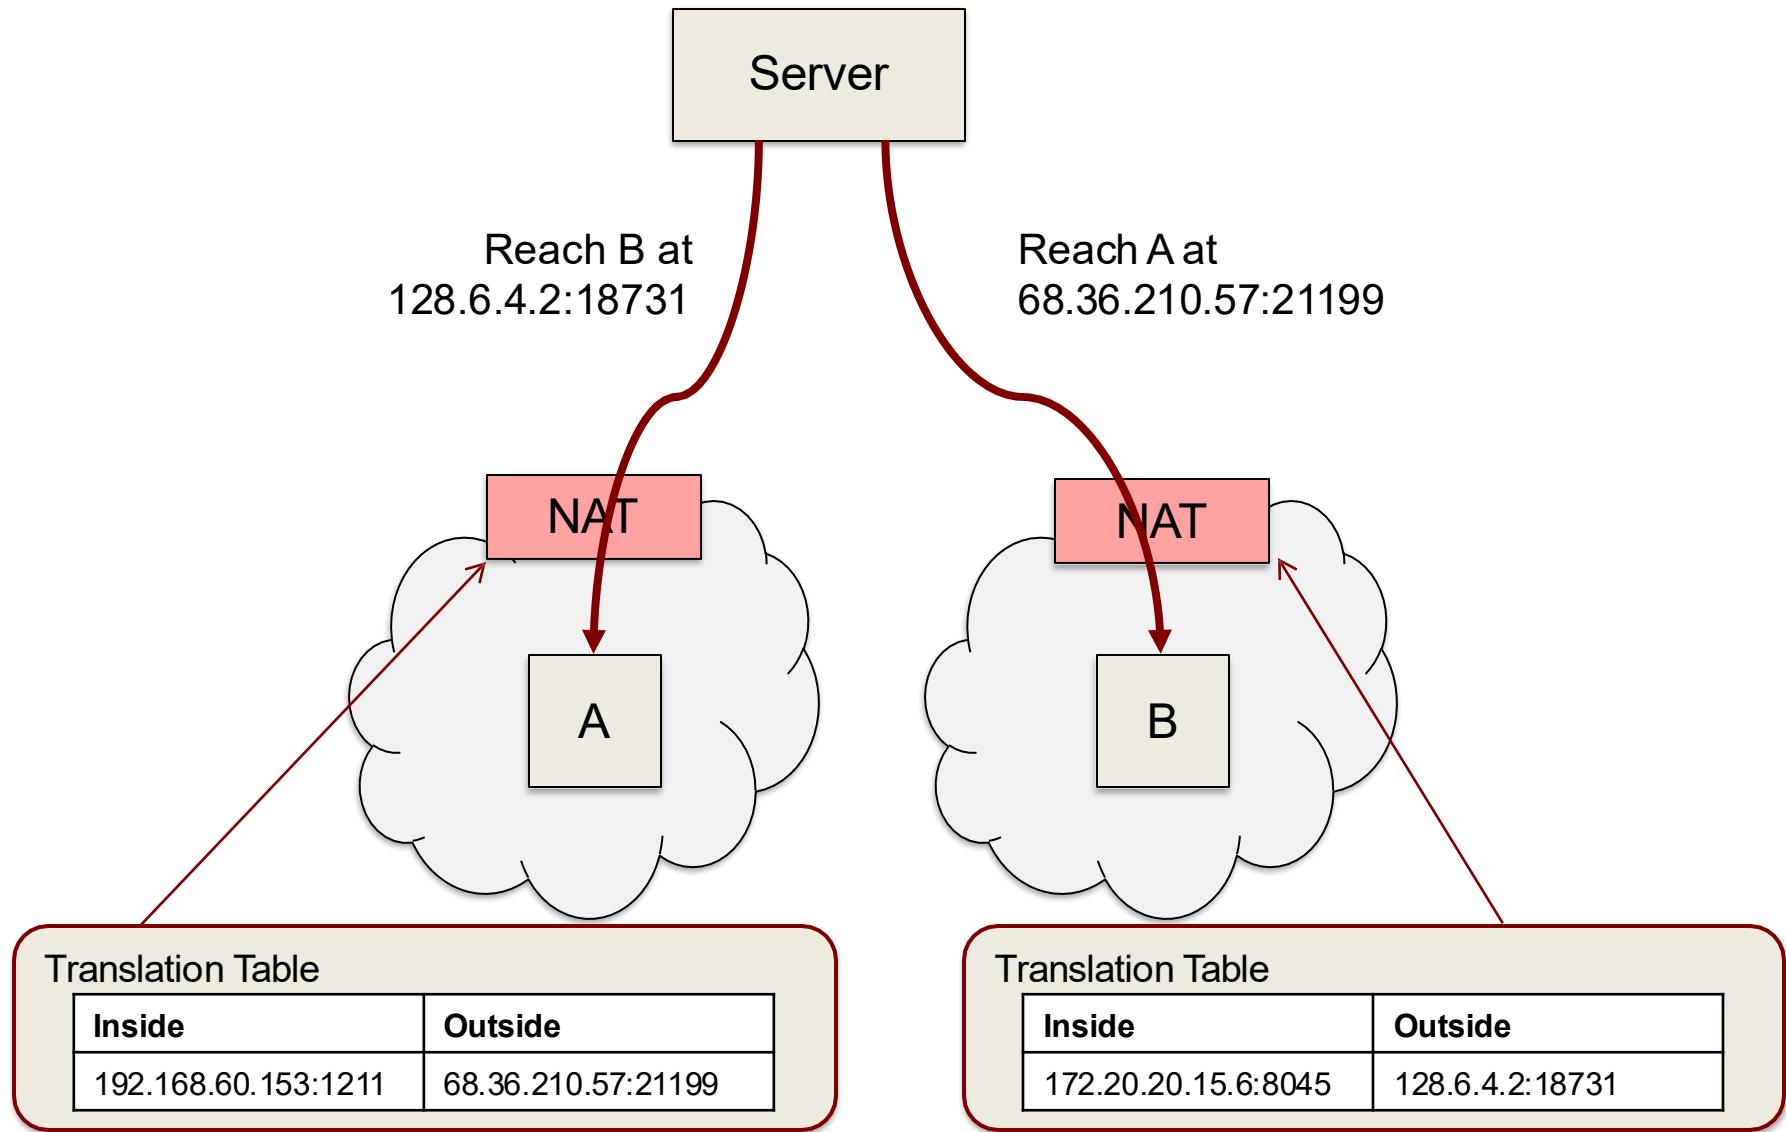
Translation Table

Inside	Outside
192.168.60.153:1211	68.36.210.57:21199

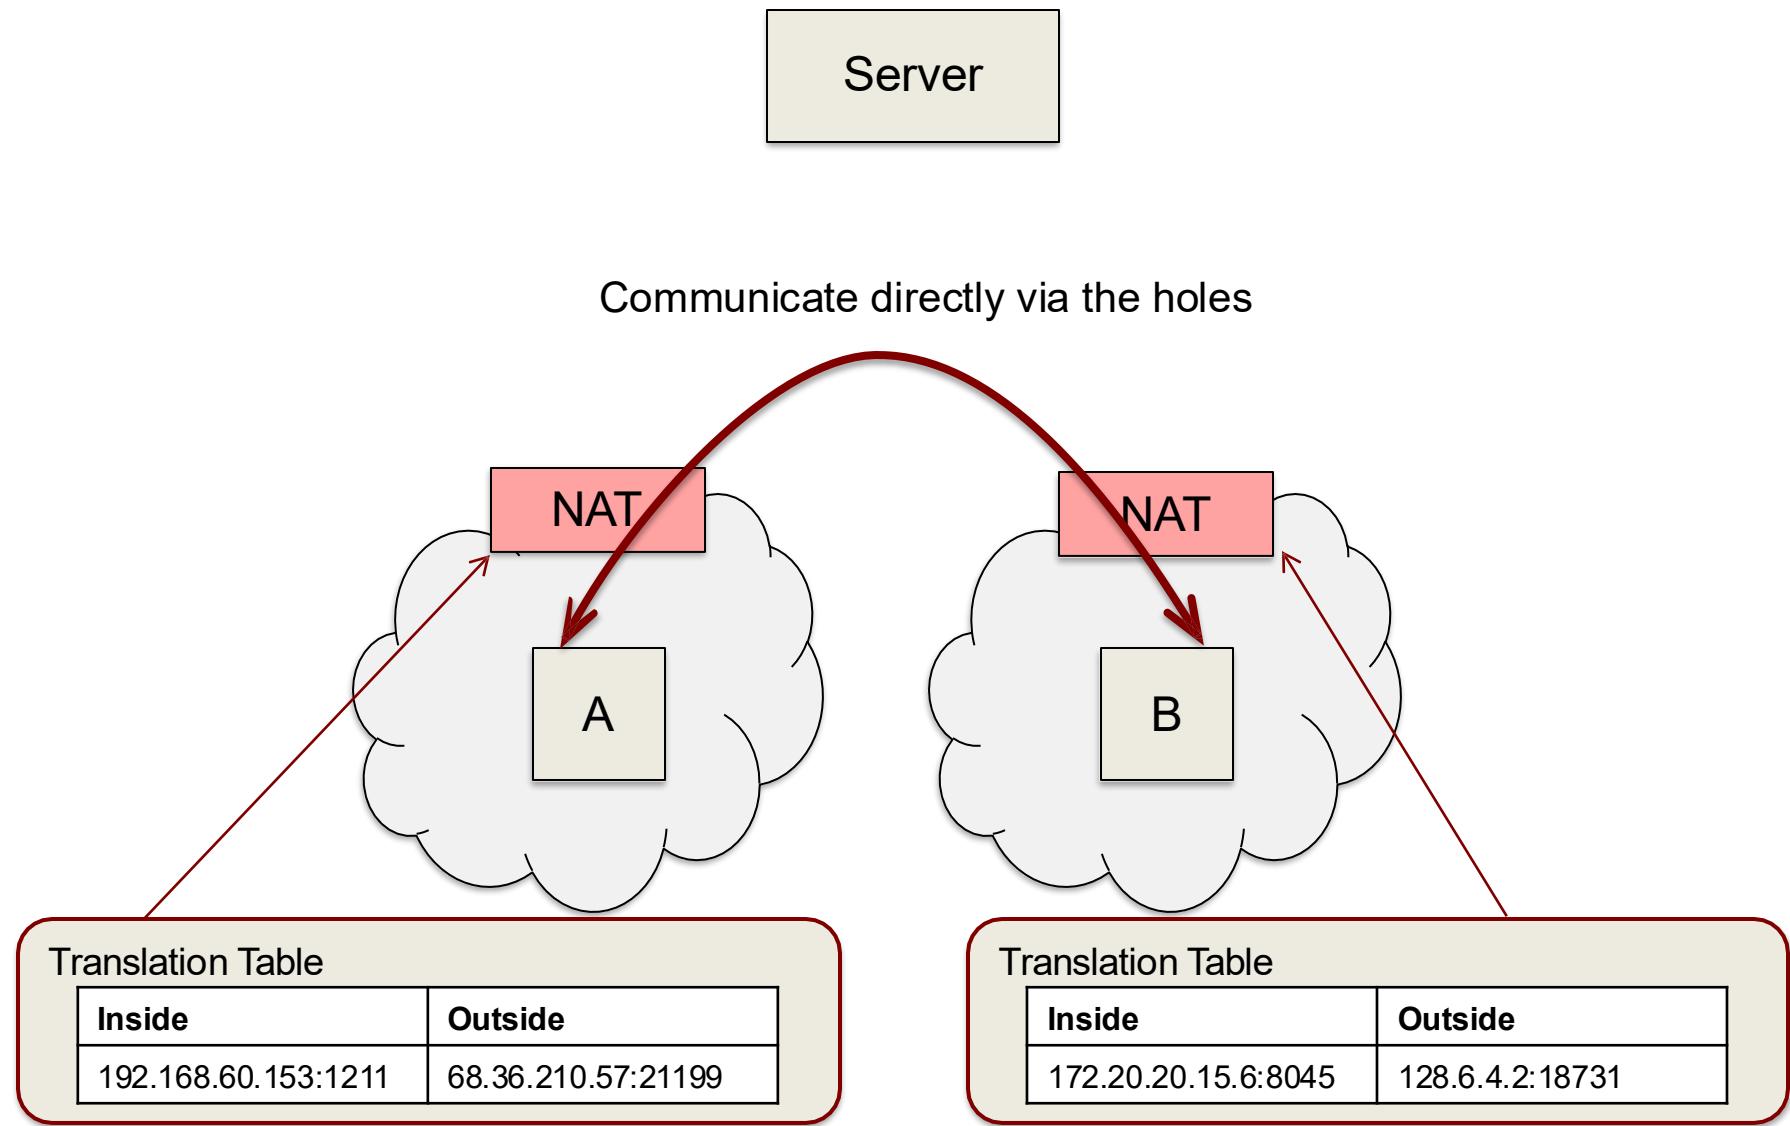
Translation Table

Inside	Outside
172.20.20.15.6:8045	128.6.4.2:18731

UDP hole punching



UDP hole punching



TCP hole punching

- Same principle (tell other host of your address:port) – BUT
 - Use **TCP Simultaneous Open**
 - Both hosts will try to connect to each other
 - Each NAT creates a translation rule
 - At least one of the SYN messages during connection set up will go through the NAT translation on the other side
 - The remote side will send a SYN-ACK
 - Need to re-use the same port # that the remote side knows about
 - Socket option to reuse an address:
SO_REUSEADDR
 - Not guaranteed to work with all NAT system

FURTHER READING: <https://bford.info/pub/net/p2pnat/>

Questions to Ponder:

Q: Why simultaneous Open?

Q: When (which type of NAT system) TCP hole punching will fail?



Summary

- ▶ IPv4 Format
- ▶ Fragmentation
- ▶ IP-Options: Record Route
- ▶ IP-Options: Strict/Loose Source Route
- ▶ ICMP – Error Message
 - ▶ Destination unreachable, Source-quench, time-exceeded, redirect.
- ▶ ICMP – Query Message
 - ▶ Echo req/rep (**ping**), Time Stamp req/rep.
- ▶ Trace Route
- ▶ NAT Traversal -- Port Forwarding & Hole Punching

END